

IBM watsonx Assistant for Z for Technical Sales Level 4 Lab Guide

Description	IBM watsonx Assistant for Z for Technical Sales Level 4 Lab Guide
Author(s)	Andrew R. Jones (andrewj@us.ibm.com) and Max Weiss (maxwell.g.weiss@ibm.com)
Copyright	Copyright © 2024 IBM - V 2.0

Table of Contents

Welcome

- Support
- Using the demonstration guide

IBM Technology Zone environment

- Create a reservation request
- Extend the reservation
- Join the
- Accessing the environments
- Troubleshooting

Pilot setup

Creating an assistant and configuring conversational search

- Access the
- Create your Assistant
- Configure conversational search
- More configurations
- Configure the base large language model
- Testing conversational search
- Troubleshooting

Ingesting client documents

Creating a stand-alone OpenSearch instance for document ingestion

- Install the Red Hat OpenShift command-line interface utility
- Prepare to ingest documents
 - Login to the OpenShift cluster from your local terminal
 - Create a working directory
 - Install IBM Certificate Manager on Red Hat OpenShift
 - Install the watsonx Assistant for Z Operator (for OpenSearch)
 - Deploy required secrets and the custom bring-your-own-search (BYOSearch) resource
 - Verify all the required pods are running
 - Update your assistant with the new
- Troubleshooting

Installing and using zassist to ingest client documents

- Install the zassist utility
- Ingest client documentation using
- Verify the document that is ingested is now returned as a source file for a query.

- Adjusting the search behavior
- Guidance for the pilot environment

Adding skills to the assistant

Getting started with skills and actions

- Environments
 - Watsonx Orchestrate
 - Ansible Automation Platform and Wazi as a Service

Explore Ansible Automation Platform

- Access the

Importing skills from Ansible Automation Platform

- Import skills into your assistant

Connecting skills to your assistant

Creating actions for your assistant

Configure the number of input fields

- Create actions

Verify the job in the Ansible Automation Platform console

Creating skill flows

Add the utility skill

- Add the skills to your Personal skills
- Create the skill flow
- Map the outputs and inputs of the skills
- Enhancing the skill flow
- Enable the skill flow in your assistant

Creating custom-built actions

Importing pre-packaged z/

Publishing and deploying your assistant

- Publish the assistant
- Configure the
- Deploy the assistant

Adding other integrations

Next steps

- Earn the badge

Welcome

Welcome to the IBM watsonx Assistant for Z for Technical Sales Level 4 Lab Guide (lab guide). The lab guide is part of the IBM watsonx Assistant for Z for Technical Sales Level 4 learning plan for IBM and Business Partner Technical Sales and related badge. The learning plan is intended to teach technical sellers and Business Partners how to conduct a proof of experience (PoX) for a client.



In a fast-paced world, things change

The products and services may appear differently than what is shown in the lab guide. This can occur if the product or service is updated with a new version.

Responses generated by IBM watsonx Assistant for Z are likely to change over time. The responses you see when you run the queries in this section may differ from the screen images captured in the lab guide.

This lab guide covers the setup, configuration, and usage of watsonx Assistant for Z. This lab guide uses the [IBM watsonx Assistant for Z Velocity collection](#) and the 3 Velocity Pilot lab environments in IBM Technology Zone (ITZ).

The lab guide also enables dedicated lab environments for customized client Proof of Experiences (PoX) and demonstrations. If you are preparing for an actual pilot engagement, refer to the [Pilot Scoping Guide for watsonx Assistant for Z](#) for additional information.

The lab guide provides guidance to:

- Provisioning the lab environments
- Creating an assistant and configuring conversational search
- Configuring assistant settings
- Testing conversational search
- Deploying a dedicated instance of OpenSearch for client document ingestion (Optional)
- Importing skills for z/OS automations
- Connecting apps to assistants
- Creating assistant actions
- Creating skill flows and custom-built actions
- Importing pre-packaged z/OS skills
- Publishing and deploying your assistant

Support

Think something is down? Check the applicable status pages for any known issues like a site or service not available:

- [IBM Technology Zone](#)

For issues with provisioning the ITZ environment for this lab (for example, a failed reservation request due to insufficient quota capacity) open a ticket with ITZ support:

- Web: [IBM Technology Zone](#)
- Email: techzone.help@ibm.com

For issues related to specific steps found in the demonstration guide after the ITZ environment is provisioned, contact the authors:

- Slack:
- [#watsonx-assistant-z-technical](#) - IBM only
- [#wxo-practitioners](#) - IBM only - for questions that are related to the software as a service (SaaS) instance of watsonx Orchestrate
- Email: andrewj@us.ibm.com; maxwell.g.weiss@ibm.com

Business Partners, use the IBM Training live Chat Support service or other support methods that are found on the IBM Training portal [here](#).

Using the demonstration guide

Use these helpful tips to take full advantage of the IBM watsonx Assistant for Z for Technical Sales Level 4 Lab Guide.

🖨 Printing the demonstration guide

⚠ Printed or saved copies can be out of date

The IBM watsonx Assistant for Z for Technical Sales Level 4 Lab Guide changes regularly to match the IBM watsonx Assistant for Z offering and associated ITZ environment. Printed or saved copies of the demonstration guide can become out-of-date quickly and result in failed steps.

A ready-to-print PDF version of the IBM watsonx Assistant for Z for Technical Sales Level 4 Lab Guide is [here](#).

🖼 Viewing images

Images in the demonstration guide can be enlarged by clicking on the image. Press the **Esc** key or click the **X** to dismiss the enlarged image.

Reserve the IBM Technology Zone environment

IBM watsonx Assistant for Z for Technical Sales Level 4 Lab Guide

Welcome

Reserve the IBM Technology Zone environment

Pilot setup

Next steps

2. Click **Reserve now**.
The **Reserve now** option creates a reservation for immediate use. Optionally, schedule the reservation for a later date, like when you will be at your client's office.



3. Complete the reservation request and click **Submit**.
The first two reservations will be similar to the first image below and have fields **a-e** that will need to be completed.

- a. Optionally, change the **Name** field for the reservation.
- b. Select the **Education** purpose tile.
- c. Enter a **Purpose description**.
- d. Select the region nearest your physical location in the **Preferred Geography** drop-down.
- e. The **End date and time** will be set to 2 days after the current date and time.
- f. Accept the IBM Technology Zone's terms and conditions and security policies.
- g. When satisfied with the parameters, click **Submit**.



Image highlighting

In some images, the following styles of highlighting are used:

- **Solid highlight box:** This style of box highlights where to click, enter, or select an item.



- **Dash highlight box:** This style of box highlights one of two things: the path to follow to get to a specific location in the user interface, or areas to explore on your own.



Copying commands and prompts

Copying and pasting commands and prompts from this demonstration guide is easy and can eliminate typographical errors.

Click the highlighted copy icon and then use your operating system's paste function. For example, ^ Ctrl + V or right-click and select Paste.

Prompt 1

What is the APF list in z/OS? Provide a detailed explanation.





Acronyms and terminology



IBM employees, and the tech industry in general, enjoy using acronyms. In the demonstration guide, most acronyms will appear with a dashed underline. Hover over the acronym to learn its meaning. A question mark (?) icon will first appear and after a second the tool tip with the acronym's meaning is displayed. Try it here: LPAR.

Guidance for delivering a demonstration

IBM watsonx Assistant for Z

Welcome

Reserve the IBM Technology Zone environment

Guidance for delivering a demonstration

Demonstration scenarios >

Next steps

Specific guidance for IBM watsonx Assistant for Z and the ITZ environment

Following the scripts provided in the IBM watsonx Assistant for Z for Technical Sales Level 3 Demonstration Guide will help guarantee a successful demonstration. Use these tips to help insure success with IBM watsonx Assistant for Z and the ITZ environment:

- Follow the scripts in the IBM watsonx Assistant for Z for Technical Sales Level 3 Demonstration Guide for the automations and skills to execute as expected.

Table of contents

General demonstration guidance

Specific guidance for IBM watsonx Assistant for Z and the ITZ environment



The Lab Guide table of contents

▼

This **Demonstration Guide** uses a responsive browser-based interface to ensure a pages are usable on various devices with different screen sizes. The Demonstration Guide table of contents may be displayed as highlighted in the green dashed box in this image:

The screenshot shows a browser window for the "IBM watsonx Assistant for Z for Technical Sales Level 3 Demonstration Guide". The title bar includes the course name, a refresh icon, and a search bar. On the left, a green dashed box highlights the Table of Contents sidebar. The sidebar lists chapters: Welcome, Guidance for delivering a demonstration, Reserve the IBM Technology Zone environment, Demonstration scenarios (with sub-points: Introduction to scenarios, Scenario 1: Authorized Program Facility, Scenario 2: Certificate renewal, Scenario 3: Db2 versioning, Scenario 4: Initial Program Load on Z, Additional IBM Z related prompts, Summary, Next steps, Printing the guide, and Instructions). The main content area displays the "Welcome" page, which introduces the course goal of providing knowledge and tools for demonstrating IBM watsonx Assistant for Z. It also notes that the course focuses on end-user querying and automation skills, while other advanced topics like Velocity Pilot are covered in the ITZ.

However, if the browser window is sized smaller, the table of contents can be accessed by clicking the main menu icon (≡):

The screenshot shows the same browser window after the main menu icon (≡) was clicked. The sidebar is now fully expanded, displaying the entire Table of Contents. The main content area remains the "Welcome" page, identical to the previous screenshot.

Click the main menu icon (≡) to expand the table of contents.

Continue to the [Reserve the IBM Technology Zone environments](#) section to begin the journey to obtain the IBM watsonx Assistant for Z Technical Sales Advanced badge.

IBM Technology Zone environment

To enable sellers to both learn how to deliver client pilots of IBM Watsonx Assistant for Z, multiple environments are available in IBM Technology Zone (ITZ). The environments that are used for the Watsonx Assistant for Z Velocity lab environment can be found in the [IBM Watsonx Assistant for Z](#) collection and consist of:

- **Watsonx Assistant for Z lab – Watsonx Orchestrate:** a dedicated environment with Watsonx Orchestrate on IBM Cloud that is used for creating and configuring the assistant, configuring conversational search, importing skills, and configuring actions.
- **Ansible Automation Platform (AAP) & z/OS:** a pre-configured instance of AAP and Wazi z/OS. The environment includes preinstalled Ansible playbooks that can be imported as skills within Watsonx Orchestrate and connected to your assistant. The environment provides preinstalled templates for various use cases that are covered in a later section. Learn more about AAP [here](#). Learn more about Wazi, [here](#).
- **Single Node OpenShift with NFS storage:** used to provision a single-node Red Hat OpenShift cluster (SNO) on IBM Cloud. The cluster is used to install a dedicated instance of [OpenSearch](#) for Watsonx Assistant for Z. This environment enables ingestion of client-supplied documents.



All activities in this lab guide are required.

While all 3 environments may not be required for every client pilot, to complete the Level 4 learning plan and earn the IBM Watsonx Assistant for Z Technical Sales Advanced badge, you must provision all three ITZ environments and complete all sections in the lab guide. **Any statements in the ITZ collection regarding optional environments or tasks should be ignored.**

Follow the instructions in to create new reservation requests, extend the reservations, and access the ITZ demonstration environments. Provisioning the SNO environment in ITZ can take several hours, while the other two environments typically provisioning in under 30 minutes.

Create a reservation request

1. Click each of the links that follow to open a browser to the reservation pages of the **IBM Watsonx Assistant for Z** ITZ environments.

You may be asked to authenticate to IBM Technology Zone

The steps to authenticate to ITZ are not detailed here as they may vary between users.

[Watsonx Assistant for Z lab – Watsonx Orchestrate - reservation page](#)

[Ansible Automation Platform \(AAP\) & z/OS - reservation page](#)

Single Node OpenShift with NFS storage - reservation page

 The next two steps are for 1 of the 3 environments. Repeat for all three environments.

Be sure to follow these steps to create a reservation in ITZ for all three environments!

2. Click **Reserve now**.

The **Reserve now** option creates a reservation for immediate use. Optionally, schedule the reservation for a later date, for example, when you are at your client's office to start a pilot.



The screenshot shows the 'Create a reservation' interface. At the top, there are tabs for 'Select a environment/infrastructure', 'Select a reservation type', 'Fill out your reservation', and 'Complete'. Below these tabs, it says 'Select your reservation type. Do you need this now or later? Single environment reservation options:' with two radio buttons: 'Reserve now' (selected) and 'Schedule for later'. The 'Reserve now' button is highlighted with a green border.

3. Complete the reservation request and click **Submit**.

The first two reservations are similar to the first image and have fields **a-e** that need to be completed.

- a.** Name: specify a name for the reservation (optional).
- b.** Purpose: select the **Education** purpose tile.



For actual client pilots...

For client pilots, set the **Purpose** field in the reservation to **Pilot** and provide an opportunity number to receive a longer reservation.

- c.** Purpose description: enter a description, for example: Level 4 education.
- d.** Preferred geography: select the region nearest to your physical location for improved performance and reduced network latency.
- e.** End date and time: the initial maximum will be set to 2 days after the current date and time. Instructions follow to extend the reservation end date.
- f.** Accept the IBM Technology Zone's terms and conditions and security policies.
- g.** When satisfied with the parameters, click **Submit**.

a. Name: watsonx Assistant for Z Pilot - watsonx Orchestrate

b. Purpose: Education (Getting experience with specific technology, product, or solutions)

c. Sales Opportunity number: Enter an opportunity number (e.g., IBM Sales Cloud opportunity ID, Project Relationship ID, or a Project Work ID) to allow you to extend your reservation date.

d. Preferred Geography: ITZ-watsonx-2 - AMERICAS - us-south region - dat10 datacenter

e. End Date & Time: Select a date (11/01/2024) and time (11:51 AM, America/Chicago).

f. Notes: Enter any notes you would like to attach to this reservation.

g. I agree to IBM Technology Zone's Terms & Conditions and End User Security Policies Submit

In addition to the preceding fields, the reservation for the **Single Node OpenShift with NFS storage** has these additional fields:

- h.** OCP/Kubernetes cluster network: leave the default setting of **10.128.0.0/14**.
- i.** Enable FIPS security: leave the default setting of **No**. Learn more about the Federal Information Processing Standards (FIPS) [here](#).
- j.** Master single node flavor: select **16 vCPU x 64 GB - 300 GB ephemeral storage**.
- k.** OpenShift version: select **4.14**.
- l.** OCP/Kubernetes service network: leave the default setting of **172.30.0.0/16**.

h. OCP/Kubernetes Cluster Network: 10.128.0.0/14

i. Enable FIPS Security: No

j. Master Single Node Flavor: 16 vCPU x 64 GB - 300 GB ephemeral storage

k. OpenShift Version: 4.14

l. OCP/Kubernetes Service Network: 172.30.0.0/16

m. Notes: Enter any notes you would like to attach to this reservation.

n. I agree to IBM Technology Zone's Terms & Conditions and End User Security Policies Submit

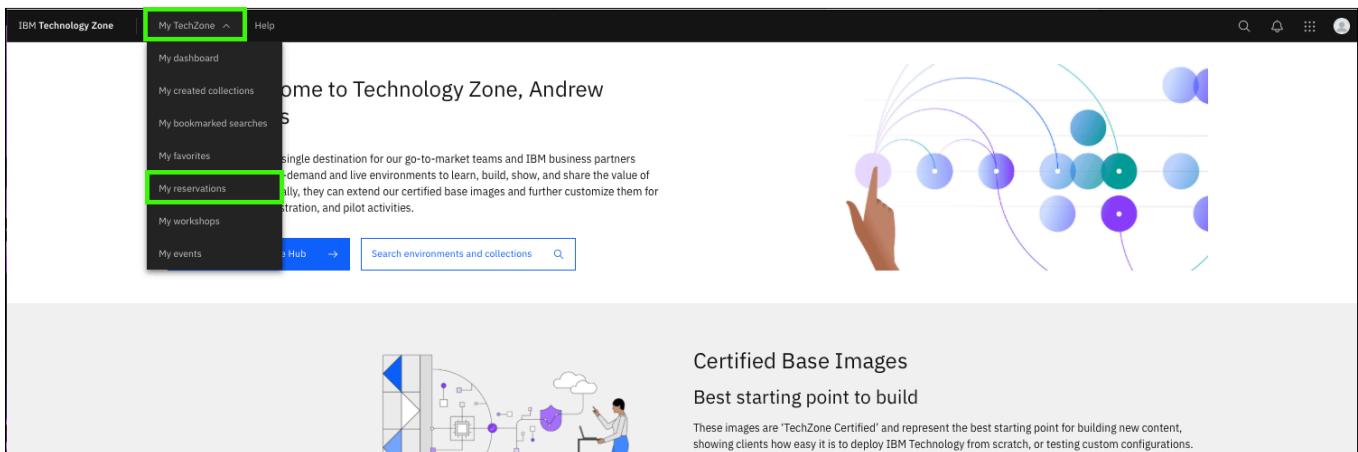
During the provisioning process, multiple emails are sent to you from ITZ as the provisioning process runs. One email states the reservation is provisioning and the other email states that the environment is **Ready**.

In rare cases, the provisioning process may fail. If you receive an email stating the reservation failed, try again by repeating Steps 1-3 for the environment that failed to provision. In addition, review the [Troubleshooting](#) section that follows. If issues continue, open an ITZ support ticket by using the methods that are mentioned in the [Support](#) section.

Extend the reservation

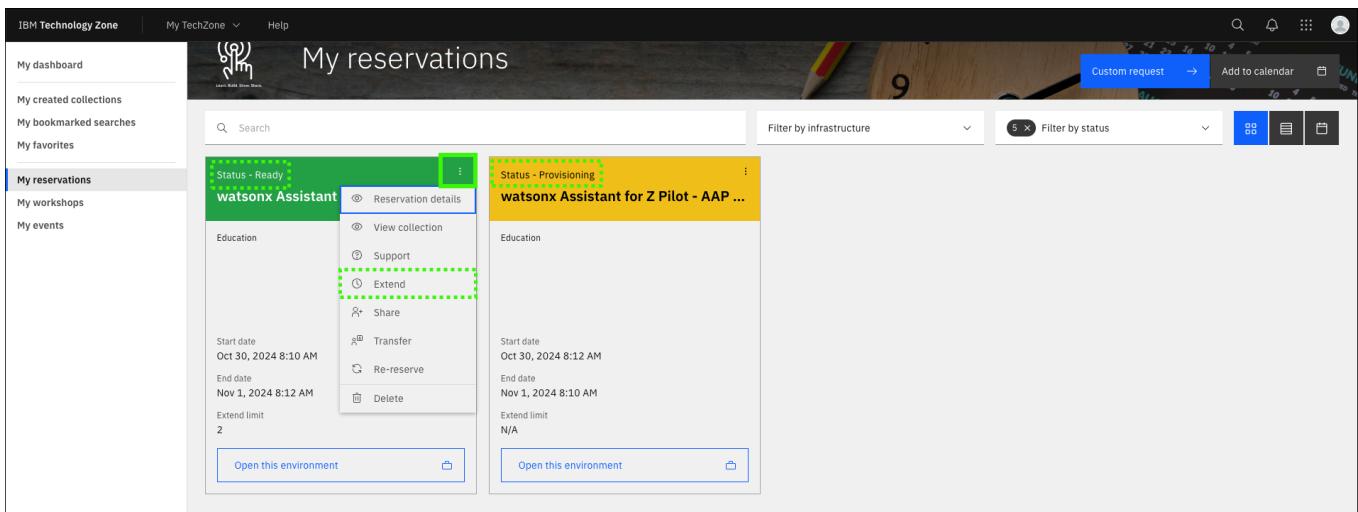
When the reservations are in the **Ready** state, you can extend each reservation to a total of 6 days.

4. In the IBM Technology Zone portal, expand **My TechZone** and select **My Reservations**.



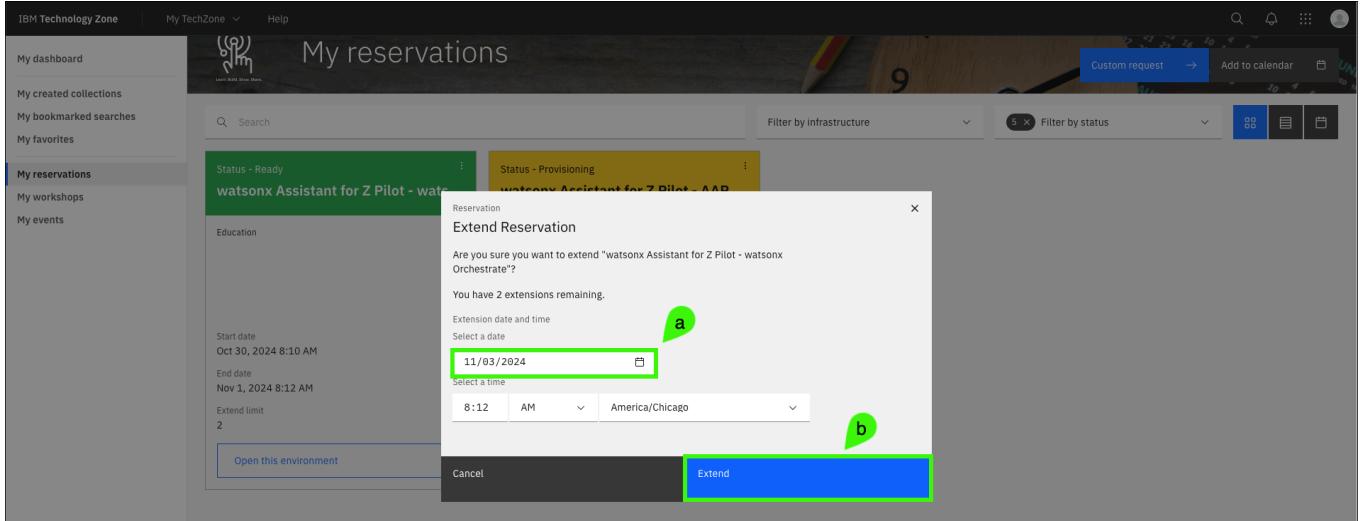
The screenshot shows the IBM Technology Zone homepage. At the top left, there's a navigation bar with 'IBM Technology Zone', 'My TechZone ▾', and 'Help'. Below this is a sidebar with links: 'My dashboard', 'My created collections', 'My bookmarked searches', 'My favorites', 'My reservations' (which is highlighted with a green box), 'My workshops', and 'My events'. To the right of the sidebar, the main content area displays a welcome message 'Welcome to Technology Zone, Andrew' and a graphic of a hand pointing at a network of nodes. Below this, there's a section titled 'Certified Base Images' with a sub-section 'Best starting point to build' and a note about TechZone Certified images.

5. Click the overflow icon (≡) on the reservation tile and select **Extend**.



The screenshot shows the 'My reservations' page. On the left, there's a sidebar with 'My dashboard', 'My created collections', 'My bookmarked searches', 'My favorites', 'My reservations' (which is selected and highlighted with a blue box), 'My workshops', and 'My events'. The main content area shows two reservation tiles. The first tile, for 'watsonx Assistant', has a status of 'Status - Ready'. A context menu is open over this tile, with the 'Extend' option highlighted (also with a green box). Other options in the menu include 'Reservation details', 'View collection', 'Support', 'Transfer', 'Re-reserve', and 'Delete'. The second tile, for 'watsonx Assistant for Z Pilot - AAP ...', has a status of 'Status - Provisioning'. It also has a context menu with similar options. At the bottom of the page, there are buttons for 'Custom request', 'Add to calendar', and navigation icons.

6. Click the **Select a date** option, (a) specify the date to extend to, and then (b) click **Extend**.

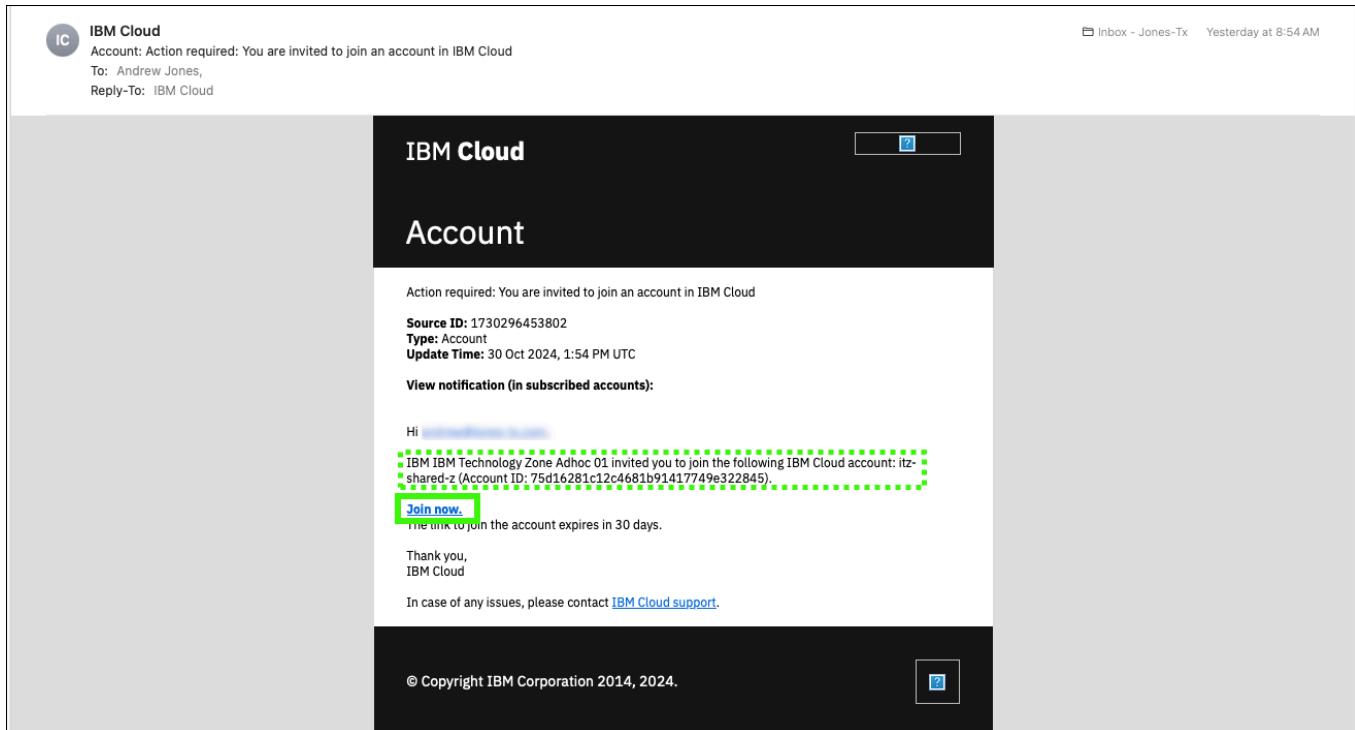


If you anticipate needing more time, repeat Steps 5 and 6 to extend the reservation to the maximum of 6 days. Repeat these steps for the other two reservations.

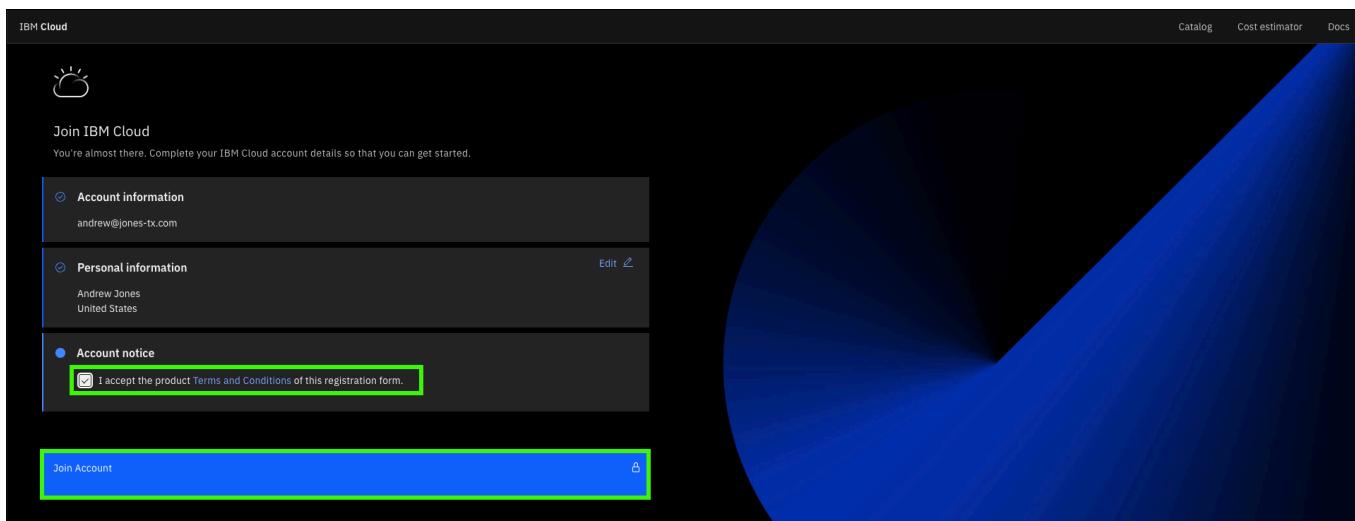
Join the ITZ IBM Cloud accounts

Both the **watsonx Assistant for Z lab – watsonx Orchestrate** and the **Ansible Automation Platform (AAP) & z/OS** environments add you to an IBM Cloud account while your reservation is active. During the provisioning process of these ITZ environments, you receive two emails from IBM Cloud. To access the environment, you must first accept the invitations to join both of the IBM Cloud accounts.

7. Open the emails from **IBM Cloud** and click the **Join now** links.



8. In the **Join IBM Cloud** browser windows that open, select the **I accept the product Terms and Conditions of this registration form**, and then click **Join Account****.



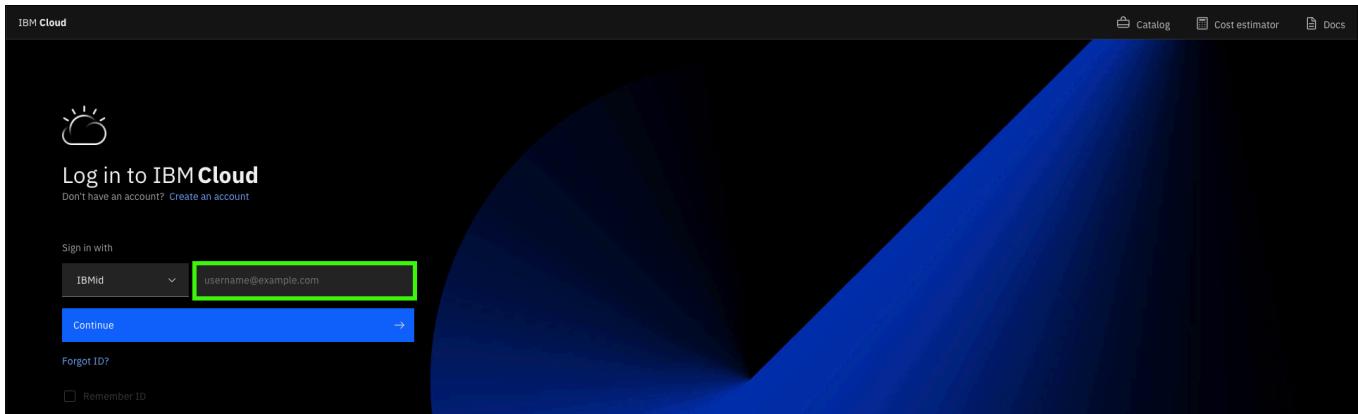
Repeat steps 7 and 8 for the second invitation.

After joining both accounts, verify that both accounts appear in your available account list in the IBM Cloud portal.

- Click the following link to open a browser to the IBM Cloud portal.

IBM Cloud portal

- Follow the directions to complete the authentication to IBM Cloud using the same email address you used to login to ITZ. The login steps very depending on any two-factor authentication methods enabled.



- Click the **account** menu and verify access to the two IBM Cloud accounts listed in your ITZ reservations.



These accounts may change within ITZ.

Over time, the accounts may change for the environments. The accounts names should align with the accounts named in the invitation emails you received.

 Does your IBM Cloud portal view look different?

If your IBM Cloud portal looks different from the images above, it could be because the IBM Cloud portal has done through a design change, or your browser window is set to smaller size. Instead of the current selected account appearing in the top menu, you may see this **change account** icon:  Click this icon to view the list of accounts you can access.



Accessing the environments

Each reservation provides access to its respective environment. Details for accessing each environment are provided in the **Pilot setup** sections that follow in the lab guide.

After all three reservations are in the **Ready** state and you accepted the invitations to the IBM Cloud accounts, proceed to the next section to complete the pilot setup.

Troubleshooting

 If your reservation for the Single Node OpenShift environment fails...

If your reservation for the Single Node OpenShift environment fails, try selecting one of the **eu-gb region** options as the **Preferred Geography**.

Pilot setup

Creating an assistant and configuring conversational search

To create your watsonx Assistant for Z, you use [watsonx Orchestrate](#) to create the assistant and configure conversational search. You can configure your assistant to use conversational search by using a hosted [OpenSearch](#) instance. The pre-configured instance in IBM Technology Zone (ITZ) has over 220 knowledge sources and supports Retrieval Augmented Generation (RAG). The large language model (LLM) providing the conversational AI augments this knowledge based on IBM Z documentation. All of these elements create IBM Z context-aware responses to queries with its content-grounded knowledge.

Access the ITZ IBM Cloud account for the watsonx Assistant for Z Pilot environment

1. In the IBM Technology Zone portal, expand **My TechZone** and select **My Reservations**, or click the following link.

ITZ My reservations

2. Click the **watsonx Assistant for Z Pilot - watsonx Orchestrate** tile.

3. Record the ITZ IBM Cloud account name associated with the reservation.

Purpose

- Purpose
- Education
- Opportunity Product(s)
- Customer(s)

Environment

Reservation ID	Type
672a091a8f85062f891e081	IBM Cloud
Request method	Transaction ID
watsonx-orchestrate	115897c9-58a1-4f17-af9c-b16dc2a97590
ITZ-WATSONX-036	Geo
Region	americas
us-south	Datacenter
Customer data	dal10
false	Environment
Idle runtime limit	watsonx-orchestrate-wusdf
10800	Timeout action

Reservation Details

IBM Cloud Login
<https://cloud.ibm.com/resources>

4. Click the IBM Cloud Login link.

Purpose

- Purpose
- Education
- Opportunity Product(s)
- Customer(s)

Environment

Reservation ID	Type
672a091a8f85062f891e081	IBM Cloud
Request method	Transaction ID
watsonx-orchestrate	115897c9-58a1-4f17-af9c-b16dc2a97590
Cloud Account	Geo
ITZ-WATSONX-036	americas
Region	Datacenter
us-south	dal10
Customer data	Environment
false	watsonx-orchestrate-wusdf
Idle runtime limit	Timeout action
10800	

Reservation Details

IBM Cloud Login
<https://cloud.ibm.com/resources>



Steps to authenticate to IBM Cloud are not illustrated here.

You may need to authenticate to IBM Cloud after clicking the link. These steps are not shown here as they may vary by individual.

5. Verify that the current IBM Cloud account is the same as the account name recorded in step 3, if not the same, switch to the proper account.

Note: the formatting of the name can appear differently than what is shown in the ITZ reservation.

The screenshot shows the IBM Cloud Resource list interface. At the top right, there is a dropdown menu labeled "2953593 - ltz-watsonx-036". This dropdown is highlighted with a green border. Below the dropdown, there is a "Create resource" button.

If the proper account is not listed, click the account drop down and select the proper account. Note, if your browser window is narrow, the account drop down can be depicted with the switch account icon ().

This screenshot is similar to the previous one, but the account dropdown menu is shown in a darker, more compact form. A small blue icon with three dots and a switch account symbol is visible next to the dropdown button, indicating it can be used to change accounts.

Create your Assistant

6. Click the **resources** icon ().

The screenshot shows the IBM Cloud Dashboard. On the left side, there is a sidebar with several icons. The "resources" icon, which looks like a grid with three horizontal lines, is highlighted with a green border. Other icons include "Dashboard", "Watson Assistant", "Watson Studio", and "Watson Assistant".

7. Expand the **AI / Machine Learning** section and click the **watsonx Orchestrate** instance listed (the instance name is different than shown in the following image).

Resource list

Name	Group	Location	Product	Status	Tags
itwxo-2700039nft-erspw	watsonx-orchestrate-erspw	Dallas	watsonx Orchestrate	Active	-

8. Click Launch watsonx Orchestrate.

itwxo-2700039nft-erspw Active Add tags

Manage Service credentials

Start by launching the tool

Launch watsonx Orchestrate Getting started tutorial

Credentials

API key: [Download](#) [Show credentials](#)

URL: <https://api.us-south.assistant-builder.watson.cloud.ibm.com/instances/a7675d8e-e885-4891>

Plan Essentials Plan

9. Click the **AI assistant builder** tile to start creating a new assistant.

Welcome, Andrew Jones!

Take productivity to the next level.

Try Skills in Chat →

AI assistant builder Boost productivity and customer care by creating conversational experiences.

Skill studio Build the skills your team needs to get their work done more quickly.

Upgrade Standard plan | Skill studio

Automate how your business works

Learn more →

10. Enter a name and optional description for your assistant and click **Next**.

Welcome to AI assistant builder

Create your first assistant

Let's get your assistant up and running. Name your assistant, add a description, and choose a language. In following steps we'll gather more information, show you basic customizations, and give you a preview of what your assistant will look like.

Assistant name
Zeeves

Your assistant name will be kept internally and not visible to your customers

Description (optional)
Zeeves, the Jeeves for everything IBM Z.

40/128

Assistant language
English (US)

This is the language your assistant will speak.

Next

11. Complete the **Personalize your assistant** form and click **Next**.

Explore the personalization options. In creating an assistant for a client pilot, consider specifying attributes that align with the client's business.

- Select **Web**.
- Select the industry of your choice.
- Select the role of your choice.
- Select the need of your choice.

Welcome to AI assistant builder

Personalize your assistant

Tell us where your assistant will live
You may add multiple channels from your dashboard.

Where do you plan on deploying your assistant?
a Web

Tell us about yourself
This information will be used to personalize your onboarding experience.

Which industry do you work in?
b Software

What is your role on the team building the assistant?
c Developer

Which statement describes your needs best?
d I want to automate common tasks in a natural way

This is what your customers will experience

watsonx Assistant

Do you have the Speed Demons in stock?

The Speed Demons are in stock at our Downtown and Northgate locations, which are both within 5 miles of you.

What size and color do you need?

I'm looking for a size 9 in white

Great news! The Speed Demons are available in white in a size 9.

You can purchase them for curbside pickup or we can ship them to you. Which would you prefer?

I'll pick them up! Ship them to me!

Type something...

Next

12. Complete the **Customize your chat UI** form and click **Next**.

Explore the customization options. When creating an assistant for a client pilot, consider specifying attributes that align with the client (for example, colors and logos).

13. Preview your assistant and then click **Create**.

The assistant is now created.

Configure conversational search

The next step will be to configure **conversational search** for your assistant that uses a hosted instance of OpenSearch.

14. Click **Generative AI** menu item (💡) in the left navigation.

15. Review the base large language model (LLM) settings.

Notice the other LLM models available. For most pilots, the **granite-3-8b-instruct** model is appropriate.

16. Click Set up your Search Integration.

By default, conversational search is not enabled when an assistant is created. Conversational search takes priority over general-purpose answering if both are enabled. Learn more about conversational search in watsonx [here](#).

17. Click Custom service.

18. Complete the **Custom service form and then click **Next**.**

- Select **By providing credentials**.
- Enter the following value in the **URL** field (use the copy icon to avoid typographical errors).

https://wxa4z-opensearch-wrapper-wxa4z-demo.wxo4z-opc-opensearch-clus-47e063e6a3ad1f71bf2e58f91c3b4c2e-0000.us-south.containers.appdomain.cloud/v1/query

- Select **None** in the **Choose an authentication type** drop-down list.

19. Enable **conversational search and then click **Save**.**

20. Update the conversational search **custom service settings based on your requirements.**

Note: the **Settings** page is divided into two sections in the following images to enhance the visibility of the screen captures. Learn more about these settings [here](#).

The following settings are proven to work well. You can experiment with these settings to see how they affect queries for your client's pilot.

a. Enable Conversational search.

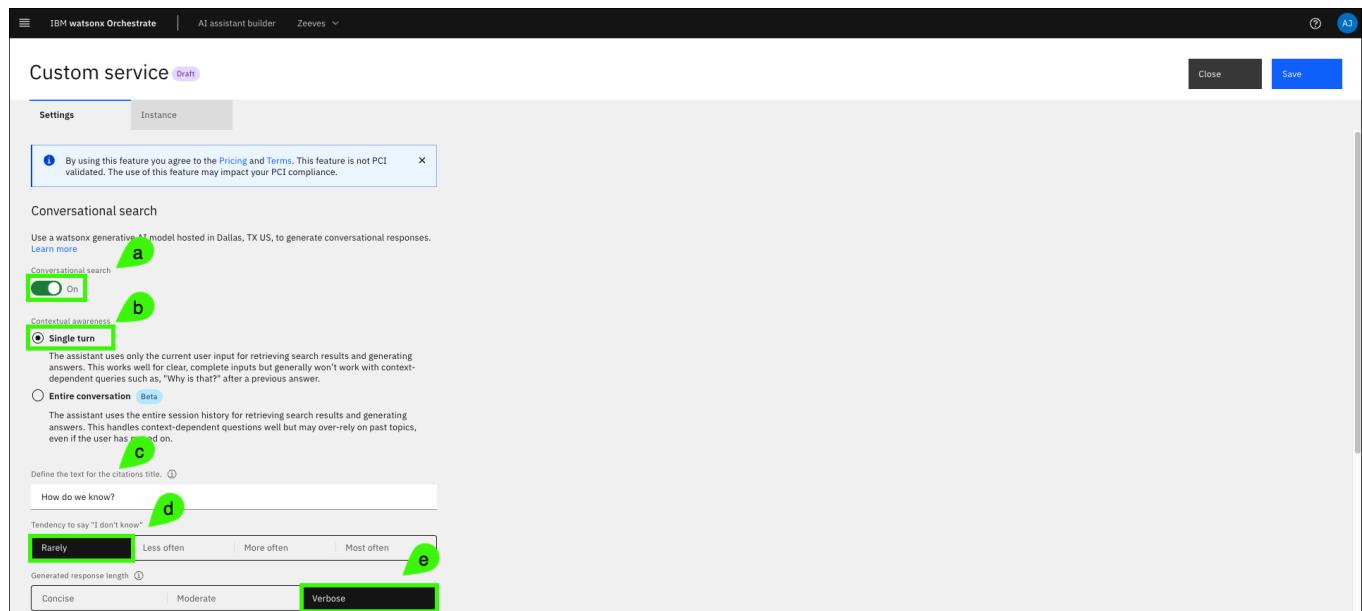
b. Select Single turn. Enabling multi-turn conversation (by selecting Entire conversation) is not yet supported for the solution on-premises. Be mindful in using this option and help ensure that the client understands what is supported in the solution.

c. Specify the text appears to expand the list of citations in the assistant (except web chat client).

d. Select Rarely for the tendency to say "I don't know" setting.

e. Select Verbose for the generated response length. This setting affects the average response length.

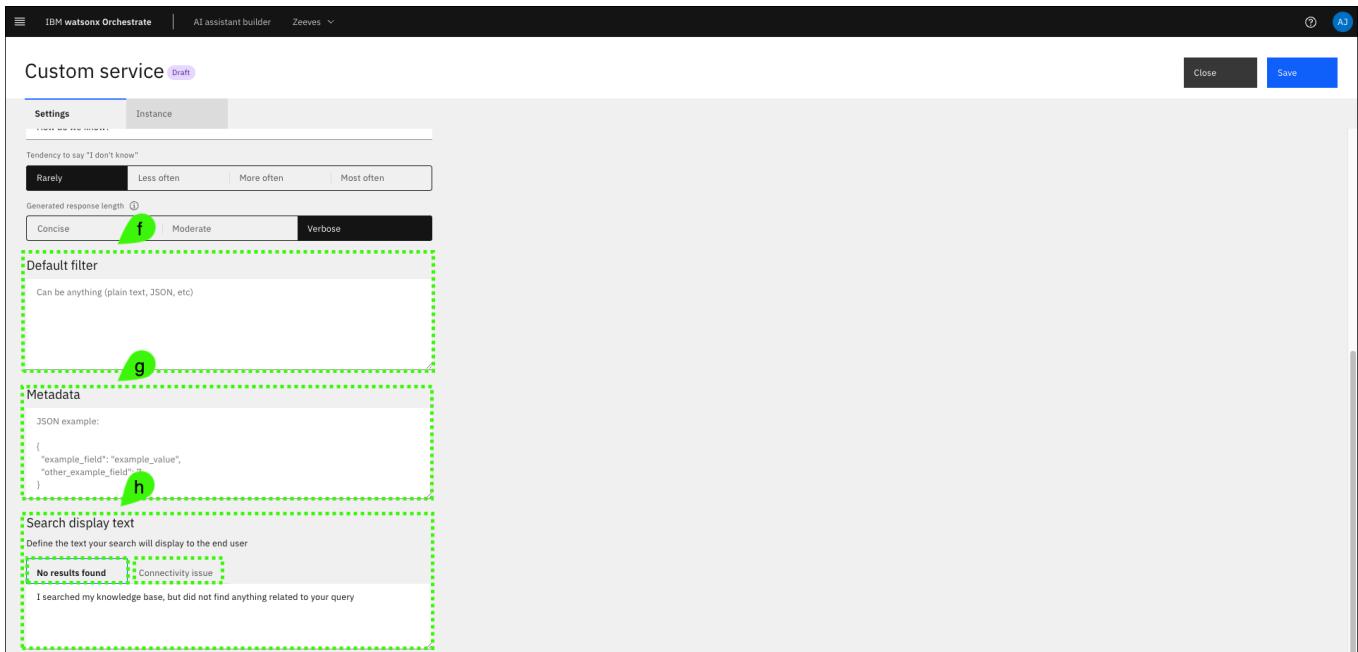
Depending on user input, variations from the selected length can occur.



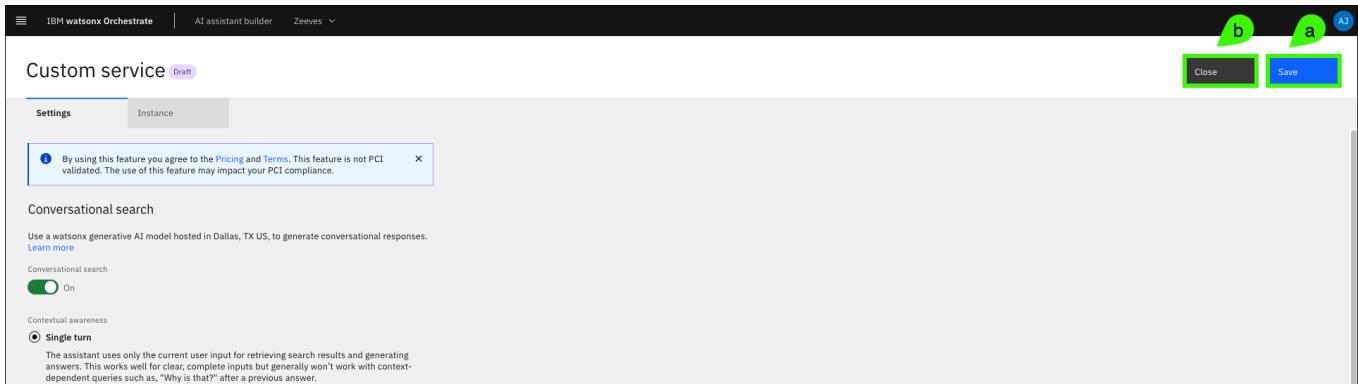
f. Leave the Default filter field empty.

g. The Metadata field provides a way to adjust your assistant's behavior during conversational search for your OpenSearch instance. This option is explored in detail in the [Installing and using zassist to ingest client documents](#). Leave the field empty for now.

h. The Search display text options specify the default text displayed when no results are found or when connectivity issues to the backend search service occur. You can keep the defaults or customize.



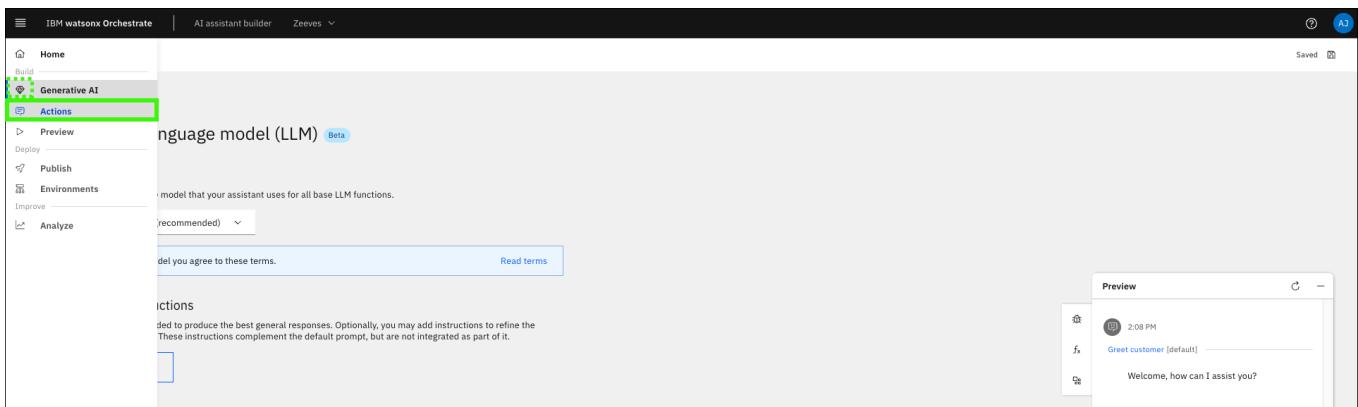
21. Click (a) **Save** and then click (b) **Close**.



More configurations

After you save and close the **Conversational search** configuration page, a few more configurations are needed to get the best experience from your conversational chat. Details on these settings are available [here](#).

22. Hover over the **Generative AI** icon (💡) in the left navigation and click **Actions**.



23. Click **Set by assistant** under the **All items** menu.

24. Click no matches.

Name	Last edited	Examples Count	Status
Greet customer	3 hours ago	1	Green
Trigger word detected	3 hours ago	1	Green
No matches	3 hours ago	1	Green
Fallback	3 hours ago	1	Green

25. Click Step 1 under Conversation steps.

26. Select (a) without conditions in the Is taken drop-down menu and then click (b) Clear conditions.

Note, the **Is taken** value does not change from **with conditions** after selecting **without conditions**.

27. Delete the default text in the Assistant says entry field.

The screenshot shows the AI assistant builder interface. On the left, there's a sidebar with 'Conversation steps'. Step 1 is highlighted with a green dashed border and contains the text 'This step has no content' and 'Action complete'. Step 2 is below it with the text 'This step has no content' and 'Go to subaction: Fallback'. In the main panel, under 'Step 1', 'Is taken' is set to 'without conditions'. Under 'Assistant says', there's a text area with placeholder text 'For example: Please select from the following options:' and a 'Define customer response' button. Below that, 'And then' has a dropdown menu with 'End the action' selected.

28. Select **Search for the answer** in the **And then** drop-down menu.

This screenshot is similar to the previous one, but the 'And then' dropdown menu is open. The 'Search for the answer' option is highlighted with a green box. Other options in the menu include 'Continue to next step', 'Re-ask previous step(s)', 'Go to a subaction', 'Use an extension', 'Connect to agent', and 'End the action'.

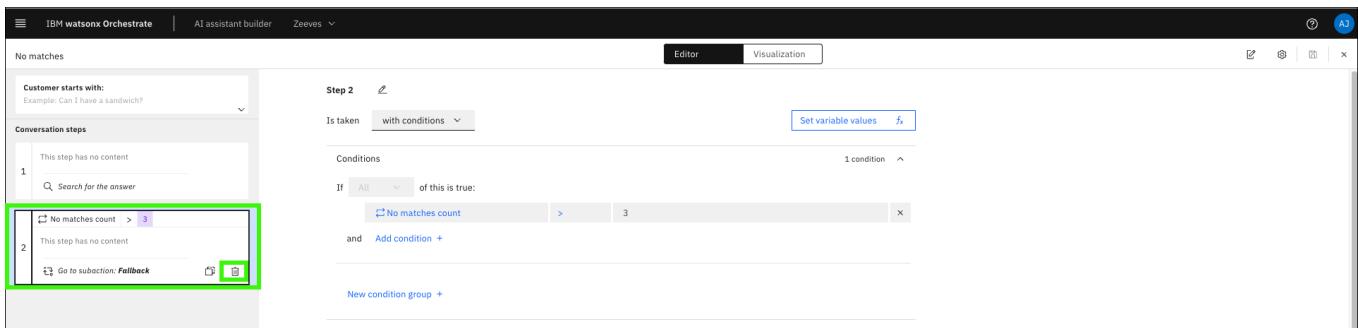
29. Click **Edit settings**.

In this screenshot, the 'And then' dropdown is open again, but now the 'Edit settings' button at the bottom of the menu is highlighted with a green box. The other options remain the same as in the previous screenshot.

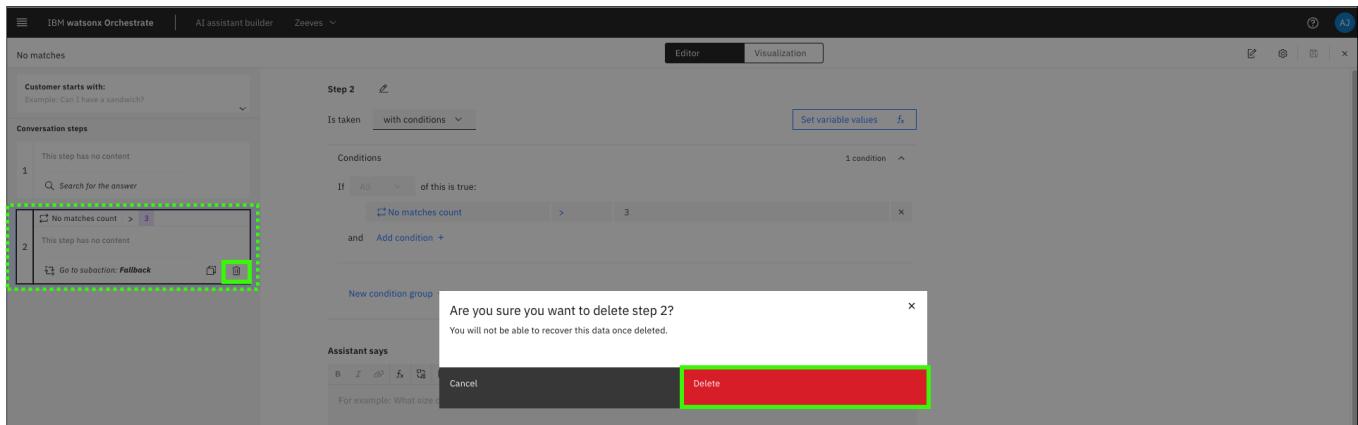
30. Select **End the actions after returning results** and then click **Apply**.



31. Select **Step 2 (No matches count)** under **Conversation steps** and click the delete icon (>Delete).



32. Click **Delete** in the confirmation dialog to delete step 2.



33. Click the **x** to close the **Editor** window.



34. Click Fallback in the Actions table.

Name	Last edited	Examples Count	Status
Greet customer	3 hours ago	1	Green
Trigger word detected	3 hours ago	1	Green
No matches	a minute ago	1	Green
Fallback	3 hours ago	1	Green

35. Delete all of the Conversation steps.

Note: the following image is edited. Only 5 steps are shown, but all 6 need to be deleted. You need to select each step individually, click the delete icon (trash), and confirm the deletion.

36. Verify that all Conversation steps are deleted and then click the x to close the Editor window.

37. Click the global settings (⚙️).

Name	Last edited	Examples Count	Status
Greet customer	4 days ago	1	Green
Trigger word detected	4 days ago	1	Green
No matches	4 days ago	1	Green
Fallback	4 days ago	1	Green

38. Click **No matches** under the **Conversation routing** tab.

The screenshot shows the 'Global settings' page with the 'Conversation routing' tab selected. In the 'No matches' section, there is a note about handling user input when no relevant options are found. A toggle switch labeled 'Enable disambiguation' is set to 'On'. Below it, a text field contains the placeholder 'Did you mean:'. Another text field below it contains 'None of the above'. To the right, there is an example of a 'Bank Bot' interaction where a user asks 'open an account' and the bot responds with 'Did you mean:' followed by a list of suggestions: 'Open a new savings account', 'Open a new checking account', 'I want to apply for a mortgage loan', and 'None of the above'.

39. Move the slider to **More often** (or select **More often** in the drop-down).

The setting helps ensure that actions are triggered less often unless the user's query specifically matches the action's input.

The screenshot shows the 'Global settings' page with the 'Conversation routing' tab selected. The 'No matches' section includes a note about triggering the 'No matches' action for unrecognized input. A slider is shown with a green box around its track, and the dropdown menu next to it is also highlighted with a green box, showing the option 'More often'.

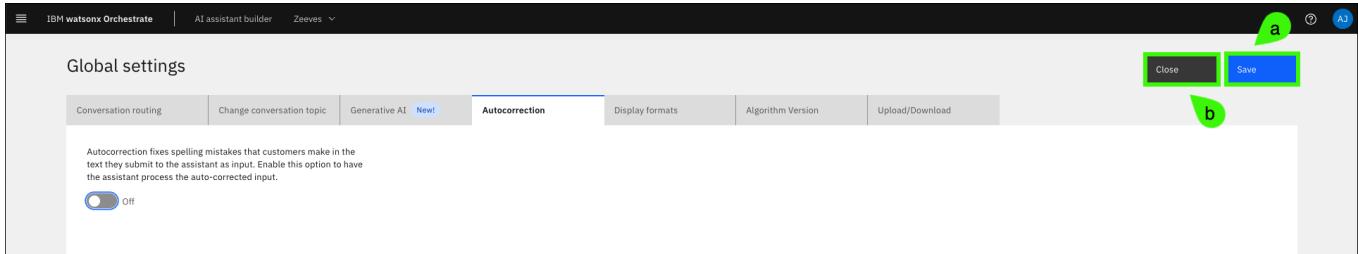
40. Click **Autocorrection**.

The screenshot shows the 'Global settings' page with the 'Autocorrection' tab selected. There is a note about fixing spelling mistakes in customer input. A toggle switch labeled 'Off' is highlighted with a green box.

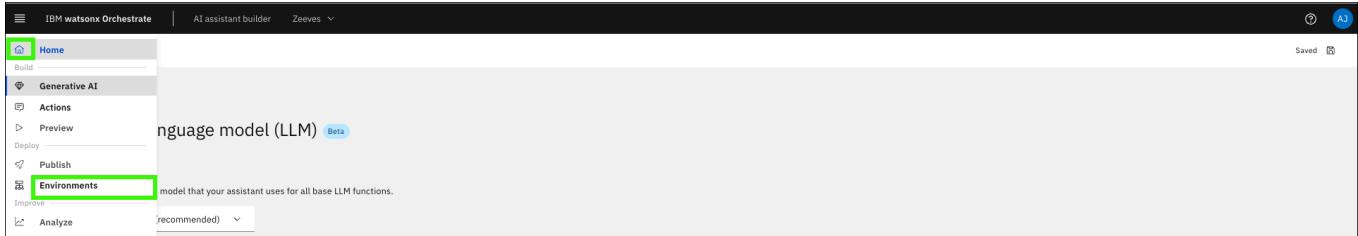
41. Click the **autocorrection** toggle to turn the feature **off**.

The screenshot shows the 'Global settings' page with the 'Autocorrection' tab selected. A note about autocorrection is present. A toggle switch labeled 'Off' is highlighted with a green box.

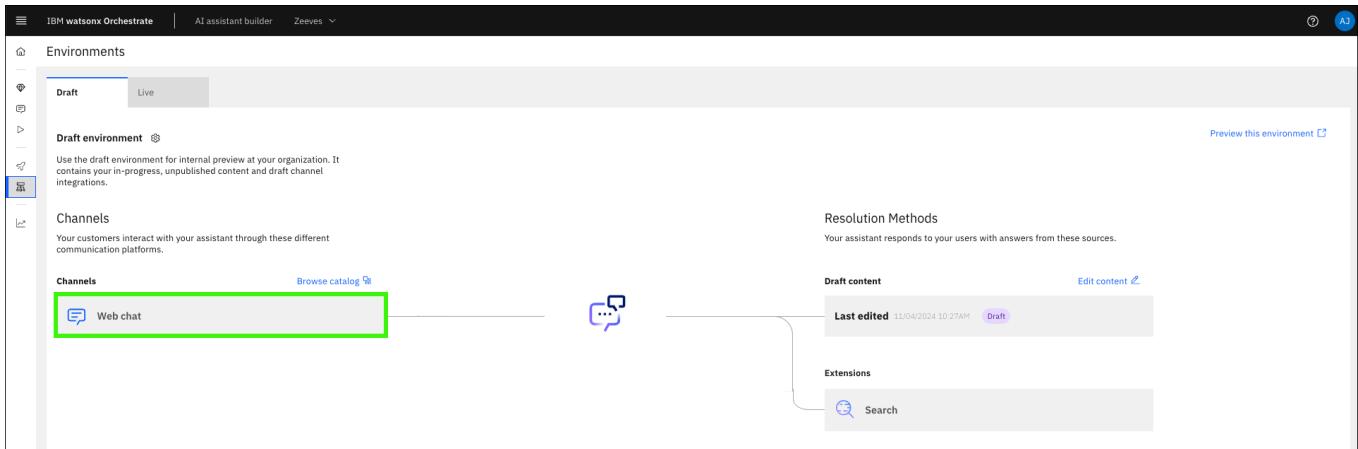
42. Click (a) **Save** and then (b) **Close**.



43. Hover over the **home** (🏠) and click **Environments**.



44. Click **Web chat**.



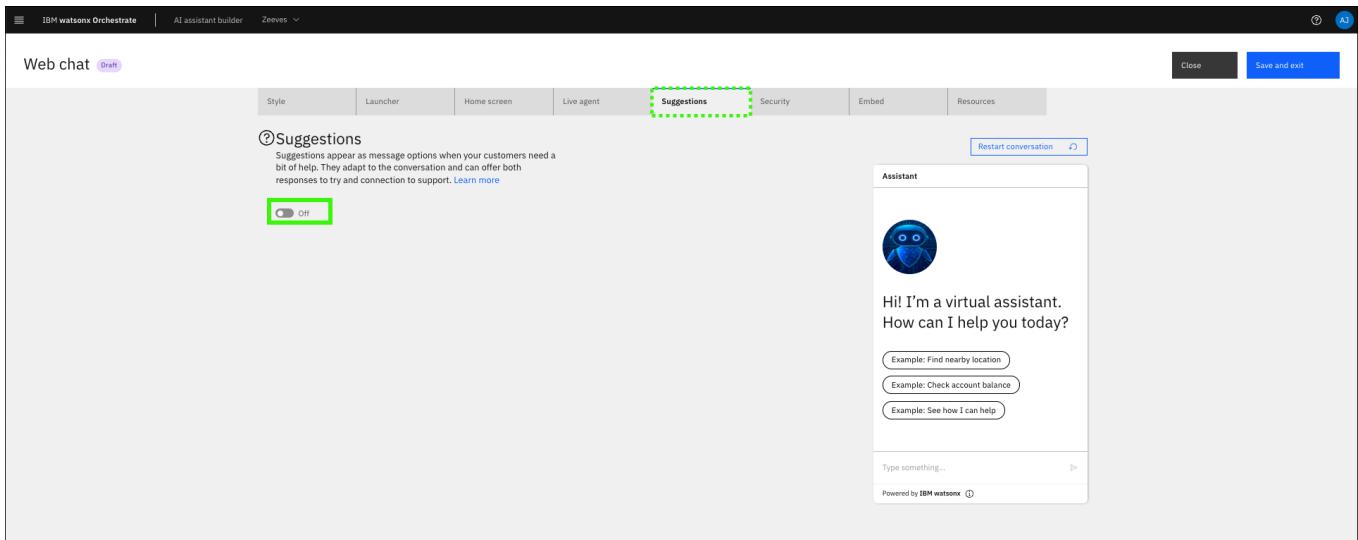
45. On the **Style** tab, click the **Streaming** toggle to enable streaming.

The screenshot shows the 'Style' tab selected in the AI assistant builder. On the left, there's a configuration panel for 'Customize your chat UI'. It includes fields for 'Assistant's name known by customers' (set to 'Assistant'), 'Intended purpose' (set to 'Standard'), 'Choose a theme' (set to 'Light'), color swatches for 'Primary color' (#FFFFFF) and 'Secondary color' (#3D3D3D), and sections for 'Size' (width 380px, height 640px), 'IBM Watermark' (disabled), and 'Streaming' (disabled). On the right, a preview window shows a blue owl-like assistant with the message 'Hi! I'm a virtual assistant. How can I help you today?'. Below the preview is a text input field with placeholder 'Type something...' and a 'Powered by IBM watsonx' link.

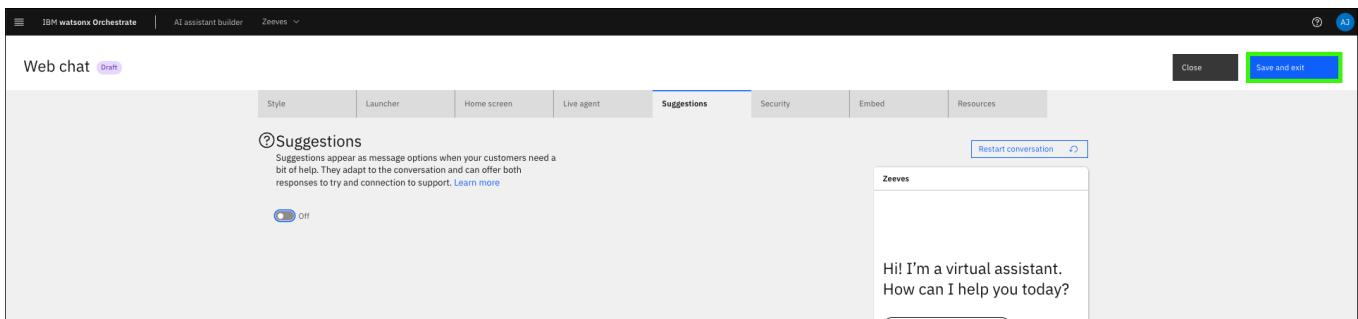
46. Click Suggestions.

The screenshot shows the 'Suggestions' tab selected in the AI assistant builder. The configuration panel on the left remains the same as in the previous screenshot. However, the 'Suggestions' toggle in the 'Streaming' section is now turned on, indicated by a green border around the button. The preview window on the right shows the same blue owl-like assistant with the same message and interface.

47. Click the Suggestions toggle to turn this feature off.



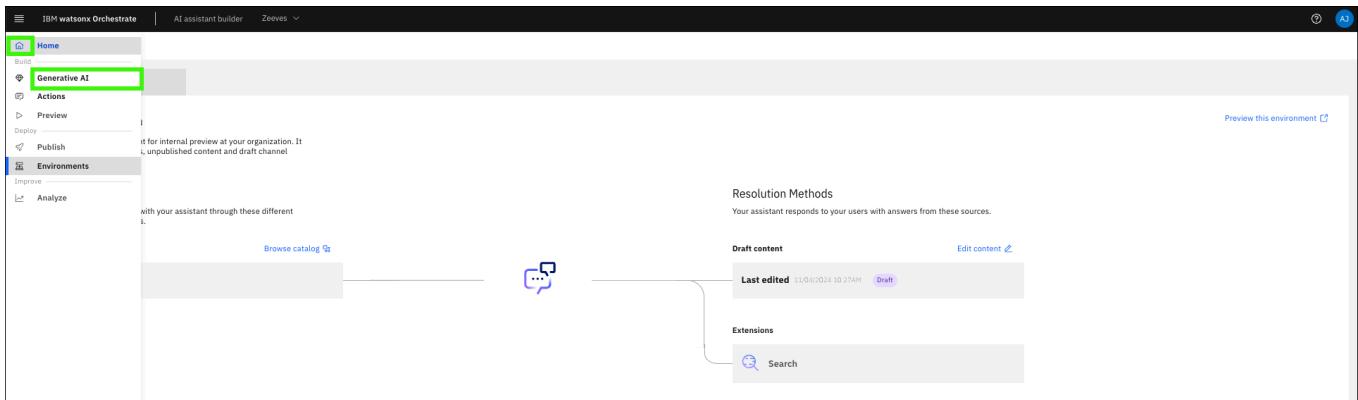
48. Click Save and exit.



Configure the base large language model

After the preceding steps are completed, there are enhancements that can be made to configure how the large language model (LLM) responds to your queries. Including adding prompt instructions and configuring the LLM's answer behavior. These options can be summarized [here](#).

1. Hover over the home () and click Generative AI.



2. Click Add instructions.

3. Enter a prompt instruction.

This option instructs the LLM in your assistant to give refined responses by adding prompt instructions. The instructions help the LLM guide the conversations with clarity and specifically to achieve the end-goal of an action.

Enter the prompt instructions in the field. The maximum number of characters you can enter in the prompt instruction field is 1,000.

The following is an example prompt instruction that works well. Experiment with different prompt instructions.

You are a subject matter expert on mainframe systems. Please respond to all prompts with truth and accuracy. Keep all answers short and concise, unless requested to provide details.

Note: When the instructions are typed in, they are automatically saved and the LLM is immediately trained on them.

4. Toggle General-purpose answering to off.

The ability exists to configure the answering behavior of your assistant to provide responses that are based on the preinstalled content or general content.

On the **Generative AI** page (under **Prompt Instructions**), you see the **Answer behavior** section. After you configure **Conversational search**, you see that it is enabled (toggled on) with the search integration added.

If you enable both General-purpose answering as well as Conversational search, the Conversational search answering takes precedence over General-purpose answering.

For purposes of retrieving Z-specific answers and responses, it is recommended that you turn off General-purpose answering and leave only Conversational search turned on.

The screenshot shows the 'Generative AI' configuration page. In the 'Answer behavior' section, the 'General-purpose answering' toggle switch is off (indicated by a greyed-out appearance). The 'Conversational search' toggle switch is on (indicated by a green background). A preview window on the right shows a sample conversation:

```

10:47 AM
Greet customer [default]
Welcome, how can I assist you?

```

Testing conversational search

Now, you can begin issuing queries to test the assistant's responses.

It is important to keep in mind that many of the settings configured earlier can be iteratively modified based on your assessment of the quality of responses. The settings can be reviewed and changed at any time. For example, adding extra prompt instructions, changing the verbosity of the responses, and modifying the indexes used for OpenSearch.

53. Hover over the **home** (🏠) and click **Preview**.

The screenshot shows the main navigation bar of the IBM Watson Assistant for Z Level 4 Lab Guide. The 'Preview' button in the top navigation bar is highlighted with a green box.

54. Experiment with different prompts and validate the answers are reasonable and related to IBM Z.

Other prompts and responses follow. Note: the responses that you receive can vary from the ones shown.

Prompt:

What is the APF list in z/OS? Provide a detailed explanation?

Example output:

The screenshot shows the IBM Watson Assistant interface. At the top, there are tabs for "IBM Watsonx Orchestrate", "AI assistant builder", and "Zeeves". Below the tabs, a sidebar on the left has icons for "Preview assistant", "Sample website", and other settings. The main area is titled "Sample website". On the right, there is a "Zeeves" chat window. A message from the user is highlighted with a green border: "What is the APF list in z/OS? Provide a detailed explanation?". The AI response follows, also with a green border: "The APF (Authorized Program Facility) list in z/OS is a system of controls which lists the data sets are authorized to run in privileged mode. It is used to ensure system security and integrity by restricting access to sensitive system resources. The APF list is located in the system's parmlib (parameter library) and can be managed using various system commands and macros. The APF list contains a list of libraries and data sets that are authorized to run in privileged mode. These libraries and data sets are typically system-critical programs and utilities that require elevated privileges to function correctly. By authorizing these". Below the AI response is a text input field with "Type something..." and a "Powered by IBM Watsonx" link.

Prompt:

Why is Db2 different than other database systems?

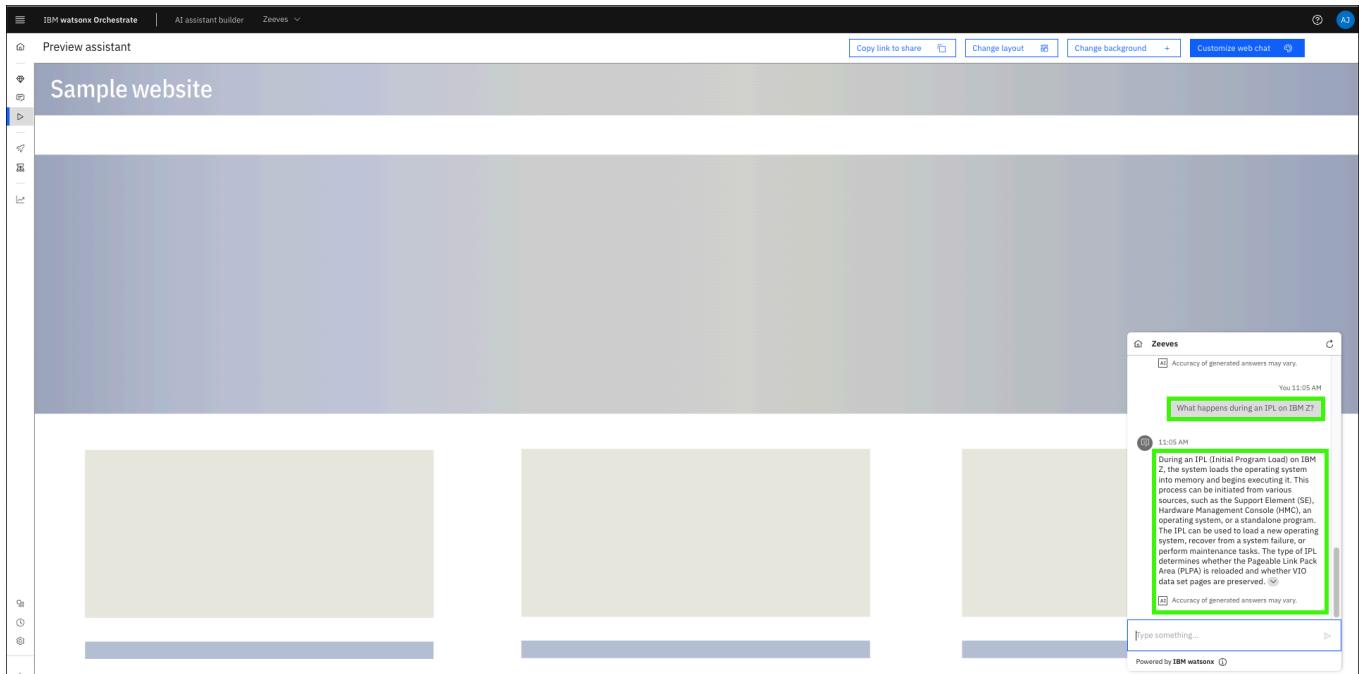
Example output:

The screenshot shows the IBM Watson Assistant interface. At the top, there are tabs for "IBM Watsonx Orchestrate", "AI assistant builder", and "Zeeves". Below the tabs, a sidebar on the left has icons for "Preview assistant", "Sample website", and other settings. The main area is titled "Sample website". On the right, there is a "Zeeves" chat window. A message from the user is highlighted with a green border: "Why is Db2 different than other databases?". The AI response follows, also with a green border: "Db2 is a relational database management system (DBMS) that runs under the z/OS operating system, which is specific to mainframe computers. Other database systems run on other operating systems like Windows, Linux, or Unix. Additionally, Db2 provides a higher level of data management and security features compared to other DBMSs. For example, it manages data stored in databases and controls access to the data. Accuracy of generated answers may vary.". Below the AI response is a text input field with "Type something..." and a "Powered by IBM Watsonx" link.

Prompt:

What happens during an IPL on IBM Z?

Example output:



You now have a working assistant that uses IBM watsonx Assistant for Z. Take time to explore different prompt instructions and settings. If you encounter any issues, the Troubleshooting section that follows can help resolve them.

Continue to the [Creating a stand-alone OpenSearch instance for document ingestion](#) learn how to configure a dedicated OpenSearch instance for ingesting client-specific documentation into the RAG model.

Troubleshooting

The following are issues that you may encounter. If the provided resolutions do not work, contact support by using the methods that are mentioned in the [Support](#) section.

✖ Assistant responds to all prompts with, "I might have information related to your query to share, but am unable to connect to my knowledge base at the moment" ✖

This Assistant is unable to connect to the custom service URL specified. This could be a network issue, the service may be down, the service may be restarting, or the service is no longer running at that URL.

Before reaching out to [Support](#), try the following:

- Wait a few minutes and try again. It may be the service was in the process of restarting.
- If you printed this demonstration guide or saved a copy, verify you are using the most current version of the [lab guide](#) and the correct service URL (<https://wxa4z-opensearch-wrapper-wxa4z-demo.wxo4z-opc-opensearch-clus-47e063e6a3ad1f71bf2e58f91c3b4c2e-0000.us-south.containers.appdomain.cloud/v1/query>). The URL may have changed since you saved or printed the lab guide.

Ingesting client documents

Creating a stand-alone OpenSearch instance for document ingestion

Next, learn to enable a client to ingest their own documentation into the Retrieval Augmented Generation (RAG) used by IBM watsonx Assistant for Z. The first step is to deploy a dedicated [OpenSearch](#) instance. The dedicated search instance will be referred to as bring-your-own-search (BYOS).

Earlier, you provisioned three IBM Technology Zone (ITZ) environments. One of which was a single-node Red Hat OpenShift (SNO) cluster. If you have not reserved this environment, or it is not in the **Ready** state, return to the [IBM Technology Zone environment](#) section to complete the reservation.

Install the Red Hat OpenShift command-line interface utility

The Red Hat OpenShift command-line interface (CLI) utility, which is known as **oc**, must be installed on your local workstation. If you already installed the **oc** utility, you can proceed to [log in to the SNO cluster](#).

1. Click the following link to open a browser window to your ITZ reservations.

[ITZ My reservations](#)

2. Click the **Single Node OpenShift** tile.

The screenshot shows the 'My reservations' section of the IBM Technology Zone interface. On the left, a sidebar lists options like 'My dashboard', 'My created collections', 'My bookmarked searches', 'My favorites', and 'My reservations'. The 'My reservations' option is selected and highlighted with a blue background. The main area displays three environment tiles:

- Single Node OpenShift (VMware on IBM Cloud)**: Status - Ready. This tile is highlighted with a green border around its title and the 'Open this environment' button below it. It shows details: Start date Oct 31, 2024 7:02 AM, End date Nov 6, 2024 6:49 AM, Extend limit 0.
- watsonx Assistant for Z Pilot - AAP & z/OS**: Status - Ready. This tile shows details: Start date Oct 30, 2024 8:12 AM, End date Nov 5, 2024 7:53 AM, Extend limit 0.
- watsonx Assistant for Z Pilot - watsonx O...**: Status - Ready. This tile shows details: Start date Oct 30, 2024 8:10 AM, End date Nov 5, 2024 7:12 AM, Extend limit 0.

At the top right of the main area, there are filters for 'Filter by infrastructure' and 'Filter by status', along with icons for 'Custom request', 'Add to calendar', and view modes ('List', 'Grid').

3. Scroll down and record the **Cluster Admin Username** and **Cluster Admin Password**.

Reservation Details

API URL
<https://api.672371d38376796fb96a6c4d.ocp.techzone.ibm.com:6443>

Bastion Password [REDACTED]

Bastion RDP Address
api.672371d38376796fb96a6c4d.ocp.techzone.ibm.com:43389

Bastion SSH connection
ssh itzuser@api.672371d38376796fb96a6c4d.ocp.techzone.ibm.com -p 40222

Bastion Username
itzuser

Cluster Admin Username
kubeadmin

Cluster Admin Password [REDACTED]

OCP Console
<https://console.openshift-console.apps.672371d38376796fb96a6c4d.ocp.techzone.ibm.com>

OCP Version
4.14

vCenter
itzeu-vc.eu.cloud.techzone.ibm.com

Download kubeconfig

4. Click the **OCP Console** link.

Note, OCP stands for OpenShift Container Platform.

Reservation Details

API URL
<https://api.672371d38376796fb96a6c4d.ocp.techzone.ibm.com:6443>

Bastion Password [REDACTED]

Bastion RDP Address
api.672371d38376796fb96a6c4d.ocp.techzone.ibm.com:43389

Bastion SSH connection
ssh itzuser@api.672371d38376796fb96a6c4d.ocp.techzone.ibm.com -p 40222

Bastion Username
itzuser

Cluster Admin Username
kubeadmin

Cluster Admin Password [REDACTED]

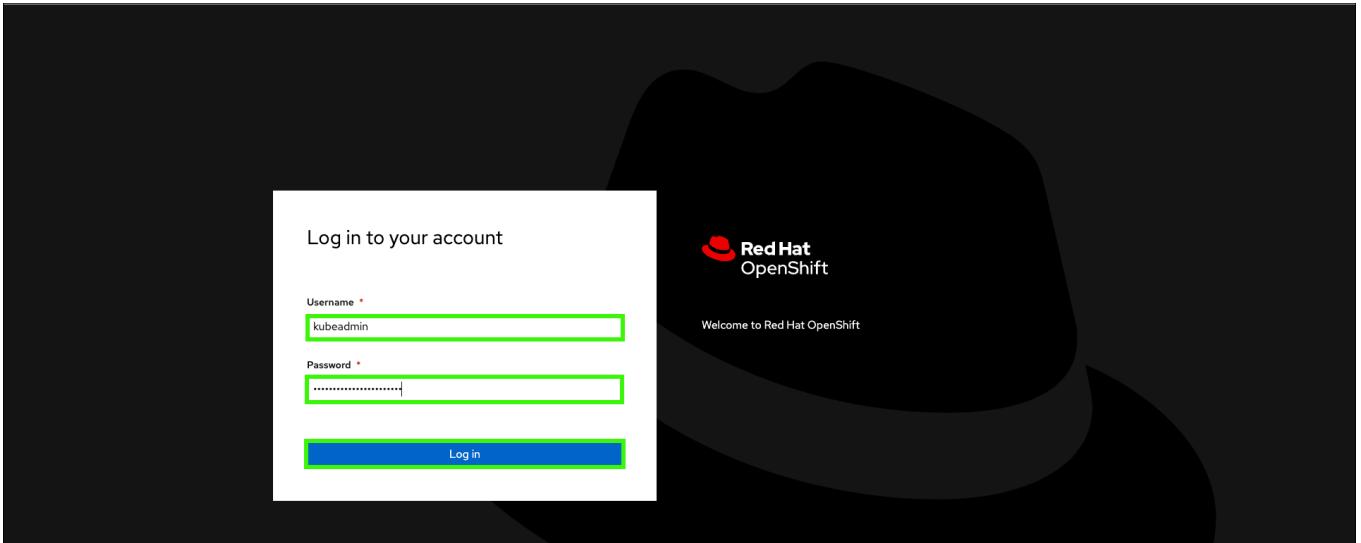
OCP Console
<https://console.openshift-console.apps.672371d38376796fb96a6c4d.ocp.techzone.ibm.com>

OCP Version
4.14

vCenter
itzeu-vc.eu.cloud.techzone.ibm.com

Download kubeconfig

5. Enter the **Cluster Admin Username** and **Cluster Admin Password** values from step 3 and click **Log in**.



6. Click the help (?) and then click **Command Line Tools**.

7. Click the link under **oc - OpenShift Command Line Interface (CLI)** for the operating system of your local machine.

Clicking the preceding link automatically downloads either a **.zip** or **.tar** file specific to your operating system. Unzip or untar the file and place the **oc** binary in a directory that is in your default PATH. The PATH will vary by operating system and individual settings, detailed instructions are not provided.

After properly installed, verify the installation by running the **oc** command on your local workstation.

```
oc --help
```

Sample output:

```
andrewjones@Andrews-MBP ~ % oc --help
OpenShift Client

This client helps you develop, build, deploy, and run your applications on any
OpenShift or Kubernetes cluster. It also includes the administrative
commands for managing a cluster under the 'adm' subcommand.

Basic Commands:
  login           Log in to a server
  new-project     Request a new project
  new-app          Create a new application
  status           Show an overview of the current project
  project          Switch to another project
  projects         Display existing projects
  explain          Get documentation for a resource

Build and Deploy Commands:
  rollout          Manage a Kubernetes deployment or OpenShift deployment
  config           Revert part of an application back to a previous deployment
```

Prepare to ingest documents

Before ingesting documents, a few setup steps must be taken.

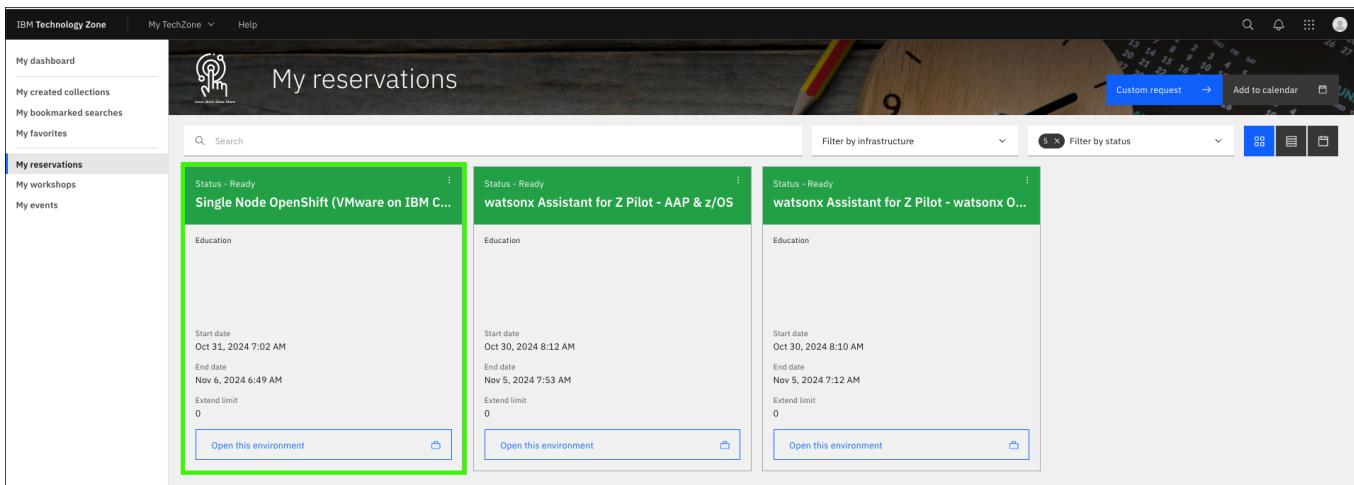
Login to the OpenShift cluster from your local terminal

Note: if you just installed the **oc** utility, you can skip the next 5 steps.

1. Click the following link to open a browser window to your ITZ reservations.

[ITZ My reservations](#)

2. Click the **Single Node OpenShift** tile.



3. Scroll to the bottom of the reservation page and record the **Cluster Admin Username** and **Cluster Admin Password**.

The screenshot shows the IBM Technology Zone interface. At the top, there are navigation links for 'IBM Technology Zone', 'My TechZone', and 'Help'. Below this, a search bar and a user icon are visible.

Two clusters are listed in the main area:

- 672371d38376796fb96a6c4d-master-1**: Red Hat Enterprise Linux 8 (64-bit), IP 192.168.252.11, Status Running. A 'Console' button is shown.
- 672371d38376796fb96a6c4d-nfs**: Red Hat Enterprise Linux 8 (64-bit), IP 192.168.252.17, Status Running. A 'Console' button is shown.

Below the clusters, the 'Reservation Details' section is expanded. It contains the following information:

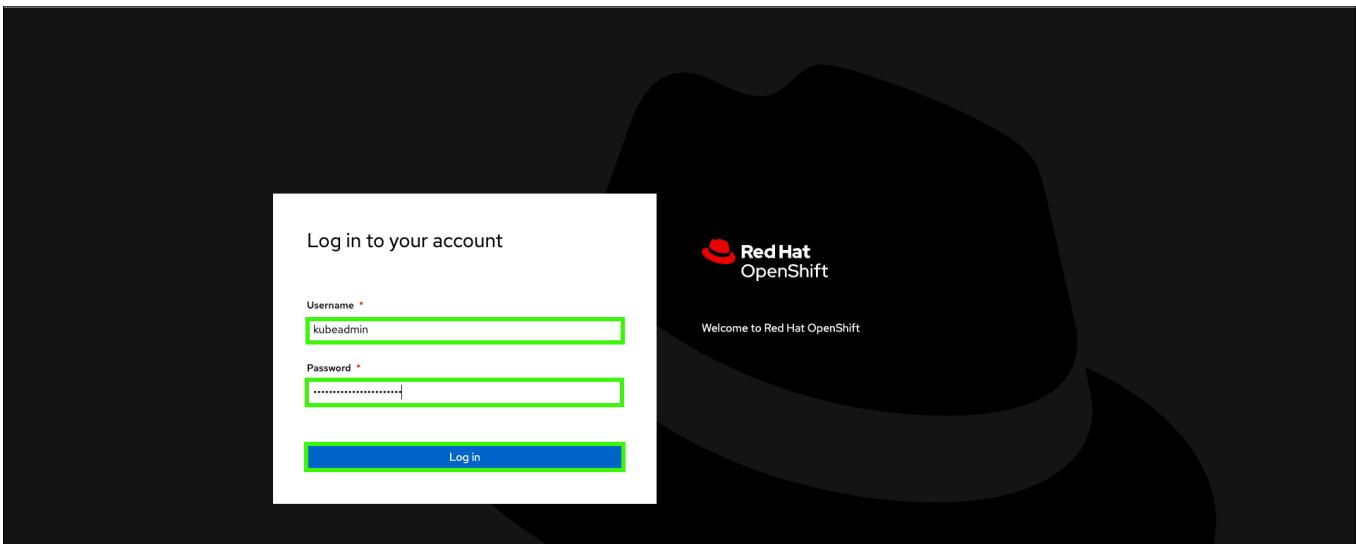
- API URL:** <https://api.672371d38376796fb96a6c4d.ocp.techzone.ibm.com:6443>
- Bastion Password:** [REDACTED]
- Bastion RDP Address:** api.672371d38376796fb96a6c4d.ocp.techzone.ibm.com:43389
- Bastion SSH Connection:** ssh itzuser@api.672371d38376796fb96a6c4d.ocp.techzone.ibm.com -p 40222
- Bastion Username:** itzuser
- Cluster Admin Username:** kubeadmin
- Cluster Admin Password:** [REDACTED]
- OCP Console:** <https://console.openshift-console.apps.672371d38376796fb96a6c4d.ocp.techzone.ibm.com> (This link is highlighted with a green box)
- OCP Version:** 4.14
- vCenter:** itzeu-vc.eu.cloud.techzone.ibm.com

A blue 'Download kubeconfig' button is at the bottom.

4. Click the OCP Console link.

This screenshot is identical to the one above, showing the same two clusters and reservation details. The 'OCP Console' link (<https://console.openshift-console.apps.672371d38376796fb96a6c4d.ocp.techzone.ibm.com>) is highlighted with a green box.

5. Enter the Cluster Admin Username and Cluster Admin Password values from step 3 and click Log in.



6. Click the **kube:admin** profile drop-down and click **Copy login command**.

You are logged in as a temporary administrative user. Update the cluster OAuth configuration to allow others to log in.

- [Copy login command](#)
- User Preferences
- Log out

7. Click **Display Token**.

8. Select and copy the **Log in with this token** string.

The steps to select and copy the value may differ by operating system, but for most, you can double-click the value and then right click and select **Copy**.

Your API token is
sha256~zuwR0KDnkYniIY0m8g8iKoUXPdFFFmou-o4s5FsrDNA

Log in with this token

```
oc login --token=sha256~zuwR0KDnkYniIY0m8g8iKoUXPdFFFmou-o4s5FsrDNA --server=https://api.672b79320c7a71b728e523b4.ocp.techzone.ir
```

Use this token directly against the API

```
curl -H "Authorization: Bearer sha256~zuwR0KDnkYniIY0m8g8iKoUXPdFFFmou-o4s5FsrDNA" "https://api.672b79320c7a71b728e523b4.ocp.techzone.ir/v1/users/~"
```

Request another token

Logout

9. Open a command prompt or terminal window on your local workstation.

10. Paste the login command and press **enter**.

```
andrewjones@Andrews-MBP ~ % oc login --token=sha256~mJ4L8K6cUMyNyk2Z69KMm3vbP1sWc8SW0eeOdVqtA94 --server=https://api.672371d38376796fb96a6c4d.ocp.techzone.ibm.com:6443
Logged into "https://api.672371d38376796fb96a6c4d.ocp.techzone.ibm.com:6443" as "kube:admin" using the token provided.
You have access to 70 projects, the list has been suppressed. You can list all projects with 'oc projects'
Using project "default".
andrewjones@Andrews-MBP ~ %
```

Create a working directory

11. Create a directory that will be used to store the configuration files that will be created in the next steps.



Instructions vary by your local workstation's operating system

The directions that follow may vary depending on your operating system. The examples provided are based upon MacOS.

```
mkdir watsonxAssistant
```

12. Change to the new directory.

```
cd watsonxAssistant
```

```
watsonxAssistant ~ % oc login --token=sha256~mJ4L8K6cUMyNyk2Z69KMm3vbP1sWc8SW0eeOdVqtA94 --server=https://api.672371d38376796fb96a6c4d.ocp.techzone.ibm.com:6443
Logged into "https://api.672371d38376796fb96a6c4d.ocp.techzone.ibm.com:6443" as "kube:admin" using the token provided.
You have access to 70 projects, the list has been suppressed. You can list all projects with 'oc projects'
Using project "default".
andrewjones@Andrews-MBP ~ % mkdir watsonxAssistant
andrewjones@Andrews-MBP ~ % cd watsonxAssistant
andrewjones@Andrews-MBP watsonxAssistant %
```

Install IBM Certificate Manager on Red Hat OpenShift

13. Create a file that is named **catalogCertManager.yaml** using a text editor and paste the following text into the file.



Formatting of the yaml file is critical!

The content of the YAML file must be formatted exactly as shown. Use the copy icon to prevent typographical errors.

File name:

```
catalogCertManager.yaml
```

File contents:

```
apiVersion: operators.coreos.com/v1alpha1
kind: CatalogSource
metadata:
  name: ibm-cert-manager-catalog
  namespace: openshift-marketplace
spec:
  displayName: ibm-cert-manager-4.2.7
  grpcPodConfig:
    securityContextConfig: restricted
  image: icr.io/cpopen/ibm-cert-manager-operator-
catalog@sha256:4dcf4ace4b5f166f83b31063f7e6404dbf78d8e98a9d4fcf52fedf576a55ca6c
  publisher: IBM
  sourceType: grpc
  updateStrategy:
    registryPoll:
      interval: 30m0s
```

14. Install the IBM Certificate Manager operator in the Red Hat OpenShift cluster.

```
oc apply -f catalogCertManager.yaml
```

The preceding command returns a message that states the **ibm-cert-manager-catalog** was created.

15. In the OpenShift web console, click **Operators** and select **OperatorHub**.

The screenshot shows the Red Hat OpenShift web console interface. The left sidebar has a navigation menu with items like Home, Operators, Workloads, and Networking. The Operators item is expanded, and its sub-item OperatorHub is selected and highlighted with a green box. The main content area displays the 'Command Line Tools' section, specifically the 'oc - OpenShift Command Line Interface (CLI)' part. It includes a 'Copy login command' button and a list of download links for the oc binary for various platforms.

16. Click the **Project** to pull-down menu and click the **Show default projects** toggle.

The screenshot shows the Red Hat OpenShift web console interface. The left sidebar has a navigation menu with items like Home, Operators, Workloads, and Networking. The Operators item is expanded, and its sub-item OperatorHub is selected and highlighted with a green box. The main content area shows a dropdown menu for 'Project: All Projects'. Inside the dropdown, there is a 'Show default projects' toggle switch, which is highlighted with a green box. Below the dropdown, a list of projects is shown, including 'All Projects', 'default', 'kube-node-lease', and 'kube-public'. A note at the bottom right of the dropdown says '624 items'.

17. Scroll down and select **openshift-marketplace**.

You are logged in as a temporary administrative user. Update the cluster OAuth configuration to allow others to log in.

Project: All Projects

openshift-kube-controller-manager-operator
openshift-kube-scheduler
openshift-kube-scheduler-operator
openshift-kube-storage-version-migrator
openshift-kube-storage-version-migrator-operator
openshift-machine-api
openshift-machine-config-operator
openshift-marketplace
openshift-monitoring
openshift-multus
openshift-network-diagnostics
openshift-network-node-identity
openshift-network-operator
openshift-node
openshift-nutanix-infra
openshift-oauth-apiserver

Community Community Marketplace Community

[DEPRECATED] CrowdStrike Operator
[DEPRECATED] Use the CrowdStrike Falcon Operator from the certified channel instead

[DEPRECATED] CrowdStrike Operator
[DEPRECATED] Use the CrowdStrike Falcon Operator from the certified channel instead

[DEPRECATED] Hazelcast Platform Operator
[DEPRECATED] Use the certified Hazelcast Operator instead

ABot
Abot Operator-v3.0.0
provided by Rehaca Technologies

Certified
Accuknox Operator
provided by Accuknox Inc.

Red Hat
Advanced Cluster Management for Kubernetes

624 items

Create Project

18. Enter IBM Cert Manager in the search field and then click the IBM Cert Manager tile.

You are logged in as a temporary administrative user. Update the cluster OAuth configuration to allow others to log in.

Project: openshift-marketplace

Discover Operators from the Kubernetes community and Red Hat partners, curated by Red Hat. You can purchase commercial software through Red Hat Marketplace. You can install Operators on your clusters to provide optional add-ons and shared services to your developers. After installation, the Operator capabilities will appear in the Developer Catalog providing a self-service experience.

All Items All Items

Q IBM Cert Manager X

ibm-cert-manager-4.2.7

IBM Cert Manager
provided by IBM
Operator for managing deployment of cert-manager service.

1 items

19. Click Install.

You are logged in as a temporary administrative user. Update the cluster OAuth configuration to allow others to log in.

Project: openshift-marketplace

Discover Operators from the Kubernetes community and Red Hat partners, curated by Red Hat. You can purchase commercial software through Red Hat Marketplace. You can install Operators on your clusters to provide optional add-ons and shared services to your developers. After installation, the Operator capabilities will appear in the Developer Catalog providing a self-service experience.

All Items All Items

Q IBM Cert Manager X

ibm-cert-manager-4.2.7

IBM Cert Manager
provided by IBM
Operator for managing deployment of cert-manager service.

IBM Cert Manager
4.2.7 provided by IBM

Install

Channel: v4.2

Version: 4.2.7

Capability level:
 Basic Install
 Seamless Upgrades
 Full Lifecycle
 Deep Insights
 Auto Pilot

Source: ibm-cert-manager-4.2.7

20. Keep the default settings and click Install.

You are logged in as a temporary administrative user. Update the [cluster OAuth configuration](#) to allow others to log in.

Install Operator

Install your Operator by subscribing to one of the update channels to keep the Operator up to date. The strategy determines either manual or automatic updates.

Update channel *

v4.2

Version *

4.2.7

Installation mode *

All namespaces on the cluster (default)
Operator will be available in all Namespaces.

A specific namespace on the cluster
This mode is not supported by this Operator

Installed Namespace *

Operator recommended Namespace: ibm-cert-manager

Select a Namespace

Namespace creation
Namespace ibm-cert-manager does not exist and will be created.

Update approval *

Automatic

Manual

Provided APIs

- IBM Cert Manager** provided by IBM
- CR CertificateRequest** Not available
- C Certificate** A Certificate resource should be created to ensure an up to date and signed x509 certificate is stored in the Kubernetes Secret resource named in `spec.secretName`. Documentation For additional details regarding install parameters check...
- CMC Cert Manager Config** CertManagerConfig is the Schema for the certmanagerconfigs API. Documentation For additional details regarding install parameters check: <https://ibm.biz/icpf39install>. License By installing this product you accept the license terms...
- Challenge** Not available
- ClusterIssuer** Not available
- Issuer** An issuer represents a certificate issuing authority which can be referenced as

Install **Cancel**

The installation process takes a few minutes. Do not continue until you see the message: **Installed operator: ready for use.**

You are logged in as a temporary administrative user. Update the [cluster OAuth configuration](#) to allow others to log in.

Installed operator: ready for use

View Operator [View installed Operators in Namespace ibm-cert-manager](#)

IBM Cert Manager
ibm-cert-manager-operator:v4.2.7 provided by IBM

Installed operator: ready for use

✓

Install the watsonx Assistant for Z Operator (for OpenSearch)

- In your command prompt or terminal window, create a new namespace called **wxa4z-byos** in the Red Hat OpenShift cluster.

```
oc create namespace wxa4z-byos
```

- Create or obtain your IBM Container Software **production entitlement key**.

A **production entitlement key** is required to pull the container images that get deployed by the operator.

To create this key, follow the [instructions here](#) to create or retrieve your existing entitlement key.

After locating your existing key or creating a new key, continue to the next step.

23. Copy the **production entitlement key**.

The screenshot shows the IBM Container Software and Cloud Pak Access Management interface. On the left, there's a sidebar with 'Entitlement keys' selected. The main area is titled 'Entitlement keys (1)'. It contains a section for 'Access your container software' and a note about entitlement keys. Below that is a 'Active entitlement keys' section with a note about using active keys to log in. A specific entitlement key is listed with an issue date of 'October 17, 2022'. At the bottom right of this key entry, there are 'Copy', 'Edit', and 'Delete' buttons, with 'Copy' being highlighted by a green box.

24. Using the following command in your command prompt or terminal window, set an environment variable with your **production entitlement key** and then create a pull secret for the Container Registry.

Substitute your **production entitlement key** for the <entitlement key> string.

```
export IBM_CS_ENT_KEY=<entitlement key>
```

```
oc -n wxa4z-byos create secret docker-registry icr-pull-secret --docker-server=cp.icr.io --  
docker-username=cp --docker-password=$IBM_CS_ENT_KEY
```

The terminal window shows the following commands and their output:

```
andrewjones@Andrews-MBP watsonxAssistant % oc create namespace wxa4z-byos  
namespace/wxa4z-byos created  
andrewjones@Andrews-MBP watsonxAssistant % export IBM_CS_ENT_KEY=  
andrewjones@Andrews-MBP watsonxAssistant % oc -n wxa4z-byos create secret docker-registry icr-pull-secret --docker-server=cp.icr.io --  
docker-username=cp --docker-password=$IBM_CS_ENT_KEY  
secret/icr-pull-secret created  
andrewjones@Andrews-MBP watsonxAssistant %
```

25. Create a file called **catalogSource.yaml** with the following content.



Formatting of the yaml file is critical!

The content of YAML files must be formatted exactly as shown. Use the copy icon to prevent typographical errors.

File name:

```
catalogSource.yaml
```

File contents:

```

apiVersion: operators.coreos.com/v1alpha1
kind: CatalogSource
metadata:
  name: ibm-wxa4z-operator-catalog
  namespace: wxa4z-byos
spec:
  displayName: "IBM watsonx Assistant for Z Operator Catalog"
  image: icr.io/cpopen/ibm-wxa4z-
catalog:2.0.2@sha256:d457aa24af07c23d7b36727a9cff95995ca1c4ed49fb980ef4f386fd09a29b4
  publisher: 'IBM'
  sourceType: grpc
  secrets:
    - icr-pull-secret

```

26. Create your document catalog in the Red Hat OpenShift operator.

```
oc apply -f catalogSource.yaml
```

```

andrewjones@Andrews-MBP watsonxAssistant % oc create namespace wxa4z-byos
namespace/wxa4z-byos created
andrewjones@Andrews-MBP watsonxAssistant % export IBM_CS_ENT_KEY=eyJ0eXAiOiJKV1QiLCJhbGciOiJIUzI1NiJ9eyJpc3MiOiJJQk0gTWFya2V0cGxhY2UiLCJpXXQiOjE2NjYwMTk1ODAsImp0aSI6IjNKOWUyMzzjZTAzMDQzMzVhNTJhYTkzMWNmOTcyMDR1In0.4M3XRD4XzkHMS0kFNJ4uKvcWZ6SnEA0Z03eL_11A2XY
andrewjones@Andrews-MBP watsonxAssistant % oc -n wxa4z-byos create secret docker-registry icr-pull-secret --docker-server=cp.icr.io --docker-username=cp --docker-password=$IBM_CS_ENT_KEY
secret/icr-pull-secret created
andrewjones@Andrews-MBP watsonxAssistant % vi catalogSource.yaml
andrewjones@Andrews-MBP watsonxAssistant % vi catalogSource.yaml
andrewjones@Andrews-MBP watsonxAssistant % oc apply -f catalogSource.yaml
catalogsource.operator.coreos.com/ibm-wxa4z-operator-catalog created:
andrewjones@Andrews-MBP watsonxAssistant %

```

27. In the Red Hat OpenShift web console, click **OperatorHub** and select the **wxa4z-byos** project.

The screenshot shows the Red Hat OpenShift web console interface. The left sidebar has a navigation menu with 'Administrator' at the top, followed by 'Home', 'Operators' (with 'OperatorHub' highlighted), 'Workloads', 'Networking', 'Storage', 'Builds', 'Observe', 'Compute', 'User Management', and 'Administration'. The main content area has a header 'Project: openshift-marketplace'. Below it, there's a search bar with 'wxa4z-byos' typed in. To the right of the search bar, it says '624 items'. There are several cards representing different operators:

- CrowdStrike**: [DEPRECATED] CrowdStrike Operator provided by CrowdStrike
- Hazelcast**: [DEPRECATED] Hazelcast Platform Operator provided by Hazelcast, Inc.
- Abot**: Abot Operator-v3.0.0 provided by Rebeca Technologies Pvt Ltd
- Accuknox**: Accuknox Operator provided by Accuknox Inc.
- Red Hat**: Advanced Cluster Management for Kubernetes provided by Red Hat

28. Enter **ibm watsonx** in the search field and the click the **IBM watsonx Assistant for Z Operator Catalog** tile.

Note, it may take a few minutes for the operator to appear. If the operator is not found initially, refresh the browser page and try the search again.

You are logged in as a temporary administrative user. Update the cluster OAuth configuration to allow others to log in.

Project: wxa4z-byos

OperatorHub

Discover Operators from the Kubernetes community and Red Hat partners, curated by Red Hat. You can purchase commercial software through Red Hat Marketplace. You can install Operators on your clusters to provide optional add-ons and shared services to your developers. After installation, the Operator capabilities will appear in the Developer Catalog providing a self-service experience.

All Items All Items

Search: ibm watsonx

1 items

IBM watsonx Assistant for Z Operator Catalog
IBM watsonx Assistant for Z provided by IBM
IBM watsonx Assistant for Z Operator

29. Click Install.

You are logged in as a temporary administrator.

Project: wxa4z-byos

OperatorHub

Discover Operators from the Kubernetes community and Red Hat partners, curated by Red Hat. You can purchase commercial software through Red Hat Marketplace. You can install Operators on your clusters to provide optional add-ons and shared services to your developers. After installation, the Operator capabilities will appear in the Developer Catalog providing a self-service experience.

All Items All Items

Search: ibm watsonx

IBM watsonx Assistant for Z
2.0.1 provided by IBM

Install

Channel: stable

Version: 2.0.1

Capability level:

- Basic Install
- Seamless Upgrades
- Full Lifecycle
- Deep Insights
- Auto Pilot

Source: IBM watsonx Assistant for Z Operator Catalog

Provider: IBM

Repository: N/A

Container Image: N/A

Created at: Oct 4, 2024, 1:31 PM

30. Select A specific namespace on the cluster under Installation mode and wxa4z-byos for the Installed Namespace, then click Install.

You are logged in as a temporary administrative user. Update the cluster OAuth configuration to allow others to log in.

Install Operator

Install your Operator by subscribing to one of the update channels to keep the Operator up to date. The strategy determines either manual or automatic updates.

Update channel * stable

Version * 2.0.1

Installation mode *

- All namespaces on the cluster (default)
Operator will be available in all Namespaces
- A specific namespace on the cluster
Operator will be available in a single Namespace only.

Installed Namespace * **wxa4z-byos**

Update approval *

- Automatic
- Manual

Install **Cancel**

The installation process takes a few minutes. Do not continue until you see the message: **Installed operator: ready for use.**

You are logged in as a temporary administrative user. Update the cluster OAuth configuration to allow others to log in.

OperatorHub

IBM watsonx Assistant for Z
ibm-wxa4z-operator:v2.0.1 provided by IBM

Installed operator: ready for use

View Operator View installed Operators in Namespace **wxa4z-byos**

- In your command prompt or terminal window, run the following commands to add the Container Registry credential to the operator's service account.

```
oc project wxa4z-byos
```

```
oc patch serviceaccount ibm-wxa4z-operator-controller-manager --type merge -p
'{"imagePullSecrets": [{"name": "icr-pull-secret"}]}'
```

```
watsonxAssistant -- cecuser@p1365-pvm1:~ -zsh - 135x37
andrewjones@Andrews-MBP watsonxAssistant % oc create namespace wxa4z-byos
namespace/wxa4z-byos created
andrewjones@Andrews-MBP watsonxAssistant % export IBM_CS_ENT_KEY=eyJ0eXAiOiJKV1QiLCJhbGciOiJIUzI1NiJ9.eyJpc3MiOiJJQk0gTWFya2V0cGxhY2UiLCJpYXQiOjE2NjYwMTk1ODAsImp0aS16IjNkOWUyMzjZTAzMDQzMzVhNTJhYTkzMWNmOTcyMDR1In0.4M3XRD4XzkHMS0kFNJ4uKVcwZ6SnEAZO3eL_11A2xY
andrewjones@Andrews-MBP watsonxAssistant % oc -n wxa4z-byos create secret docker-registry icr-pull-secret --docker-server=cp.icr.io --docker-username=cp --docker-password=$IBM_CS_ENT_KEY
secret/icr-pull-secret created
andrewjones@Andrews-MBP watsonxAssistant % vi catalogSource.yaml
andrewjones@Andrews-MBP watsonxAssistant % vi catalogSource.yaml
andrewjones@Andrews-MBP watsonxAssistant % oc apply -f catalogSource.yaml
catalogsource.operator.coreos.com/ibm-wxa4z-operator-catalog created
andrewjones@Andrews-MBP watsonxAssistant % oc project wxa4z-byos
Now using project "wxa4z-byos" on server "https://api.672371d38376796fb96a6c4d.ocp.techzone.ibm.com:6443".
andrewjones@Andrews-MBP watsonxAssistant % oc patch serviceaccount ibm-wxa4z-operator-controller-manager --type merge -p '{"imagePullSecrets": [{"name": "icr-pull-secret"}]}'
serviceaccount/ibm-wxa4z-operator-controller-manager patched
andrewjones@Andrews-MBP watsonxAssistant %
```

32. In the Red Hat OpenShift web console, click **Workloads** and select the **Pods**.

33. Verify the two pods that start with **ibm-wxa4z-operator** have a Status of **Running** and that all pods are **Ready**.

Name	Status	Ready	Restarts	Owner	Memory	CPU	Created
ld126367b1ca53dcf2b0c93acd733e33875cd4b6c382dd5eb412032ac2h6	Completed	0/1	0	ld126367b1ca53dcf2b0c93acd733e33875cd4b6c382dd5eb412032ac42b38	-	-	Nov 4, 2024, 4:44 PM
ibm-wxa4z-operator-catalog-n9m5m	Running	1/1	0	ibm-wxa4z-operator-catalog	19.0 MiB	0.003 cores	Nov 4, 2024, 4:32 PM
ibm-wxa4z-operator-controller-manager-7c7898d7d4-87hmt	Running	2/2	0	ibm-wxa4z-operator-controller-manager-7c7898d7d4	31.0 MiB	0.001 cores	Nov 4, 2024, 4:44 PM

34. Run the following command to set the administrative policy for the workspace.

```
oc -n wxa4z-byos adm policy add-scc-to-user privileged -z byos
```

```
andrewjones@Andrews-MBP watsonxAssistant % oc create namespace wxa4z-byos
namespace/wxa4z-byos created
andrewjones@Andrews-MBP watsonxAssistant % export IBM_CS_ENT_KEY=eyJ0eXAiOiJKV1QiLCJhbGciOiJIUzI1NiJ9eyJpc3MiOiJJQk0gTWFya2V0cGxhY2UiLCJpYXQiOjE2NjYwMTk1ODAsImp0aSI6ijNkOWUyMzJzTAzMDQzMzvhNTJhYTkzMWNmOTcyMDR1In0.4M3XRD4XzkHMS0kFNJ4uKVcWZ6SnEA0Z03eL_11A2xY
andrewjones@Andrews-MBP watsonxAssistant % oc -n wxa4z-byos create secret docker-registry icr-pull-secret --docker-server=cp.icr.io --docker-username=cp --docker-password=$IBM_CS_ENT_KEY
secret/icr-pull-secret created
andrewjones@Andrews-MBP watsonxAssistant % vi catalogSource.yaml
andrewjones@Andrews-MBP watsonxAssistant % vi catalogSource.yaml
andrewjones@Andrews-MBP watsonxAssistant % oc apply -f catalogSource.yaml
catalogsource.operators.coreos.com/ibm-wxa4z-operator-catalog created
andrewjones@Andrews-MBP watsonxAssistant % oc project wxa4z-byos
Now using project "wxa4z-byos" on server "https://api.672371d38376796fb96a6c4d.ocp.techzone.ibm.com:6443".
andrewjones@Andrews-MBP watsonxAssistant % oc patch serviceaccount ibm-wxa4z-operator-controller-manager --type merge -p '{"imagePullSecrets": [{"name": "icr-pull-secret"}]}'
serviceaccount/ibm-wxa4z-operator-controller-manager patched
andrewjones@Andrews-MBP watsonxAssistant % oc -n wxa4z-byos adm policy add-scc-to-user privileged -z byos
clusterrole.rbac.authorization.k8s.io/system:openshift:scc:privileged added: "byos"
andrewjones@Andrews-MBP watsonxAssistant %
```

Deploy required secrets and the custom bring-your-own-search (BYOSearch) resource

35. Create a file called **os-secret.yaml** with the following content.

File name:

```
os-secret.yaml
```

Substitute a secure password of your choosing for the string <OPENSEARCH_PASSWORD>.

File contents:

```
apiVersion: v1
stringData:
  password: <OPENSEARCH_PASSWORD>
kind: Secret
metadata:
  name: opensearch-creds
  namespace: wxa4z-byos
  type: Opaque
```

36. Create the secret by running the following command.

```
oc apply -f os-secret.yaml
```

37. Create a file called **client-ingestion-secret.yaml** with the following content.

File name:

```
client-ingestion-secret.yaml
```

Substitute a secure authentication key of your choosing for the string <CLIENT_INGESTION_AUTHKEY>. The authentication key can be a random password.

File contents:

```
apiVersion: v1
stringData:
  authkey: <CLIENT_INGESTION_AUTHKEY>
kind: Secret
metadata:
  name: client-ingestion-authkey
  namespace: wxa4z-byos
  type: Opaque
```

37. Create the secret by running the following command.

```
oc apply -f client-ingestion-secret.yaml
```

38. Obtain and record your cluster domain that is used for routes by running the following command.

```
oc -n openshift-ingress-operator get ingresscontroller default -o jsonpath=".status.domain"
```



Remove the trailing % from the cluster domain name.

The value returned for the cluster domain may include a % character at the end. Do not include the % in the next step!

Note, the output of the command will be a string similar to:

apps.672b79320c7a71b728e523b4.ocp.techzone.ibm.com

39. Create a file called **byos.yaml** with the following content.

File name:

```
byos.yaml
```

Substitute the domain name recorded in the previous step for the string <YOUR_CLUSTER_DOMAIN>. Do not include the % at the end of the domain name.

File contents:

```

apiVersion: wxa4z.watsonx.ibm.com/v1
kind: BYOSearch
metadata:
  name: byosearch
  namespace: wxa4z-byos
spec:
  imagePullSecrets:
    - name : icr-pull-secret
  namespace: wxa4z-byos
  clusterName: wxa4z-byos-cluster
  clusterDomain: <YOUR_CLUSTER_DOMAIN>

  opensearch:
    secretName: opensearch-creds

  persistence:
    enabled: true
    storageClass: "managed-nfs-storage"
    accessModes:
      - ReadWriteOnce
    size: 24Gi

  wrapper:
    createRoute: true
    resources:
      requests:
        cpu: 4
        memory: "500Mi"
      limits:
        cpu: 4
        memory: "1Gi"

  clientIngestion:
    secretName: client-ingestion-authkey

    resources:
      limits:
        cpu: "500m"
        memory: 2Gi
        nvidia.com/gpu: "0"
      requests:
        cpu: "500m"
        memory: 1Gi
        nvidia.com/gpu: "0"
    pvc:
      storageClass: "managed-nfs-storage"
      enabled: true
      size: 24Gi

```

40. Run the following command to deploy BYOS on your cluster.

```
oc apply -f byos.yaml
```

The BYOS deployment begins and might take up to 20 minutes or more for the images to download and the deployment to complete. You can check the status by looking at the Pods view in the OCP console.

Verify all the required pods are running

42. Verify that all pods have the status of **Running** or **Completed**.

Name	Status	Ready	Restarts	Owner	Memory	CPU	Created
da80909aab1563ff0c15b7c6d	Completed	0/1	0	da80909aab1563ff0c15b7c6da81 a8bfefabaa605065bb6c6b490e87 0e87f0k96l	-	-	Nov 6, 2024, 9:26 AM
ibm-wxa4z-operator-catalog-cfsv	Running	1/1	0	ibm-wxa4z-operator-catalog	25.0 MiB	0.004 cores	Nov 6, 2024, 9:25 AM
ibm-wxa4z-operator-controller-manager-556fcf98bb-4bhmk	Running	2/2	0	ibm-wxa4z-operator-controller-manager-556fcf98bb	142.8 MiB	0.023 cores	Nov 6, 2024, 9:26 AM
wxa4z-byos-cluster-0	Running	1/1	0	wxa4z-byos-cluster	1,095.8 MiB	0.017 cores	Nov 6, 2024, 10:22 AM
wxa4z-byos-cluster-1	Running	1/1	0	wxa4z-byos-cluster	1,169.1 MiB	0.016 cores	Nov 6, 2024, 10:22 AM
wxa4z-byos-cluster-2	Running	1/1	0	wxa4z-byos-cluster	1,113.3 MiB	0.015 cores	Nov 6, 2024, 10:22 AM
wxa4z-client-ingestion-7f9b8d6c58-9bzth	Running	1/1	0	wxa4z-client-ingestion-7f9b8d6c58	501.8 MiB	0.071 cores	Nov 6, 2024, 12:18 PM
wxa4z-opensearch-wrapper-5cb879f5fb-qw7qt	Running	1/1	0	wxa4z-opensearch-wrapper-5cb879f5fb	547.1 MiB	0.031 cores	Nov 6, 2024, 10:22 AM
wxa4z-snapshot-setup-job-nsqlz	Completed	0/1	0	wxa4z-snapshot-setup-job	-	-	Nov 6, 2024, 10:22 AM

When deployment completes and all the pods have a status of "Ready", "Running", or "Completed", the next step is to retrieve your BYOS endpoint URL.

43. In the OCP console, click **Networking** and then **Routes**.

Name	Status	Ready	Restarts	Owner	Memory	CPU	Created
da80909aab1563ff0c15b7c6d	Completed	0/1	0	da80909aab1563ff0c15b7c6da81 a8bfefabaa605065bb6c6b490e87 0e87f0k96l	-	-	Nov 6, 2024, 9:26 AM
ibm-wxa4z-operator-catalog-cfsv	Running	1/1	0	ibm-wxa4z-operator-catalog	25.0 MiB	0.004 cores	Nov 6, 2024, 9:25 AM
ibm-wxa4z-operator-controller-manager-556fcf98bb-4bhmk	Running	2/2	0	ibm-wxa4z-operator-controller-manager-556fcf98bb	141.2 MiB	0.022 cores	Nov 6, 2024, 9:26 AM
wxa4z-byos-cluster-0	Running	1/1	0	wxa4z-byos-cluster	1,095.8 MiB	0.016 cores	Nov 6, 2024, 10:22 AM
wxa4z-byos-cluster-1	Running	1/1	0	wxa4z-byos-cluster	1,169.1 MiB	0.016 cores	Nov 6, 2024, 10:22 AM
wxa4z-byos-cluster-2	Running	1/1	0	wxa4z-byos-cluster	1,113.3 MiB	0.015 cores	Nov 6, 2024, 10:22 AM
wxa4z-client-ingestion-7f9b8d6c58-9bzth	Running	1/1	0	wxa4z-client-ingestion-7f9b8d6c58	534.6 MiB	0.087 cores	Nov 6, 2024, 12:18 PM
wxa4z-opensearch-wrapper-5cb879f5fb-qw7qt	Running	1/1	0	wxa4z-opensearch-wrapper-5cb879f5fb	547.1 MiB	0.030 cores	Nov 6, 2024, 10:22 AM
wxa4z-snapshot-setup-job-nsqlz	Completed	0/1	0	wxa4z-snapshot-setup-job	-	-	Nov 6, 2024, 10:22 AM

44. Copy and record the location for the **wxa4z-opensearch-wrapper** route.

Name	Status	Location	Service
wxa4z-client-ingestion	Accepted	https://wxa4z-client-ingestion-wxa4z-byos.apps.672b79320c7a71b728e523b4.ocp.techzone.ibm.com	wxa4z-client-ingestion
wxa4z-opensearch-wrapper	Accepted	https://wxa4z-opensearch-wrapper-wxa4z-byos.apps.672b79320c7a71b728e523b4.ocp.techzone.ibm.com	wxa4z-opensearch-wrapper

Update your assistant with the new BYOS search instance

45. Configure your assistant with the route recorded in the previous step.

First, append the string `**/v1/query**` to complete the URL endpoint. The URL should look similar to:

```
**https://wxa4z-opensearch-wrapper-wxa4z-
byos.apps.672b79320c7a71b728e523b4.ocp.techzone.ibm.com/v1/query**
```

Note, the above URL will not work for you. Use the value of your specific OpenSearch instance that is recorded in the previous step.

Next, you need to return to your assistant in the watsonx Orchestrate AI assistant builder and update the custom search integration URL. The steps to update the URL are illustrated in the animated gif that follows. You can see the steps [here](creatingAssistant-configuringConvoSearch.md#configureCustomSearchURL).

??? Tip "How to set the custom search integration URL."

Troubleshooting

The following are issues that you may encounter. If the provided resolutions do not work, contact support by using the methods that are mentioned in the [Support](#) section.

✖️ Pods have a status of ErrImagePull or ImagePullBackoff

If the pods starting with **ibm-wxa4z-operator** have a status of "ErrImagePull" or "ImagePullBackoff", you can delete the pod and it will automatically restart and pull the image successfully. Wait until the pod is re-created successfully.

 The wxa4z-client-ingestion pod does not start 

Did you include teh % character in the **clusterDomain** name when creating the **byos.yaml**? To resolve, edit the **byos.yaml** file and run the following command again. The current pod will be terminated and a new one started. This will take about 20 minutes to start.

```
oc apply -f byos.yaml
```

Installing and using zassist to ingest client documents

With bring-your-own-search (BYOS) installed and configured in your assistant, you can now prepare for document ingestion. Currently, only PDF, HTML, and DOCX file formats are supported for ingestion.

To prepare for document ingestion, you can also reference the setup instructions that are located [here](#).

Install the zassist utility

The **zassist** utility is an executable program that automates the ingestion of client documentation into the RAG for watsonx Assistant for Z. The utility is available to clients through [IBM Passport Advantage](#).

A version of zassist is available for download for IBMers and Business Partners for conducting pilots. Follow the steps that follow to download and install **zassist**.

1. Click the following link and download the **zassist.zip** file.

<https://ibm.box.com/s/j3nt5iw4fqd5w2jgcqwxnjlsu8bpvl77>

The screenshot shows a file browser interface for Box@IBM. At the top, there's a header with 'Box@IBM' and a search bar. Below the header, a folder named 'zassist.zip' is displayed. To the right of the folder name is a 'Download' button, which is highlighted with a green box. The main area shows a list of files and folders under 'zassist.zip'. There is one item listed: 'zassist' (a folder). The list includes columns for Name, Modified, and Size.

Name	Modified	Size
zassist	Sep 5 at 9:15 AM	--

2. Extract the **zassist.zip** file.

3. Locate the appropriate file for your local workstation's operating system.

The screenshot shows a file browser interface for Watsonx Assistant. The title bar says 'watsonxAssistant'. The main area displays a file tree. A folder named 'zassist' is highlighted with a green box. Inside the 'zassist' folder, there are sub-folders: 'linux', 'mac', 'windows', and 'mac-arm'. A specific file within the 'mac' folder, 'zassist', is also highlighted with a green box. The file tree has a blue selection bar underneath it. The columns in the file list are 'Name', 'Date Modified', 'Size', and 'Kind'.

Name	Date Modified	Size	Kind
byos.yaml	Today at 12:17 PM	956 bytes	YAML
ca.crt	Today at 11:36 AM	Zero bytes	certificate
catalogCertManager.yaml	Today at 9:15 AM	476 bytes	YAML
catalogSource.yaml	Today at 9:25 AM	397 bytes	YAML
client-ingestion-secret.yaml	Today at 10:05 AM	155 bytes	YAML
os-secret.yaml	Today at 10:04 AM	149 bytes	YAML
zassist	Sep 5, 2024 at 11:15 AM	--	Folder
linux	Sep 18, 2024 at 8:15 AM	--	Folder
mac	Sep 18, 2024 at 8:15 AM	--	Folder
zassist	Sep 18, 2024 at 8:17 AM	10.2 MB	Unix Executable File
mac-arm	Sep 18, 2024 at 8:15 AM	--	Folder
windows	Sep 18, 2024 at 8:15 AM	--	Folder
zassist.zip	Today at 1:47 PM	23 MB	ZIP archive

4. Either copy the appropriate **zassist** file to a directory in your PATH, or copy it to a directory and add that directory to your PATH environment variable.

Additional information for running the preceding tasks can be found [here](#).

5. Verify that the **zassist** utility is working.



```
andrewjones@Andrews-MacBook-Pro watsonxAssistant % PATH=$PATH:.
andrewjones@Andrews-MacBook-Pro watsonxAssistant % zassist
zassist: error: expected one of "version", "init", "login", "ingest", "load", ...
andrewjones@Andrews-MacBook-Pro watsonxAssistant %
```

Ingest client documentation using **zassist**

With the **zassist** command installed, you are now able to begin ingesting data.

Step-by-step guidance for ingesting documents using zassist are provided in the IBM watsonx Assistant for Z documentation. In the directions, when running the log in command in step 7, use the **CLIENT_INGESTION_AUTHKEY** you created in the [previous section](#).

6. Follow the directions [here](#) to ingest documents using zassist.

These steps are not repeated in this lab guide. The following video illustrates the steps to ingest a single document. This video has no audio.

Don't see the video in the PDF version of the lab guide?



If you are viewing the [PDF](#) of the lab guide, you can access the video [here](#).

The document that is ingested in the video is a compressed [PDF](#) of the **IBM z/OS Continuous Delivery Red Piece**. You can download a copy of this document [here](#).

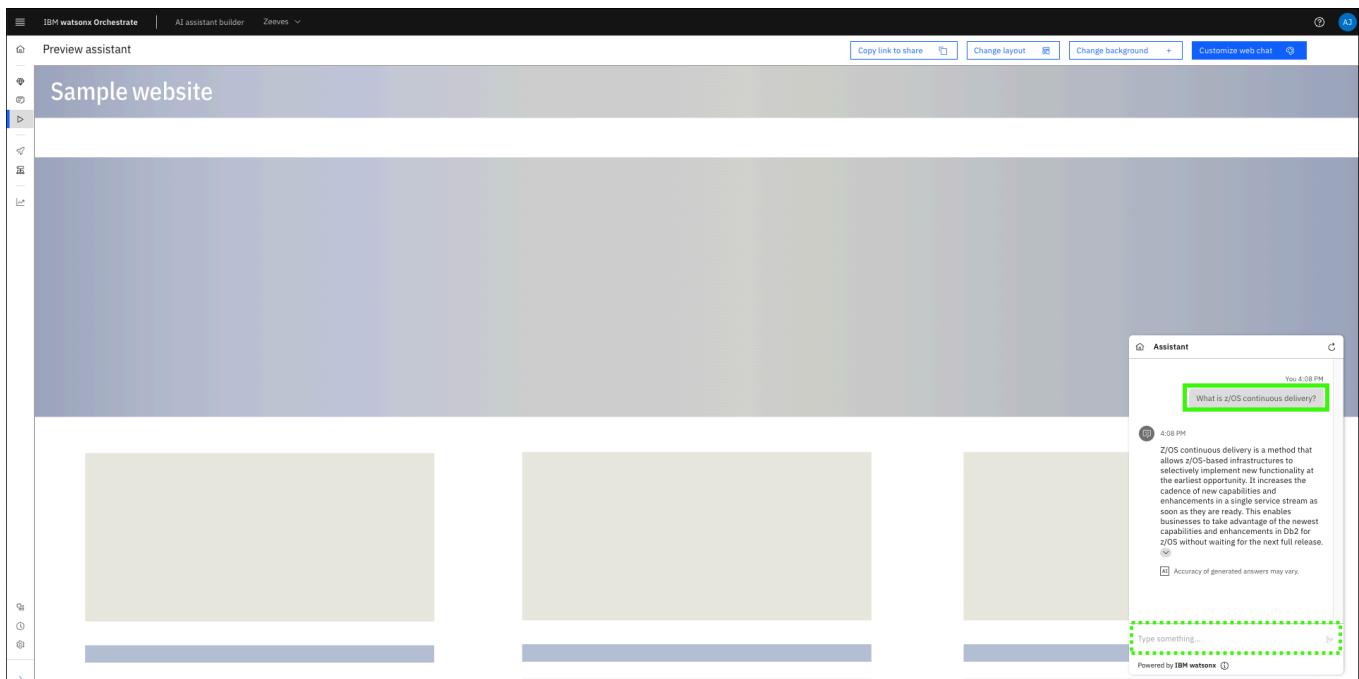
Verify the document that is ingested is now returned as a source file for a query.

Use the watsonx Orchestrate AI assistant builder to verify your document ingestion.

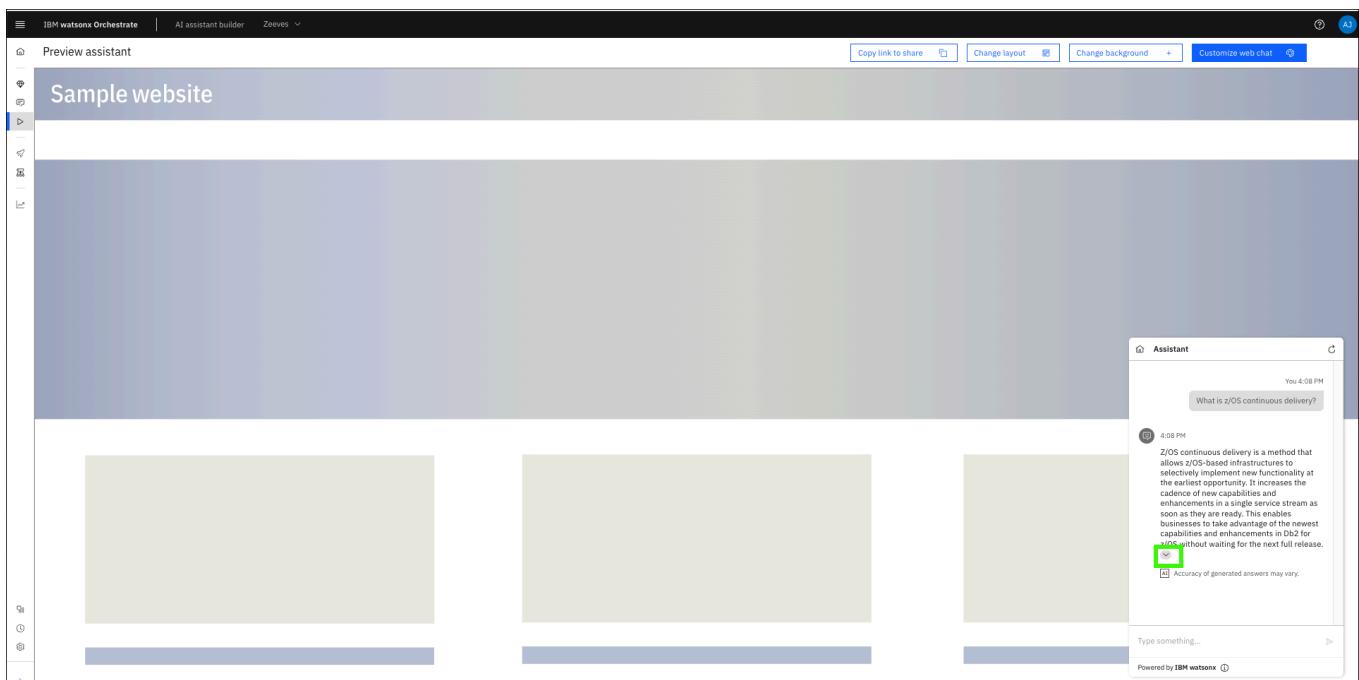
7. Hover over the home () icon and click **Preview**.

8. Enter the following prompt in your assistant.

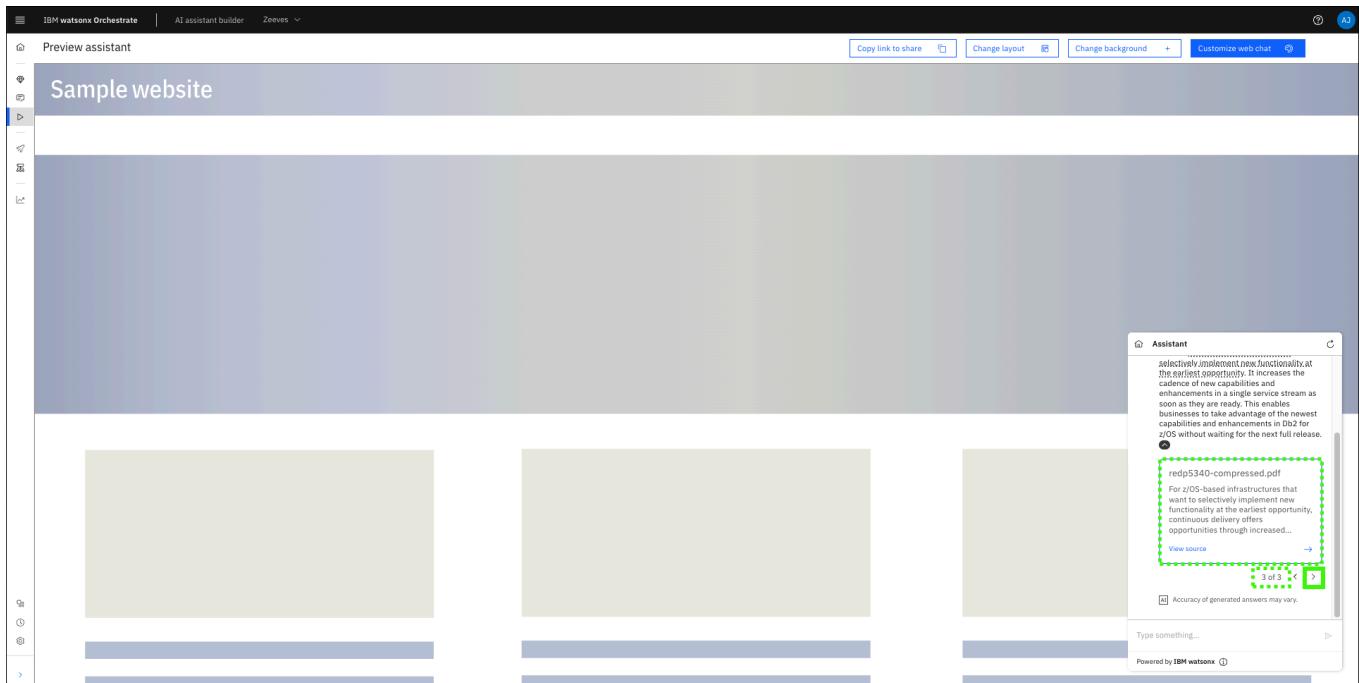
What is z/OS continuous delivery?



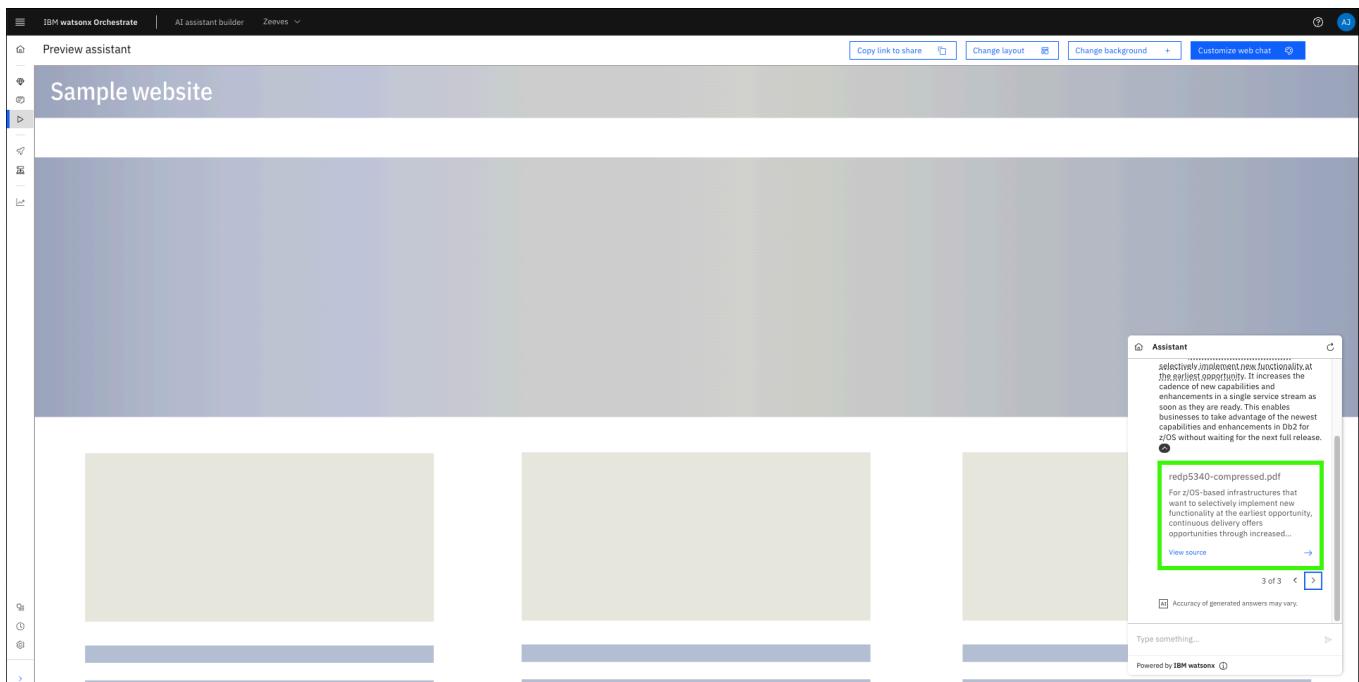
9. Expand the sources section by clicking the (▼).



10. Click through the list of resources and find the reference to the Red Piece document you ingested.

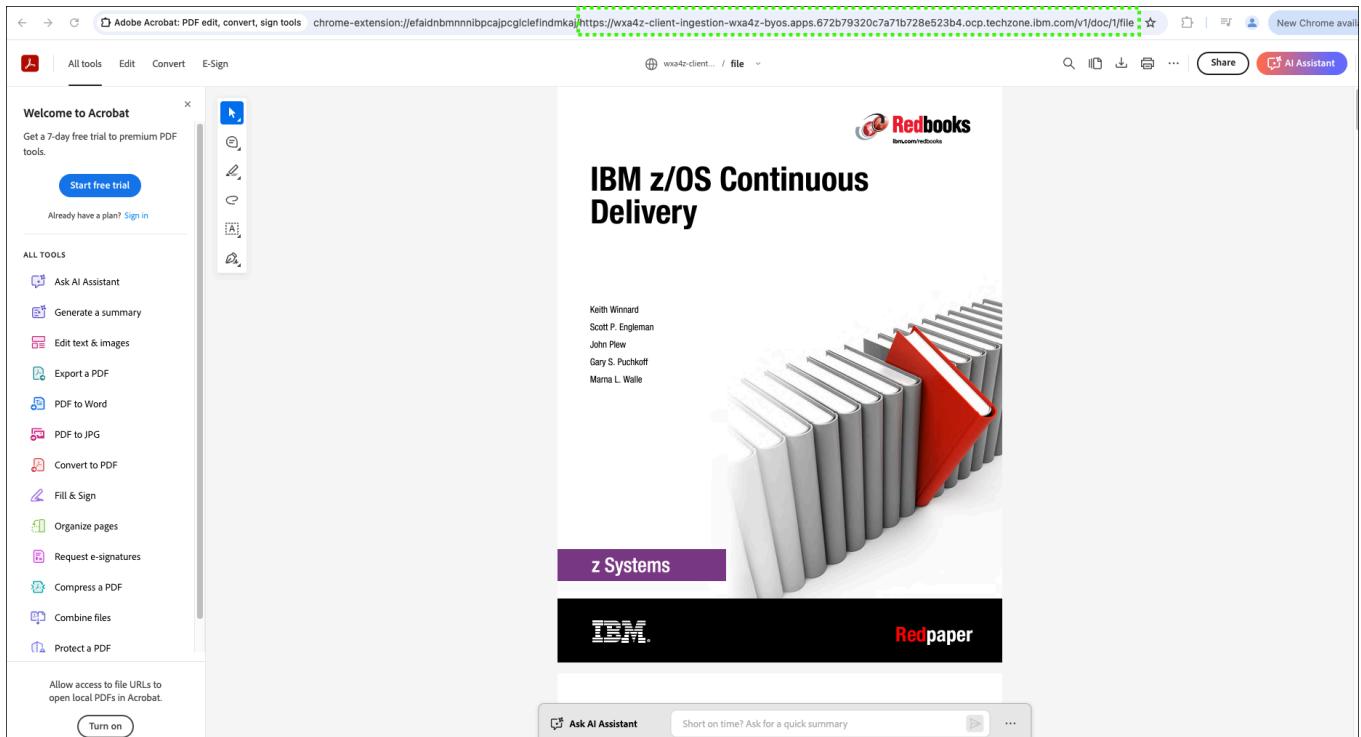


11. Click the ingested document reference.



12. Accept the security risk to view the source document.

The steps to accept the security risk for the document are not shown. The risk occurs because the certificate for the connection to the SNO instance is not secure. Notice that the URL contains the path to your SNO instance route.



Adjusting the search behavior

Do you recall the **Metadata** field when you configured your assistant?

The screenshot shows the 'Custom service' configuration page in the IBM watsonx Orchestrate interface. The 'Metadata' field is highlighted with a green dashed box. The JSON example provided is:

```
{
  "example_field": "example_value",
  "other_example_field": 7
}
```

The Metadata field provides a way to adjust your assistant's behavior during conversational search for your OpenSearch instance. Now that you have your own docs that are ingested for conversational search, you can set the metadata field for your assistant to use those documents in its content-grounded search. If you leave the metadata field empty, then it defaults to settings found to perform well. This replaces having to paste a complicated search string. By default (without any string in the Metadata field), it searches all the default IBM provided documentation and all ingested customer documentation using the following value:

```
{"ibm_indices": "*_ibm_docs_slate",
"customer_indices": "customer_*"}
```

Replacing the wildcard string with an explicit list of indices allows personalization. The metadata setting is where you can input specific indices (pointing to the underlying documentation) that you want your assistant to use for the content-grounded search. Out of the box there are over 220 products and topics that the OpenSearch instance has IBM Documentation for. You can find those indices and products [here](#).

You can input a subset of indices into the “Metadata” field in cases where you only want your assistant to gather context for specific IBM products or topics. The specific indices can be listed out in this format:

```
{"ibm_indices": "<comma separated index values>", "customer_indices": "customer_*"}
```

For example, if you only want your assistant to reference documentation for “Db2 Analytics Accelerator for z/OS” and no ingested client documentation, you can enter the following into the metadata field:

```
{"ibm_indices": "ss4lq8_ibm_docs_slate"}
```

If you have a mix of IBM Documentation and client documentation ingested, then there's an optional search string that you can use to set the “weights” used for each.

For example:

```
{"doc_weight": {
"product_docs": 0.5,
"customer_docs": 0.5},
"ibm_indices": "*_ibm_docs_slate",
"customer_indices": "customer_*"
}
```

In this case, “product_docs” is the weight that is assigned to “ibm_indices” and “customer_docs” is the weight that is assigned to “customer_indices”.

After you have configured all the settings for Conversational Search on the page, click “Save” in the upper-right of the page.

For more information on customizing the metadata field for conversational search, refer to this supplemental video found [here](#).

You are encouraged to experiment with the metadata field!

Try setting the metadata field to the following, which weights ingested docs higher than the product docs:

```
{"doc_weight":  
{"product_docs":0.2,  
"customer_docs":0.8},  
"ibm_indices":"*_ibm_docs_slate",  
"customer_indices":"customer_**"  
}
```

Now, repeat steps 6 through 8 (make sure you start a fresh preview of the assistant). Notice that the ingested Red Piece document is now the first sited reference!

Guidance for the pilot environment

If you or your client have other documents to ingest, you can do so by repeating the steps using zassist. The Velocity Pilot ITZ environment is limited in compute and storage capacity. The following limits should be adhered to:

- Greater than 15,000 documents is not supported.
- Individual documents should be less than 10 megabytes (MB).
- Larger documents greater than 50 MB should work if the document contains mostly images.
- Loading documents can take a long time, especially with > 100 MB of text.
- It is recommended to run large loads late at night.
- When loading, ensure your workstations does not sleep during the process.
- If you receive a **batch time error**, set the batch size to a lower number for that command. For example:

```
zassist ingest . -s 50
```

After ingesting all the your additional documents, proceed to the next section to learn about adding skills to your assistant.

Adding skills to the assistant

Getting started with skills and actions

Watsonx Assistant for Z can import skills to automate a range of IBM Z related tasks through assistant interactions. Each skill is a pre-defined automation that runs tasks to accomplish some unit or units of work. For example, you can use a skill to view z/OS IPL information, or work with z/OS datasets.

IBM Watsonx Assistant for Z provides an extension within Watsonx Orchestrate to build new skills from Ansible Automation platform or z/OS Management Facility (z/OSMF). The Z Skills Accelerator extension connects Ansible and z/OS application programming interfaces (APIs) and can import automation as Ansible Playbooks, JCL, or REXX as skills. Learn more importing and building skills [here](#).

Environments

Watsonx Orchestrate

The Z Skills Accelerator extension is already configured in your Watsonx Orchestrate IBM Technology Zone (ITZ) environment. You can use this component to import new skills.

Ansible Automation Platform and Wazi as a Service

To import skills for automations, you use Ansible Automation Platform (AAP) and Wazi as a Service (Wazi aaS) to serve as the z/OS back-end. Learn more about AAP [here](#). Learn more about Wazi, [here](#).

The two resources are provisioned together in the ITZ environment that you reserved earlier. This environment enables the ability to manage and automate z/OS tasks and subsystems with various preinstalled ansible playbooks. It includes a z/OS back-end (Wazi as a Service) with all needed prerequisites.

The playbooks provided cover various use cases for automating z/OS management. Ansible's capabilities for automating various Z-specific tasks are not limited to the use cases that are preinstalled in the AAP instance. The preinstalled playbooks are tasks from the 'IBM z/OS core collection'. Using this environment accelerates the ability to showcase the value of Watsonx Assistant for Z, and to get started with simple automations that can be expanded.

The ITZ environment gives you access to AAP, which is preconfigured to target the accompanying z/OS Wazi system, along with web-based access to AAP to experiment with different playbook templates. These templates are imported into Watsonx Orchestrate as skills and connected to your assistant.

For more information on the AAP and Wazi z/OS environments, refer to this document.

The playbook templates that are preinstalled in AAP cover various use cases, which you can explore, including:

- z/OS Certificate Management (create, delete, list, and renew certificates)
- dataset management (create, delete, fetch datasets)
- Submit JCL
- Run Operator commands

- Run TSO commands
- And more

Each of the sections that follow build upon each other. Complete each to successfully enhance your assistant by starting with [Explore Ansible Automation Platform](#).

Explore Ansible Automation Platform

After you reserved the Ansible Automation Platform (AAP) and Wazi z/OS environment in IBM Technology Zone (ITZ) and the reservation is in the **Ready** state, follow these steps to explore AAP.

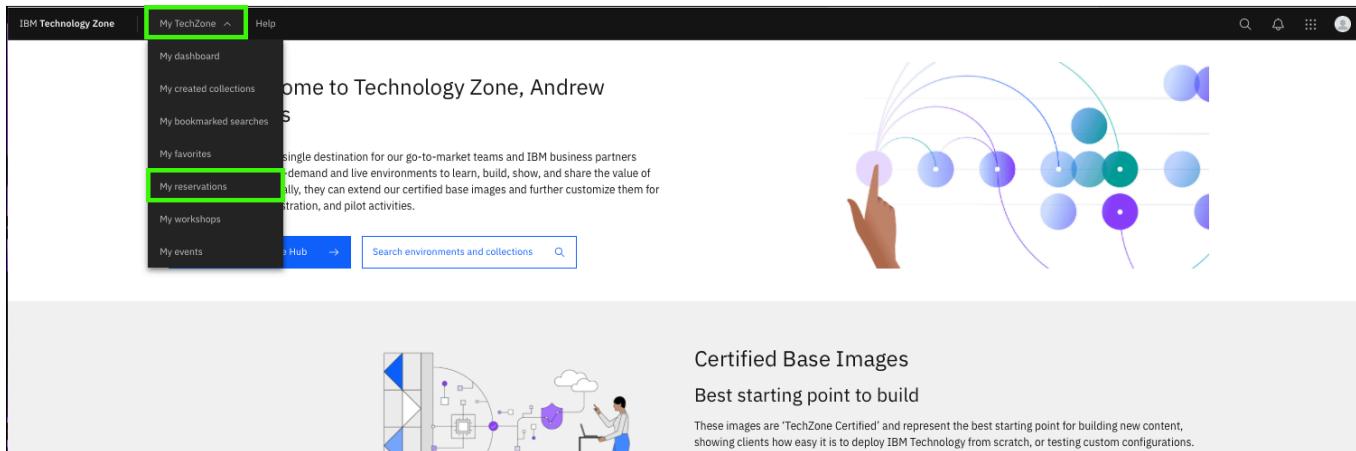
Access the AAP and Wazi as a Service environment

Be sure to record the information as instructed

Several of the steps below instruct you to record values from your ITZ reservation. Be sure to do this as they will not only be used in this section, but also in later sections of the lab guide.

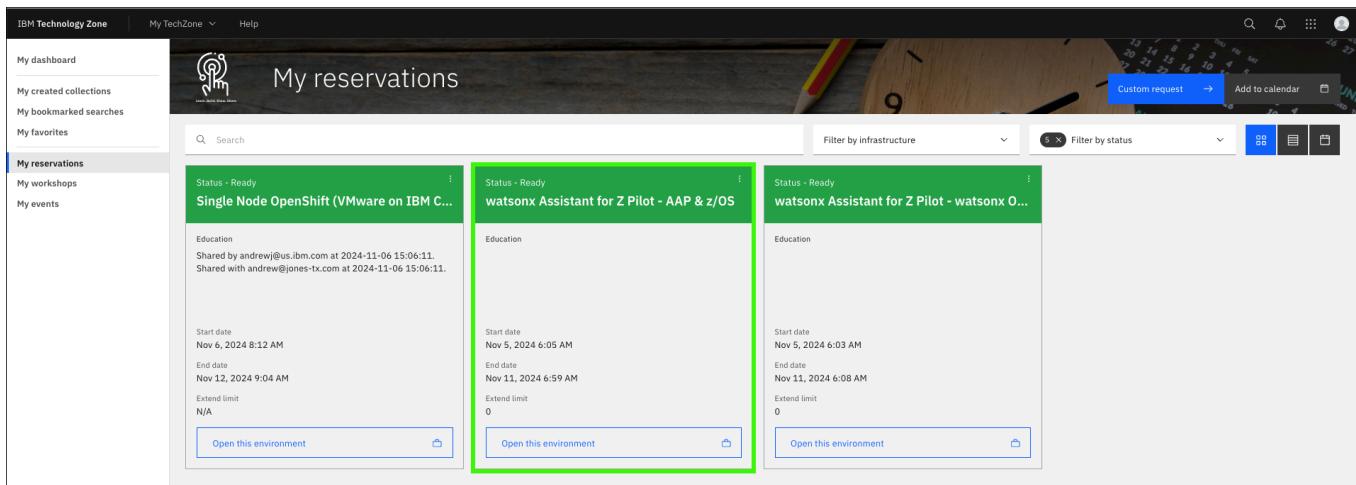
1. In the IBM Technology Zone portal, expand **My TechZone** and select **My Reservations**, or click the following link.

[ITZ My reservations](#)



Certified Base Images
Best starting point to build
These images are "TechZone Certified" and represent the best starting point for building new content, showing clients how easy it is to deploy IBM Technology from scratch, or testing custom configurations.

2. Click the **watsonx Assistant for Z Pilot - AAP & z/OS** tile.



3. Locate and record the **AAP User Name (For UI access)** and **AAP User Password** fields.

The screenshot shows the IBM Technology Zone interface with the following details:

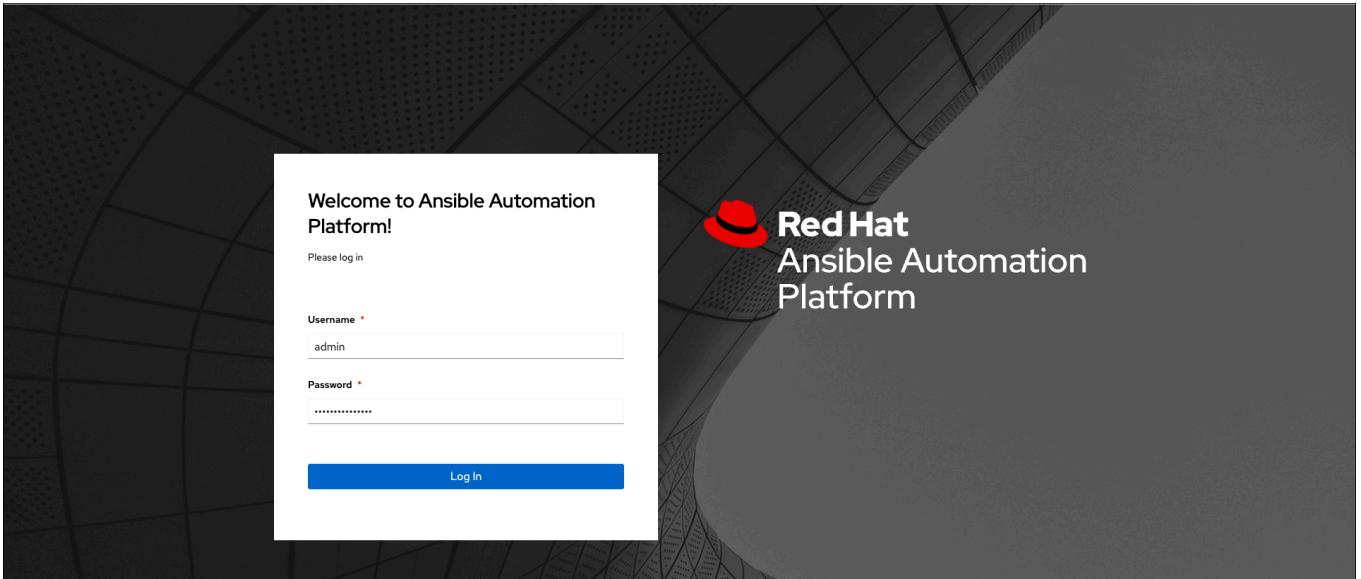
- Idle runtime limit:** 10800
- Reservation Details:**
 - Ansbile Automation Platform Cockpit UI: <https://itvsi-aap-phdhuz4.techzone.ibm.com:9090>
 - Ansbile Automation Platform Private IP: 10.241.0.6
 - Ansbile Automation Platform Public IP: 150.239.83.37
 - Ansbile Automation Platform UI: <https://itvsi-aap-phdhuz4.techzone.ibm.com>
 - AAP User Name (For SSH access): U6VO7N3
 - AAP User Name (For UI access): admin
 - AAP User Password: 39hQc-NiS458-Wa
 - AAP VSI Management: https://cloud.ibm.com/vpc-ext/compute/vs/us-east-0757_d038c79d-2fc8-45fd-8429-5bc642623771?view=edit&bs_account=7fd16281c12c4681b91417749e322845
 - z/OS Private IP: 10.241.0.7
 - z/OS Public IP: 169.63.103.211
 - z/OS User Name (Use SSH key to login, not password): IBMUSER
 - z/OS VSI Management: https://cloud.ibm.com/vpc-ext/compute/vs/us-east-0757_d729cb11-98ff-4453-8f74-acc8c595be1?view=edit&bs_account=7fd16281c12c4681b91417749e322845
 - User SSH key: [dropdown menu]

4. Record and then click the Ansible Automation Platform UI link.

The screenshot shows the IBM Technology Zone interface with the following details, identical to the previous one except for the highlighted UI link:

- Idle runtime limit:** 10800
- Reservation Details:**
 - Ansbile Automation Platform Cockpit UI: <https://itvsi-aap-phdhuz4.techzone.ibm.com:9090>
 - Ansbile Automation Platform Private IP: 10.241.0.6
 - Ansbile Automation Platform Public IP: 150.239.83.37
 - Ansbile Automation Platform UI: <https://itvsi-aap-phdhuz4.techzone.ibm.com>
 - AAP User Name (For SSH access): U6VO7N3
 - AAP User Name (For UI access): admin
 - AAP User Password: 39hQc-NiS458-Wa
 - AAP VSI Management: https://cloud.ibm.com/vpc-ext/compute/vs/us-east-0757_d038c79d-2fc8-45fd-8429-5bc642623771?view=edit&bs_account=7fd16281c12c4681b91417749e322845
 - z/OS Private IP: 10.241.0.7
 - z/OS Public IP: 169.63.103.211
 - z/OS User Name (Use SSH key to login, not password): IBMUSER
 - z/OS VSI Management: https://cloud.ibm.com/vpc-ext/compute/vs/us-east-0757_d729cb11-98ff-4453-8f74-acc8c595be1?view=edit&bs_account=7fd16281c12c4681b91417749e322845
 - User SSH key: [dropdown menu]

5. Enter the **Username** and **Password** that is recorded in step 3 and click **Log In**.



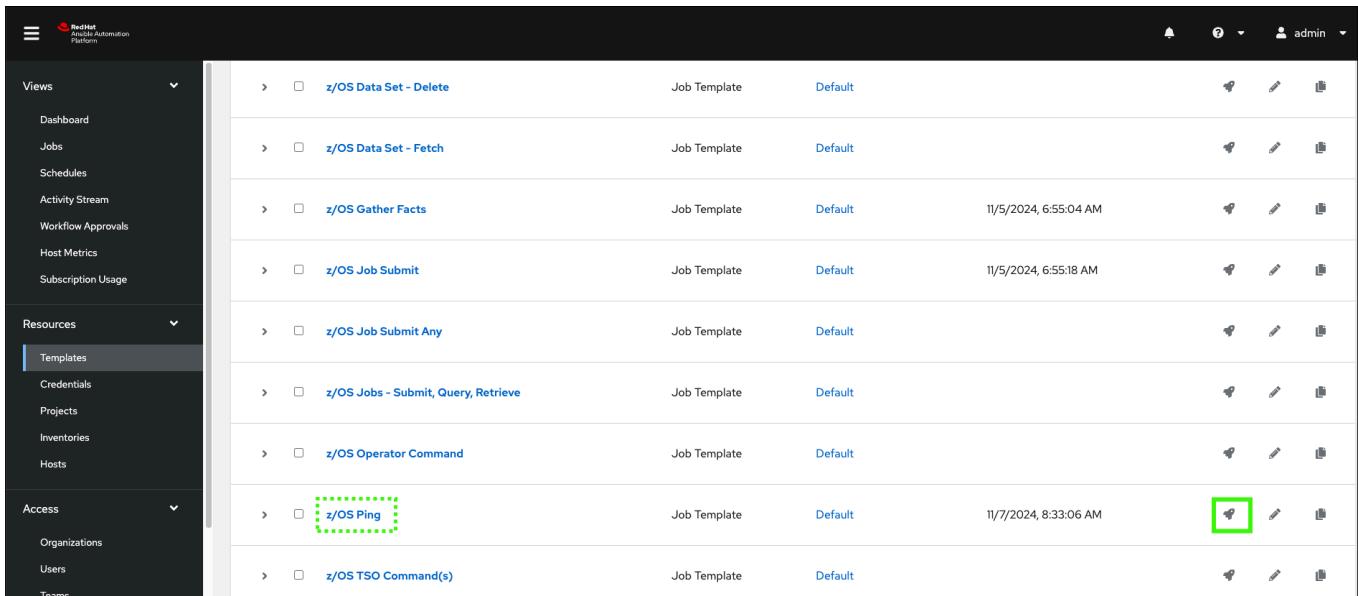
6. Click **Templates** under the **Resources** section.



The AAP instance is preconfigured to the Wazi aaS instance

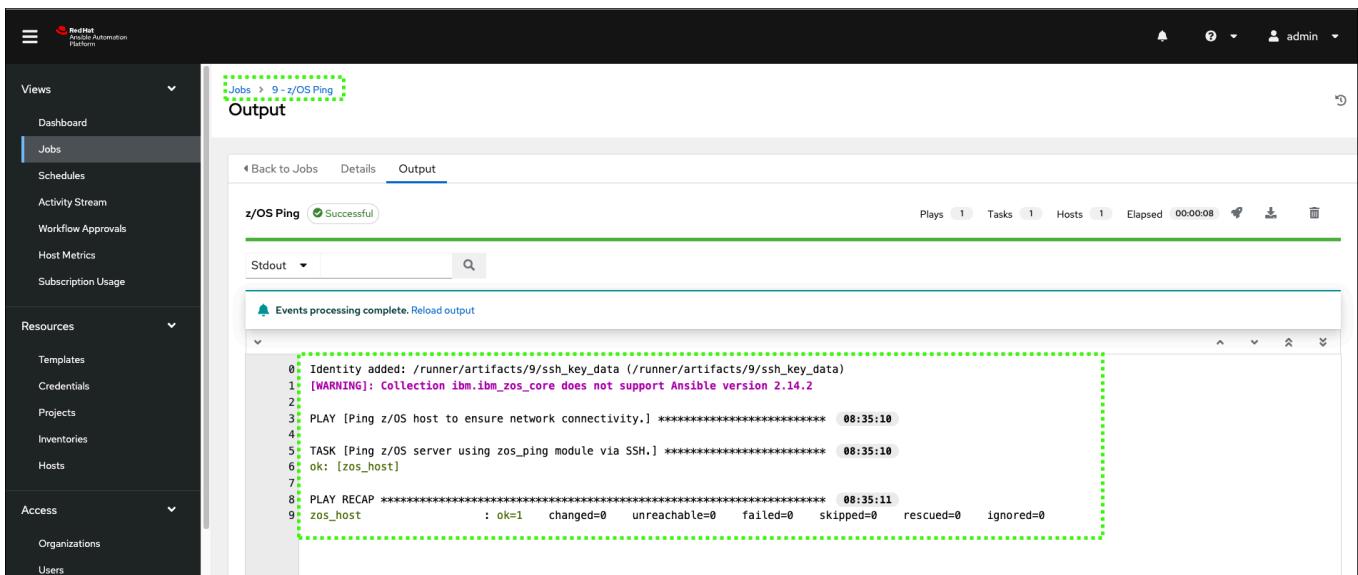
Note that because the AAP instance and the back-end z/OS system are preconfigured, no changes are needed to execute the templates and they will target your provisioned z/OS system by default.

7. Locate the **z/OS Ping** template and click the rocket (🚀) icon to start the template.



Views	z/OS Data Set - Delete	Job Template	Default			
Dashboard	z/OS Data Set - Fetch	Job Template	Default			
Schedules	z/OS Gather Facts	Job Template	Default	1/5/2024, 6:55:04 AM		
Activity Stream	z/OS Job Submit	Job Template	Default	1/5/2024, 6:55:18 AM		
Workflow Approvals	z/OS Job Submit Any	Job Template	Default			
Host Metrics	z/OS Jobs - Submit, Query, Retrieve	Job Template	Default			
Subscription Usage	z/OS Operator Command	Job Template	Default			
Resources	z/OS Ping	Job Template	Default	1/7/2024, 8:33:06 AM		
Access	z/OS TSO Command(s)	Job Template	Default			

8. Observe the z/OS Ping job run.



Jobs > 9 - z/OS Ping

Output

z/OS Ping Successful

Stdout

```

Events processing complete. Reload output
-----
0 Identity added: /runner/artifacts/9/ssh_key_data (/runner/artifacts/9/ssh_key_data)
1 [WARNING]: Collection ibm.ibm_zos_core does not support Ansible version 2.14.2
2
3 PLAY [Ping z/OS host to ensure network connectivity.] **** 08:35:10
4
5 TASK [Ping z/OS server using zos_ping module via SSH.] **** 08:35:10
6 ok: [zos_host]
7
8 PLAY RECAP **** 08:35:11
9 zos_host : ok=1    changed=0    unreachable=0    failed=0    skipped=0    rescued=0    ignored=0
-----
```

Take time to explore the other templates that are ready to use. Learn more about the automation capabilities [here](#).

Importing skills from Ansible Automation Platform

Now that you understand Ansible Automation Platform (AAP) and the preinstalled automations available, you can import them as skills into your watsonx Orchestrate instance, which is used for assistant guided actions.

Import skills into your assistant

The next steps assume that you have an active browser window to the watsonx Orchestrate ITZ cloud account. If you do not, refer to the initial steps in [Creating an assistant and configuring conversational search](#).

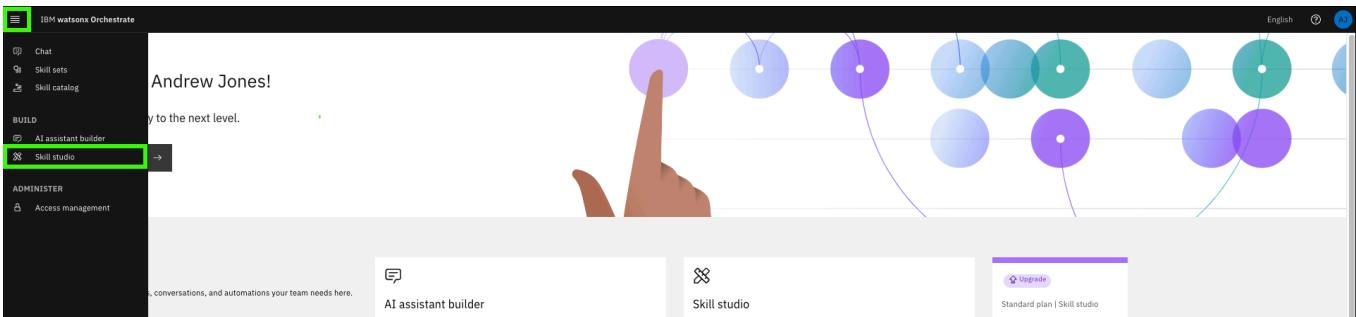
1. Click the **resources** icon (☰).

The screenshot shows the IBM Cloud dashboard with a green box highlighting the 'Build' card under the 'For you' section. The 'Build' card has a blue background and text: 'Explore IBM Cloud with this selection of easy starter tutorials and services.' Below it are other service cards: 'Use Watson Assistant', 'IBM Cloud Essential Security and Observability Services', 'Use Cloud Object Storage', 'Maximo Application Suite', 'Build with Watson', 'Use Watson Studio', and 'Get Started with Studio'. Each card includes a small icon, a title, a brief description, and a status badge (Popular, Recommended, 2 min, etc.).

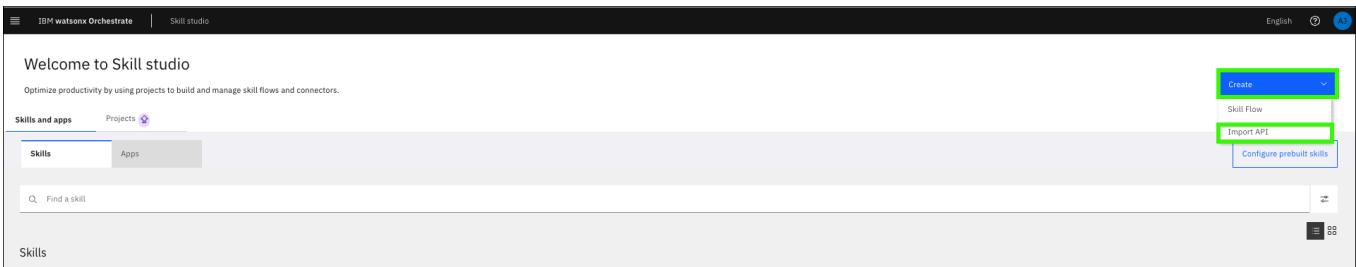
2. Click **Launch watsonx Orchestrate**.

The screenshot shows the 'Resource list' page for the watsonx instance 'itzwxo-2700039nft-wusdf'. The 'Manage' tab is selected. On the left, there's a sidebar with 'Service credentials'. The main area has a heading 'Start by launching the tool' with a 'Launch watsonx Orchestrate' button (which is highlighted with a green box) and a 'Getting started tutorial' link. Below this is a 'Credentials' section with fields for 'API key' (with a download and show credentials link) and 'URL' (with a copy link). To the right, there's a 'Plan' section titled 'Essentials Plan'.

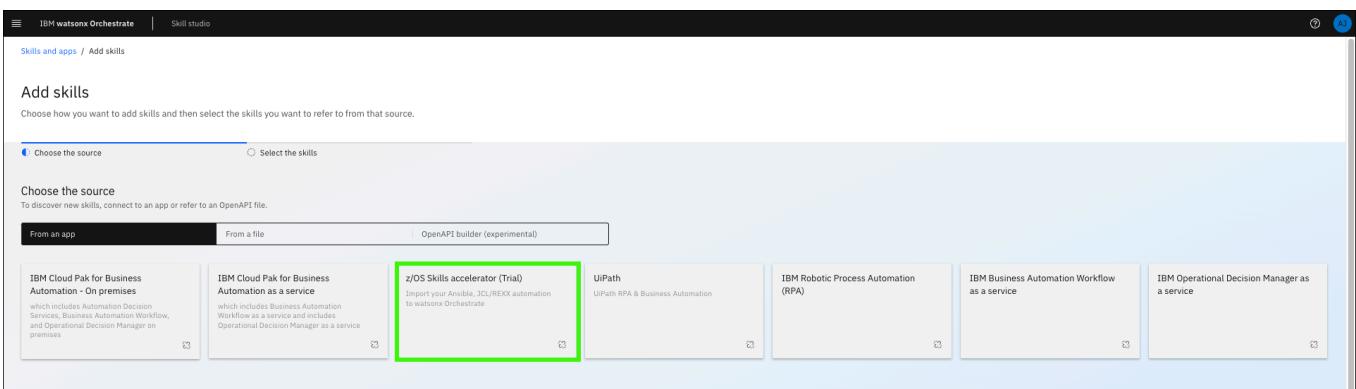
3. Click the main menu (☰) and then click **Skill studio**.



4. Expand Create and click Import API.



5. Click the z/OS Skills accelerator (Trial) tile.



6. Enter the following values in the z/OS Skills accelerator form and then click Connect.

Use the **URL**, **User Name**, and **Password** values recorded in the [Explore Ansible Automation Platform](exploreAAP.md) section earlier.

a: Connection Type: **ansible**

b: Application Name: <enter a meaningful name for the skills that you will import> - *be sure to remember this name, you will need in the next section*

c: Connection URL: <enter the URL for your AAP UI>

d: User Name: <enter the AAP User Name (for UI access)>

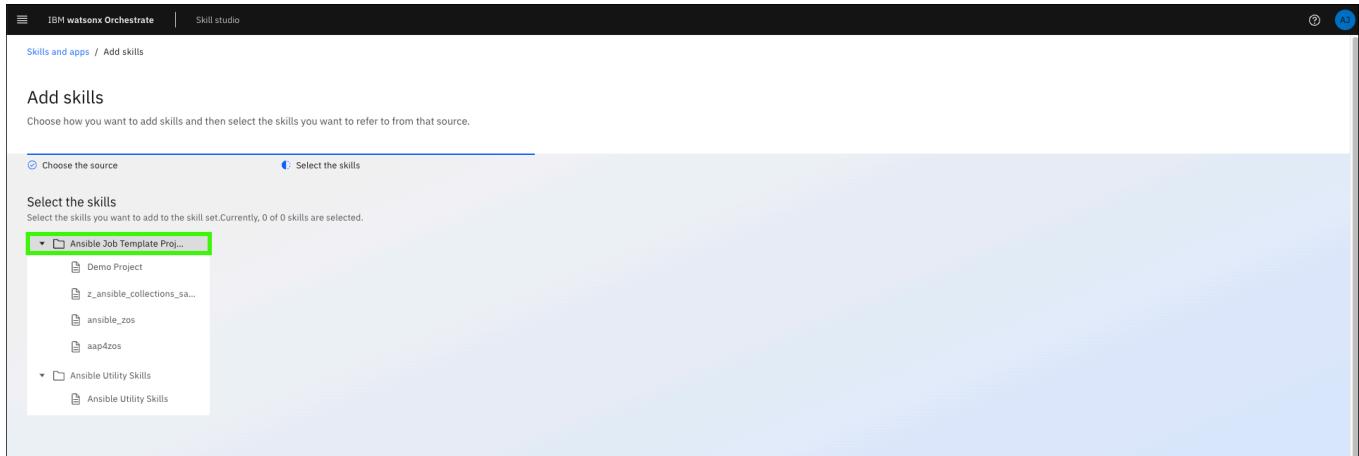
e: Password: <enter the AAP User Password>

f: Search Pattern: *****

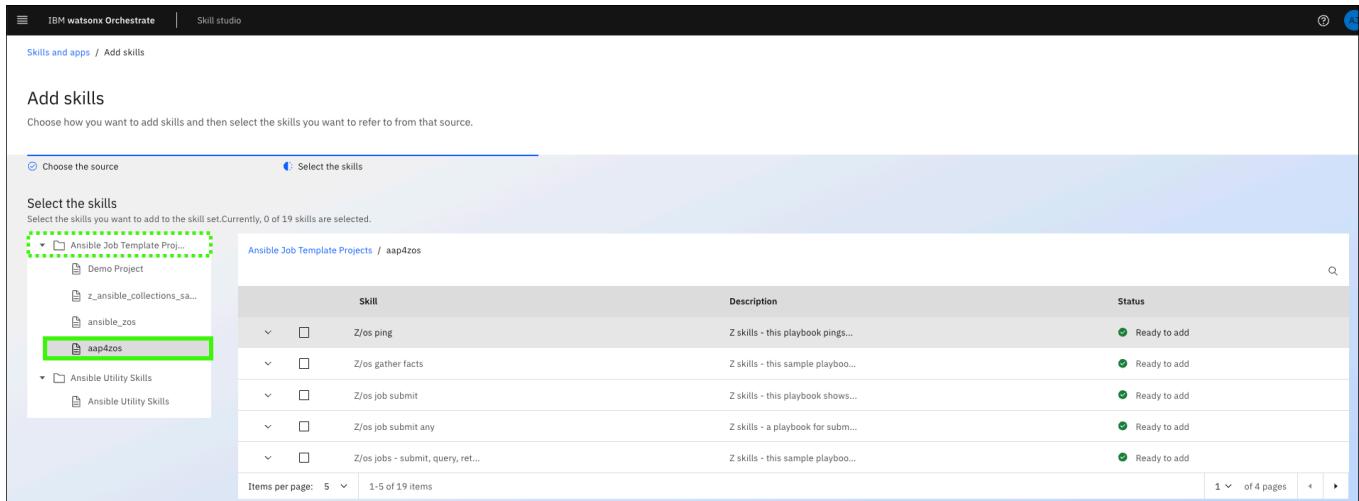
7. Expand the Ansible Job Template Proj... folder.

Browse the other available skills

There are also utility skills provided out of the box with the “Z Skills Accelerator” which are needed for actions such as retrieving the output of an Ansible skill (i.e. “Retrieve job output”). You should also import these utility skills to enable more complete automation execution flows.



8. Click **aap4zos**.



Skill	Description	Status
Z/os ping	Z skills - this playbook pings...	Ready to add
Z/os gather facts	Z skills - this sample playbook...	Ready to add
Z/os job submit	Z skills - this playbook shows...	Ready to add
Z/os job submit any	Z skills - a playbook for subm...	Ready to add
Z/os jobs - submit, query, ret...	Z skills - this sample playboo...	Ready to add

9. Select the skills you want to import into your application and then click **Save as draft**.

10. Click the ellipses (![]_attachments/ellipsesIcon.png) for the first skill and select **Enhance this skill.**

11. Click **Publish.**

On the **Enhance this skill** page, you can specify enhancements to the default skill. Refer to this documentation for more information on enhancing skills.

The screenshot shows the 'Skill studio' interface for enhancing a skill. The skill being enhanced is 'z/OS Gather Facts'. The 'Name' field contains 'z/OS Gather Facts'. The 'Description' field contains a sample playbook description. The 'API version*' field is set to '1.0'. Under 'Categories', there is an 'App' section with 'Ansible Controller Skills - z skills'. On the right, there are two preview sections: one for the catalog showing the skill card, and another for the skill set showing a single skill card. At the bottom right are buttons for 'Cancel', 'Publish' (which is highlighted with a green box), and 'Save as draft'.

12. Repeat steps 10 and 11 for each skill you imported.

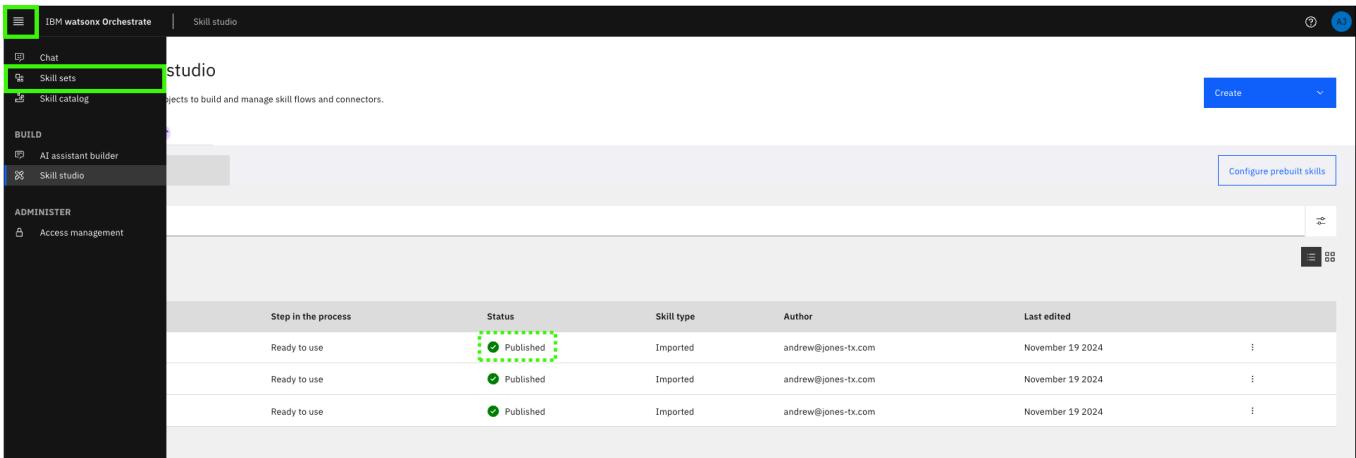
The screenshot shows the 'Skills' list in the 'Skill studio'. It lists two skills: 'z/OS Gather Facts' and 'z/OS Ping'. The 'z/OS Gather Facts' skill is marked as 'Published' (indicated by a green dashed box around the status cell) and was last edited on November 19 2024. The 'z/OS Ping' skill is marked as 'Ready to publish' and was last edited on November 19 2024. The 'Skills and apps' tab is selected. A 'Create' button is visible at the top right.

The selected skills are now ready for use and available to your assistant. In the next section, learn how to connect them to your assistant.

Connecting skills to your assistant

Once you have a subset of skills published, the application you created can be connected to your assistant.

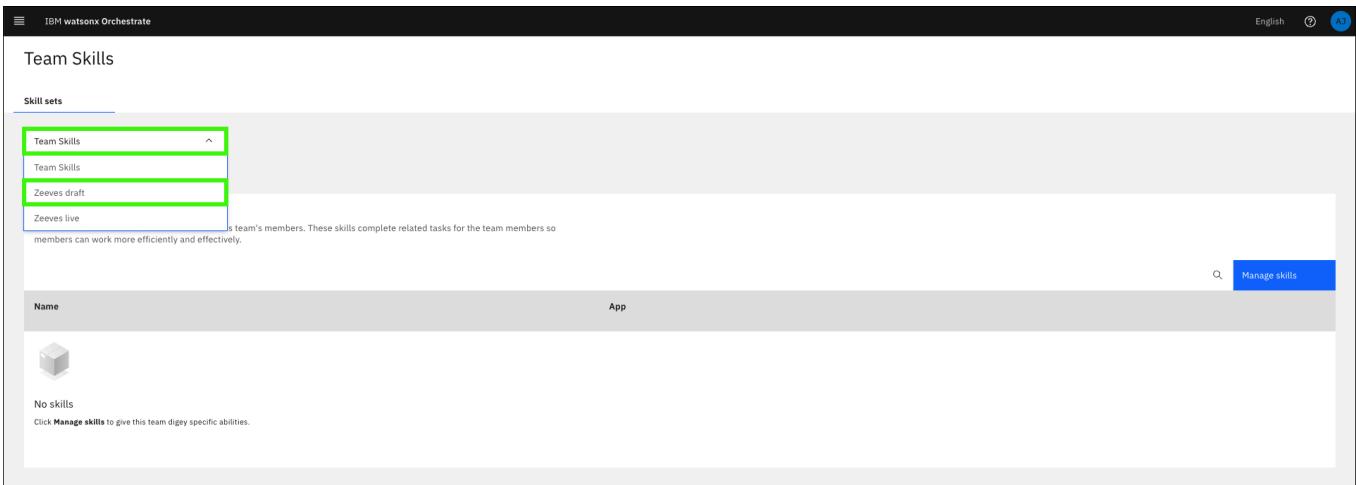
1. Expand the main menu and select **Skill sets**.



The screenshot shows the IBM Watsonx Orchestrate interface with the 'Skill studio' tab selected. The left sidebar has 'Skill sets' highlighted with a green box. The main area displays a table of published skills:

Step in the process	Status	Skill type	Author	Last edited
Ready to use	Published	Imported	andrew@jones-tx.com	November 19 2024
Ready to use	Published	Imported	andrew@jones-tx.com	November 19 2024
Ready to use	Published	Imported	andrew@jones-tx.com	November 19 2024

2. Click the **Team Skills** drop-down and select the **Draft** of your assistant.

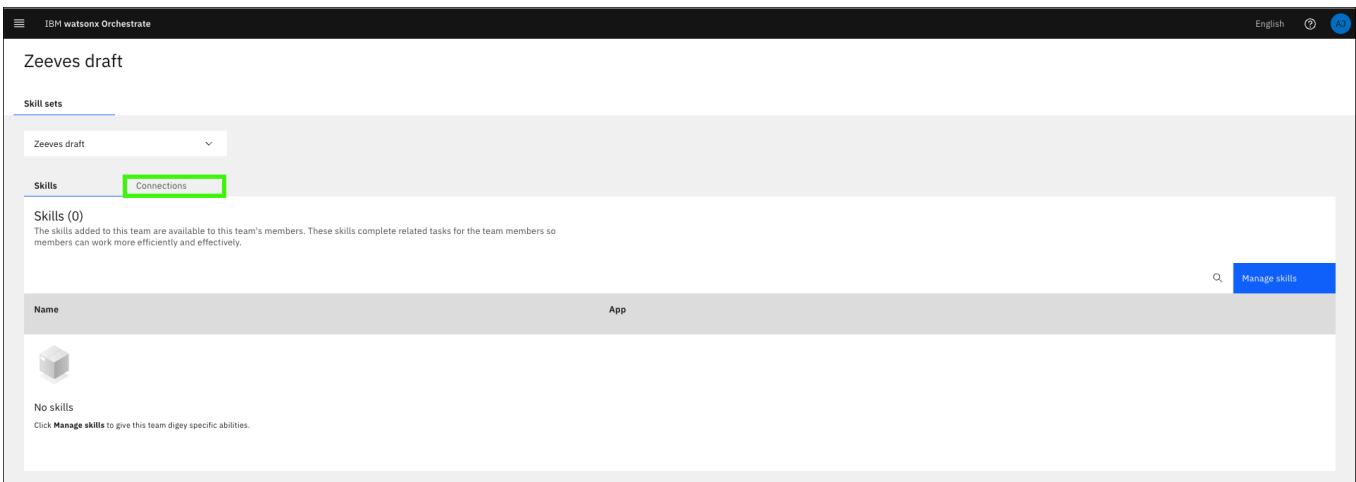


The screenshot shows the 'Team Skills' page. The 'Skill sets' dropdown is open, showing 'Team Skills' and 'Zeeves draft' (which is highlighted with a green box). Below the dropdown, a note says: 'The skills added to this team are available to this team's members. These skills complete related tasks for the team members so members can work more efficiently and effectively.' The main area shows a table with one row:

Name	App
No skills	

Below the table, there is a note: 'Click **Manage skills** to give this team digest specific abilities.'

3. Click the **Connections** tab.



The screenshot shows the 'Zeeves draft' team page. The 'Skills' tab is selected (highlighted with a green box), and the 'Connections' tab is also visible. A note above the table says: 'The skills added to this team are available to this team's members. These skills complete related tasks for the team members so members can work more efficiently and effectively.' The main area shows a table with one row:

Name	App
No skills	

Below the table, there is a note: 'Click **Manage skills** to give this team digest specific abilities.'

4. Click the **Search (🔍)** icon.

Application	Number of skills	Credential type	Connected by ⓘ	Action
Activate or deactivate attracting candidates using ThisWay Global	4	⚠️ Not specified	-	⋮
Adobe Workfront	37	⚠️ Not specified	-	⋮
Alliance Virtual Office	2	⚠️ Not specified	-	⋮
Amazon S3	8	⚠️ Not specified	-	⋮
Amazon SES	10	⚠️ Not specified	-	⋮

Items per page: 5 | 1-5 of 78 items | 1 ⚏ of 16 pages | < >

5. Search for the application name you specified in the previous section.

Application	Number of skills	Credential type	Connected by ⓘ	Action
Ansible Controller Skills - z skills	2	⚠️ Not specified	-	⋮

Items per page: 5 | 1-1 of 1 items | 1 ⚏ of 1 page | < >

6. Click the ellipses (⋮) and click **Connect app**.

Application	Number of skills	Credential type	Connected by ⓘ	Action
Ansible Controller Skills - z skills	2	⚠️ Not specified	-	⋮

Items per page: 5 | 1-1 of 1 items | Connect app | 1 ⚏ of 1 page | < >

7. On the **Connect to Ansible Controller Skills** form, keep the defaults and click **Connect app**.

The screenshot shows the IBM Watsonx Orchestrate interface. On the left, there's a sidebar with 'Skill sets' and a dropdown menu showing 'Zeeves draft'. Below it are tabs for 'Skills' and 'Connections'. A search bar 'Q z skills' is present. In the main area, there's a table with columns 'Application', 'Number of skills', and 'Credential type'. One row shows 'Ansible Controller Skills - z skills' with '2' skills and 'Not specified' as the credential type. At the bottom, there are pagination controls 'Items per page: 5' and '1-1 of 1 items'. On the right, a modal window titled 'Connect to Ansible Controller Skills - z skills' is displayed. It has two radio button options: 'Member credentials' (unselected) and 'Team credentials' (selected). Below the radio buttons, a note says 'You selected Team credentials for the credential type. Click Connect app to provide the credentials your team will use to connect to the app.' A large green box highlights the 'Connect app' button at the bottom right of the modal.

8. Enter the (a) **username** and (b) **password** using the username (**admin**) and password for your IBM Technology Zone (ITZ) watsonx Assistant for Z Pilot - AAP & z/OS reservation (**AAP User Password (Use SSH key to login, only use password for UI)**), and then click **Connect app**.

This screenshot shows the same modal window from the previous step. The 'username' field (labeled 'a') contains 'admin' and the 'password' field (labeled 'b') contains '.....'. A note below the fields says 'If the service instance uses legacy credentials for authentication, provide the password for the specified username.' A large green box highlights the 'Connect app' button at the bottom right, which is labeled 'c' with a green arrow pointing to it.

The application should now be connected to the draft version of your assistant.

This screenshot shows the 'Zeeves draft' skill set in the IBM Watsonx Orchestrate interface. The 'Connected by' column for the 'Ansible Controller Skills - z skills' row displays the email 'andrew@jones-tx.com'. A green dotted box highlights the entire row. The rest of the interface is similar to the previous screenshots, showing the skill set details and connection status.

Continue to the next section to create actions for your assistant.

Creating actions for your assistant

Once the skills in your application are connected to your assistant, you're ready to begin creating actions tied to those skills. Learn more about building actions [here](#)

Configure the number of input fields

Before configuring actions, it's important to modify a setting within watsonx Orchestrate that allows triggered skills to display as forms (versus conversational skills).

1. Click your (a) profile icon and then click (b) **Settings**

Learn more about configuring input fields [here](#).

Zeeves draft

Skill sets

Zeeves draft

Skills Connections

These are the applications that are used by the skills in team skill set. Application connections are required to execute skills. Set preference at an application level to enable skills to either use personal or team credentials.

Application	Number of skills	Credential type	Connected by
Ansible Controller Skills - z skills	2	Team	andrew@jones-tx.com

Items per page: 5 1-1 of 1 items

English AJ

Region US-SOUTH Plan Essentials Settings Provide feedback Privacy About Log out

2. Click the **Skill configurations** tab.

Settings

Data Retention Skill configurations Platform languages

Data Retention Policy

Active

The chat history for the users in the tenant are saved for 30 days. After 30 days, the chat history is deleted permanently.
Deleting the chat history doesn't affect the connected apps or the added skills.

3. Enter **0** for the **Number of form fields**.

Settings

Data Retention Skill configurations Platform languages

Skill configurations

Set default for maximum number of fields to consider a skill to be conversational

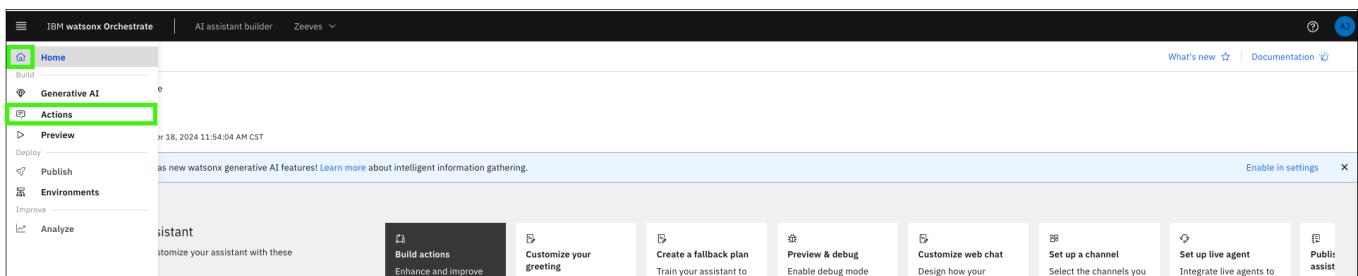
Number of form fields - +

Create actions

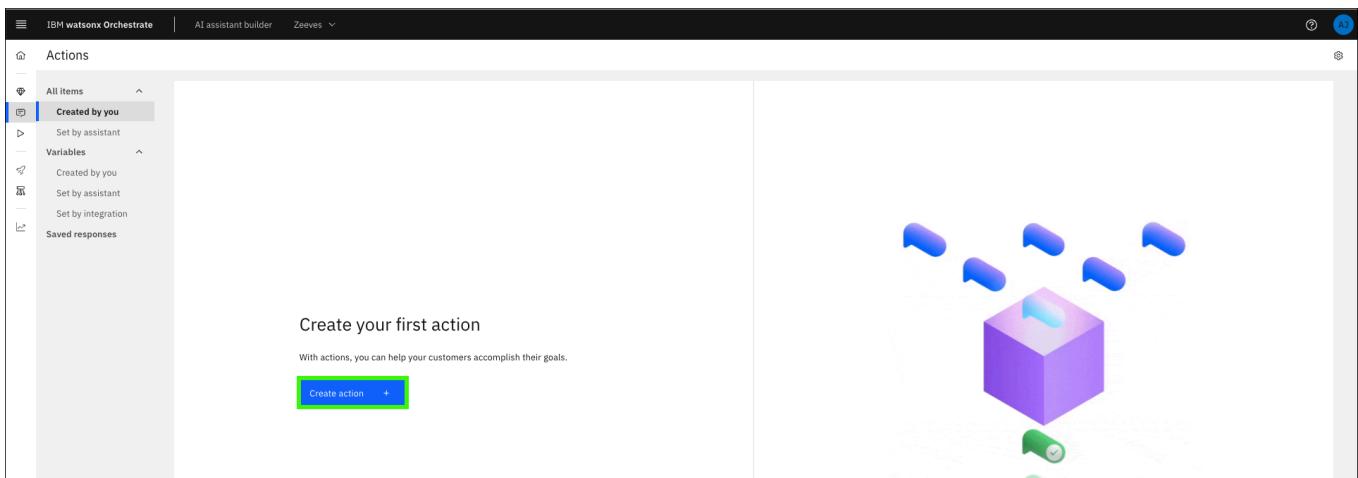
4. Click the main menu and select **AI assistant builder**.



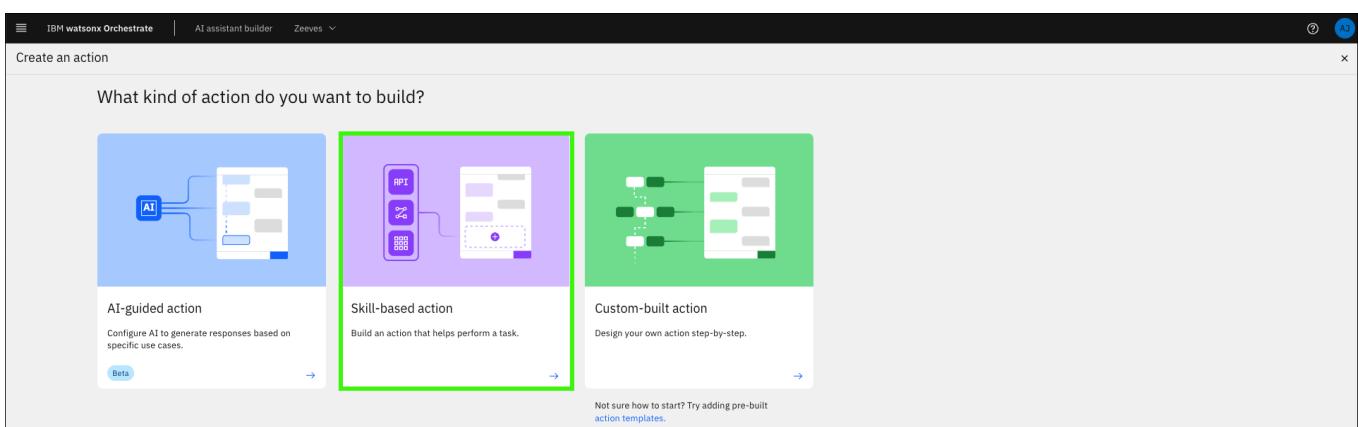
5. Hover over the **Home** icon (🏠) and click **Actions**.



6. Click **Create action**.



7. Click the **Skill-based action** tile.



8. Select the **z/OS Gather Facts** tile and click **Next**.

Note, the date shown in the tile should reflect when you add the skill to your application.

The screenshot shows a list of skills under the heading "Build an action from a skill". The "z/OS Gather Facts" skill is selected and highlighted with a green border. The "Next" button at the top right is also highlighted with a green border.

Build an action from a skill				
Select a skill Choose a conversational skill published as a foundation of your action.				
<input type="text" value="Search a skill"/> Cancel Next				
z/OS Ping <small>z skills - This playbook pings the z/OS host to test connectivity.</small> <small>Last updated: 2024-11-19T15:58:20.567Z</small>	z/OS Gather Facts <small>z skills - This sample playbook demonstrates the z/OS gather facts module, which pulls z/OS-specific information from the z/OS host.</small> <small>Last updated: 2024-11-19T15:56:26.843Z</small>			
Summarize a Servicenow incident <small>in watsonx.ai</small> <small>Last updated: 2024-11-04T10:49:05.828Z</small>	Summarize a Salesforce opportunity <small>in watsonx.ai</small> <small>Last updated: 2024-11-04T10:49:01.769Z</small>	Sharepoint document summary <small>in watsonx.ai</small> <small>Last updated: 2024-11-04T10:48:55.707Z</small>	Salessoft email summary <small>in watsonx.ai</small> <small>Last updated: 2024-11-04T10:48:51.522Z</small>	Summarize a Zendesk ticket <small>in watsonx.ai</small> <small>Last updated: 2024-11-04T10:49:09.476Z</small>
Salesforce case sentiment analyze	Outlook email summary	Github issue summarization	Github issue sentiment	Generate an email

9. On the **New action** dialog, (a) enter a prompt a user of the assistant might use to initiate the action and then (b) click **Save**.

Sample prompts:

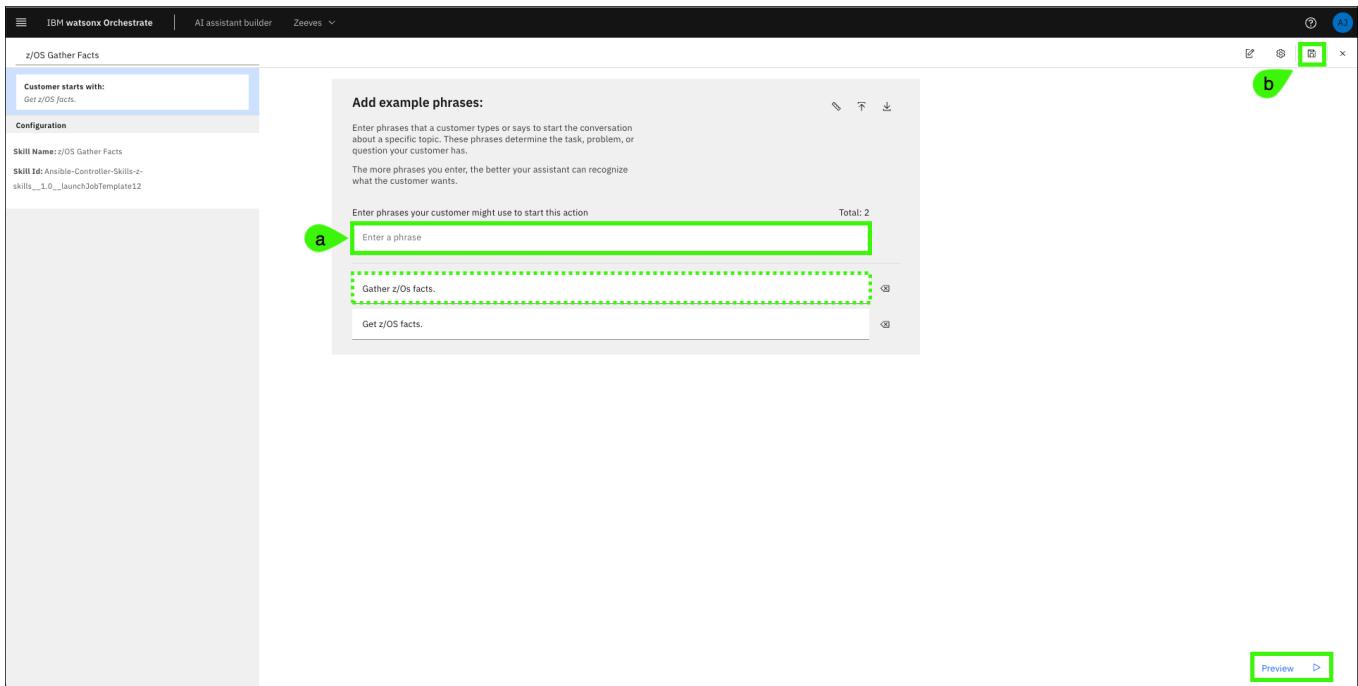
The screenshot shows the "Add example phrases" dialog. A prompt "Get z/OS facts" is entered in the input field. The "Save" button is highlighted with a green border and labeled 'b'.

Add example phrases:
Enter phrases that a customer types or says to start the conversation about a specific topic. These phrases determine the task, problem, or question your customer has.
The more phrases you enter, the better your assistant can recognize what the customer wants.

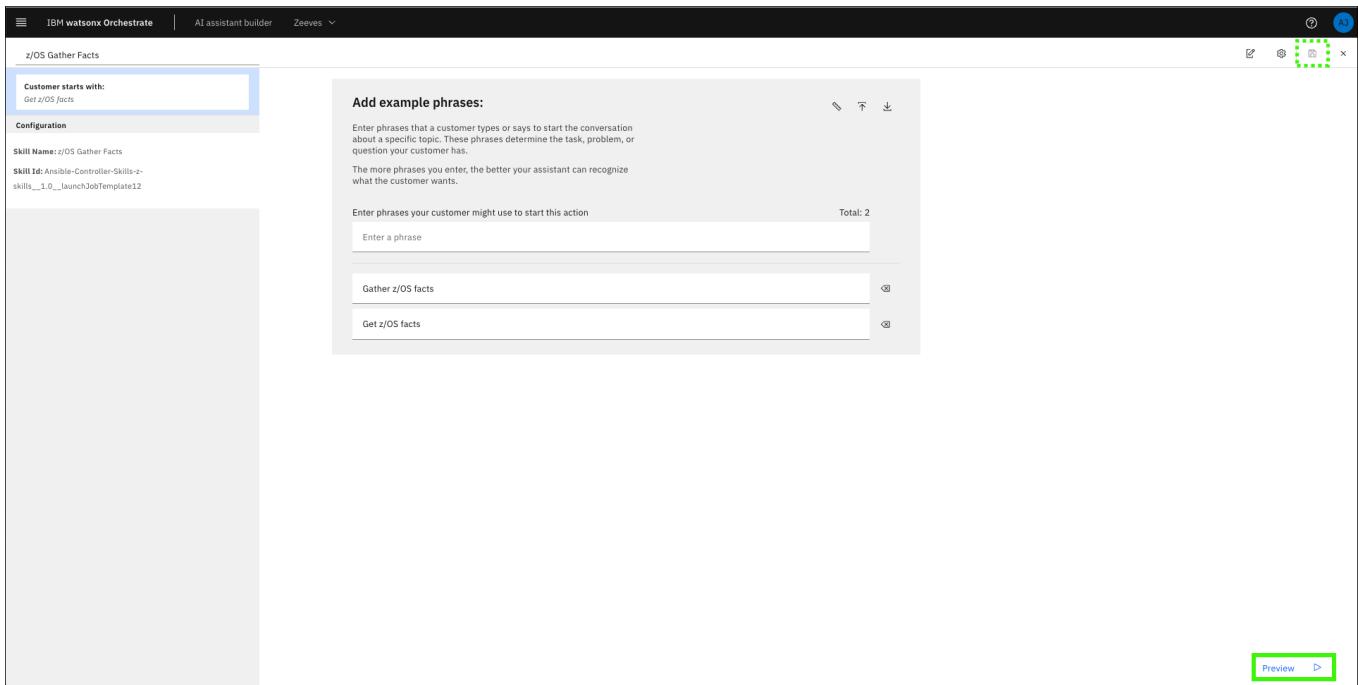
Enter phrases your customer might use to start this action
Total: 0

Example: I want to pay my credit card bill.

10. Add any (a) additional prompts and then (b) click the save ().



11. Click Preview.



12. Enter one of the prompts you specified in step 9 or 10.

Prompt:

Get z/OS facts

Customer starts with:
Get z/OS facts

Add example phrases:

Enter phrases that a customer types or says to start the conversation about a specific topic. These phrases determine the task, problem, or question your customer has.

The more phrases you enter, the better your assistant can recognize what the customer wants.

Enter phrases your customer might use to start this action
Total: 2

Enter a phrase

Gather z/OS facts
Get z/OS facts

Preview

12:55 PM
Greet customer [default]

Welcome, how can I assist you?

Get z/OS facts

13. Review the returned results.

Customer starts with:
Get z/OS facts

Add example phrases:

Enter phrases that a customer types or says to start the conversation about a specific topic. These phrases determine the task, problem, or question your customer has.

The more phrases you enter, the better your assistant can recognize what the customer wants.

Enter phrases your customer might use to start this action
Total: 2

Enter a phrase

Gather z/OS facts
Get z/OS facts

Preview

12:55 PM
Greet customer [default]

Welcome, how can I assist you?

You 1:04 PM
Get z/OS facts

1:05 PM
Conversational skill called z/OS Gather Facts recognized
job : 12
status : pending
There are no additional steps for this action. Add a new step or end the action.

Use the up arrow for prior messages

In the execution of this skill-based action, the skill executed properly and the output is the job id. Learn more about building actions [here](#).

Verify the job in the Ansible Automation Platform console

Return to the Ansible Automation Platform (AAP) console and review the job information.

14. Click **Jobs** and expand the ## - z/OS Gather Facts job.

The screenshot shows the Red Hat Ansible Automation Platform web interface. The left sidebar has a 'Views' dropdown with 'Dashboard' and 'Jobs' selected. Below it are sections for 'Resources' like 'Templates', 'Credentials', 'Projects', and 'Inventories'. The main area is titled 'Jobs' and lists 12 items. One item, '12 - z/OS Gather Facts', is highlighted with a green box. This item's details are shown in a modal below:

Launched By	admin	Job Template	z/OS Gather Facts	Inventory	AAP z/OS
Project	aap4zos	Execution Environment		aap4zos EE	
Credentials	SSH: z/OS Host SSH...				
Job Slice	0/1				

As seen in the assistant, the actual contents of the output aren't displayed. The utility skills are used to retrieve the job output. It is also possible to create a skill flow that executes the **z/OS Gather Facts** skill followed by the **Retrieve job output** utility skill in sequence; passing the job id from the first skill to the second, in order to view the output within the assistant. Creating a skill flow is covered in the next section.

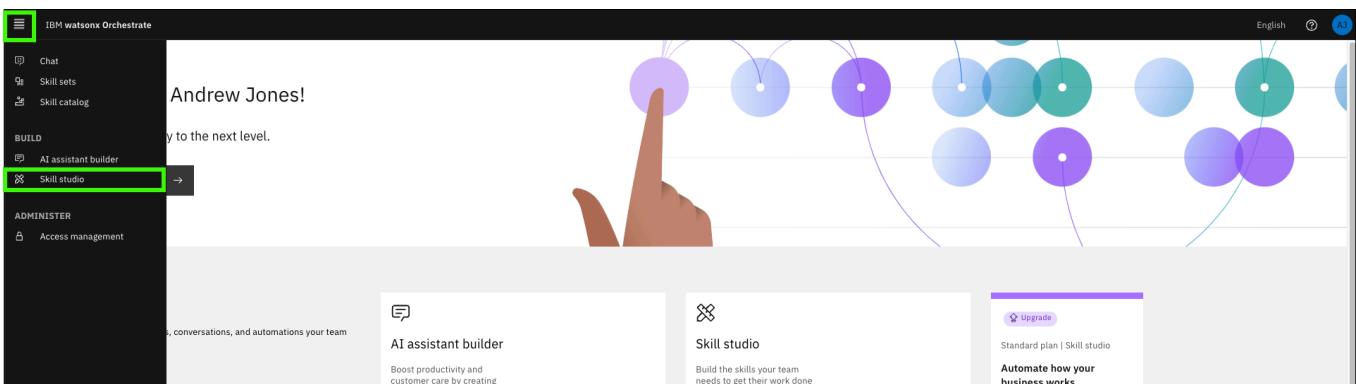
Creating skill flows

As seen in the previous section, running the Ansible skill to **Gather z/OS facts**, the skill executed successfully and was verified within the Ansible Automation Platform (AAP) console by viewing the job output. However, the output wasn't displayed by the assistant. To enable this scenario, a skill flow is needed. Skills are often more valuable when combined with other skills. You can create a skill flow to use two or more skills together to finish a task (like returning the output of a previous skill). When you create a skill flow, you map the output of one skill as the input for subsequent skills. Learn more about creating skill flows [here](#).

As mentioned in a previous section, there are some default utility skills that are provided out of the box with the "Z Skills Accelerator" which are leveraged to return the output of a skill. To accomplish this, we will import the Ansible Utility skill called **Retrieve job output**.

Add the utility skill

1. Open IBM watsonx Orchestrate **Skill studio**.



2. Expand **Create** and click **Import API**.

Name	Step in the process	Status	Skill type	Author	Last edited
z/OS Ping	Ready to use	Published	Imported	andrew@jones-tx.com	November 19 2024
z/OS Gather Facts	Ready to use	Published	Imported	andrew@jones-tx.com	November 19 2024

3. Click the **z/OS Skills accelerator (Trial)** tile.

The screenshot shows the 'Add skills' interface in IBM Watsonx Orchestrate. At the top, there are two tabs: 'Choose the source' (selected) and 'Select the skills'. Below this, a section titled 'Choose the source' with the sub-instruction 'To discover new skills, connect to an app or refer to an OpenAPI file.' contains three buttons: 'From an app' (selected), 'From a file', and 'OpenAPI builder (experimental)'. A list of skill sources is shown below, with the 'z/OS Skills accelerator (Trial)' entry highlighted by a green box. Other entries include 'IBM Cloud Pak for Business Automation - On premises', 'IBM Cloud Pak for Business Automation as a service', 'UiPath', 'IBM Robotic Process Automation (RPA)', and 'IBM Business Automation Workflow as a service'. Each entry has a small icon and a 'View details' button.

4. Enter the following values in the **z/OS Skills accelerator** form and then click **Connect**.

Use the **URL**, **User Name**, and **Password** values recorded in the **Explore Ansible Automation Platform** section earlier.

a: Connection Type: **ansible**

b: Application Name: <use the same application name as in previous section> - *

c: Connection URL: <enter the URL for your AAP UI>

d: User Name: <enter the AAP User Name (for UI access)>

e: Password: <enter the AAP User Password>

f: Search Pattern: *

This screenshot shows the configuration dialog for connecting to the 'z/OS Skills accelerator (Trial)'. The right side of the dialog is a sidebar with instructions and field descriptions. The main area contains several input fields, each labeled with a letter from a to f. The 'Connection Type' field (a) contains 'ansible'. The 'Application Name' field (b) contains 'z skills'. The 'Connection URL' field (c) contains 'https://itzvi-aap-phdhuzza.techzone.ibm.com'. The 'User name' field (d) contains 'admin'. The 'Password' field (e) contains a masked password. The 'Search Pattern' field (f) contains a single quote character. At the bottom right are 'Cancel' and 'Connect' buttons, with 'Connect' being highlighted by a green box.

5. Expand **Ansible Utility Skills** and click **Ansible Utility Skills**.

6. Select **Retrieve job output** and click **Save as draft**.

7. Click the ellipses (...) for the **Retrieve job output** skill and select **Enhance this skill**.

8. Review the skill settings and then click **Publish**.

Name: Retrieve job output

Description: z skills - Retrieve job output by job Id

API version*: 1.0

Categories: Add categories

App: Ansible Controller Skills - z skills

Preview

The skill will look like this in the catalog.

The skill will look like this in the skill set.

Publish (highlighted)

9. Select Skill sets from the main menu.

Step in the process	Status	Skill type	Author	Last edited
Ready to use	Published	Imported	andrew@jones-tx.com	November 19 2024
Ready to use	Published	Imported	andrew@jones-tx.com	November 19 2024
Ready to use	Published	Imported	andrew@jones-tx.com	November 19 2024

10. Select (a) your draft assistant in the Team Skills drop-down list and (b) click the Connections tab.

Skills (highlighted)

Connections (highlighted)

Application	Number of skills	Credential type	Connected by	Action
Activate or deactivate attracting candidates using ThisWay Global	4	Not specified	-	⋮
Adobe Workfront	37	Not specified	-	⋮
Alliance Virtual Office	2	Not specified	-	⋮
Amazon S3	8	Not specified	-	⋮
Amazon SES	10	Not specified	-	⋮

11. Click the **Search (🔍)** icon.

Application	Number of skills	Credential type	Connected by ⓘ	Action
Activate or deactivate attracting candidates using ThisWay Global	4	⚠️ Not specified	-	⋮
Adobe Workfront	37	⚠️ Not specified	-	⋮
Alliance Virtual Office	2	⚠️ Not specified	-	⋮
Amazon S3	8	⚠️ Not specified	-	⋮
Amazon SES	10	⚠️ Not specified	-	⋮

Items per page: 5 | 1-5 of 78 items | 1 ⚏ of 16 pages | ⌂ ⌃

12. Search for the application name you specified earlier.

Application	Number of skills	Credential type	Connected by ⓘ	Action
Ansible Controller Skills - z skills	2	⚠️ Not specified	-	⋮

Items per page: 5 | 1-1 of 1 items | 1 ⚏ of 1 page | ⌂ ⌃

13. Click the (a) ellipses (⋮) for your application and (b) click **Edit connection**.

Application	Number of skills	Credential type	Connected by ⓘ	Action
Ansible Controller Skills - z skills	3	⚠️ Team	andrew@jones-tx.com	⋮

a
b
Edit connection
Delete connection

Items per page: 5 | 1-1 of 1 items | 1 ⚏ of 1 page | ⌂ ⌃

14. Verify the application is (a) **Connected** and then (b) click the x to close the dialog.

Zeeves draft

Skill sets

Skills Connections

These are the applications that are used by the skills in team skill set. Application connections are required to execute skills. Set preference at an application level to enable skills to either use personal or team credentials.

Application	Number of skills	Credential type
Ansible Controller Skills - z skills	3	Team

Items per page: 5 | 1-1 of 1 items

Edit the Ansible Controller Skills - z skills connection **b**

Member credentials
Each team member uses their own credentials to connect to this app and use its skills.

Team credentials (Active)
The admin sets the credentials each team member uses to connect to this app and use its skills.

You have an active connection set using **Team credentials**. If you wish to update connection details click on the **Edit** **a** below.

Connected **a**

Add the skills to your Personal skills

15. Click **Skill catalog** in the main menu.

Chat

Skill sets

Skill catalog **a**

AI assistant builder

Skill studio

ADMINISTER

Access management

Number of skills	Credential type	Connected by a	Action
3	Team	andrew@jones-tx.com	b

1 of 1 items

16. Search for the application name you specified earlier.

Skill catalog

Skills are grouped by app. Select an app to see all the skills that use that app.

Personal skills

z skills

Most popular skills

Send an email from Gmail	Create a lead in Salesforce	Send an email using Outlook
--------------------------	-----------------------------	-----------------------------

All Apps

Coupa	ZoomInfo	HubSpot CRM	Aptio Targetprocess	Salesforce	Zendesk Service
Calendly	Square	Oracle E-Business Suite	GitLab	Toggl Track	Microsoft Teams

17. Click the tile for your application.

Note, the tile name is proceeded by **Ansible Controller Skills**.

Skill catalog

Skills are grouped by app. Select an app to see all the skills that use that app.

Personal skills

z skills

Apps

- ZoomInfo (26 skills)
- Workday HCM (36 skills)
- FreshService (25 skills)
- Seismic (22 skills)
- Reveal your existing applic... (7 skills)
- Interview top candidates u... (2 skills)
- Salesforce Chatter (5 skills)
- IBM Process Mining (2 skills)
- Webex (14 skills)
- Adobe Workfront (37 skills)
- Skill flows (167 skills)
- Cognos (8 skills)
- Salessoft (47 skills)
- Ansible Controller Skills - z ... (3 skills)**

18. Click **Add skill** for each of the skills you want to add to the flow.

Skill catalog / Ansible Controller Skills - z skills (3)

Connect app

Personal skills

Search skills

Ansible Controller Skills - z skills

- Retrieve job output (z skills - Retrieve job output by job Id) **Add skill +**
- z/OS Gather Facts (z skills - This sample playbook demonstrates the z/OS gather facts module, which pulls...) **Add skill +**
- z/OS Ping (z skills - This playbook pings the z/OS host to test connectivity.) **Add skill +**

Create the skill flow

19. Click **Skill studio** in the main menu.

IBM Watsonx Orchestrate

English

Chat

Skill sets

Skill catalog

BUILD

- AI assistant builder
- Skill studio**

ADMINISTER

- Access management

Skills - z skills (3)

z skills

- z/OS Gather Facts (z skills - This sample playbook demonstrates the z/OS gather facts module, which pulls...) **Added ✓**
- z/OS Ping (z skills - This playbook pings the z/OS host to test connectivity.) **Added ✓**

Connect app

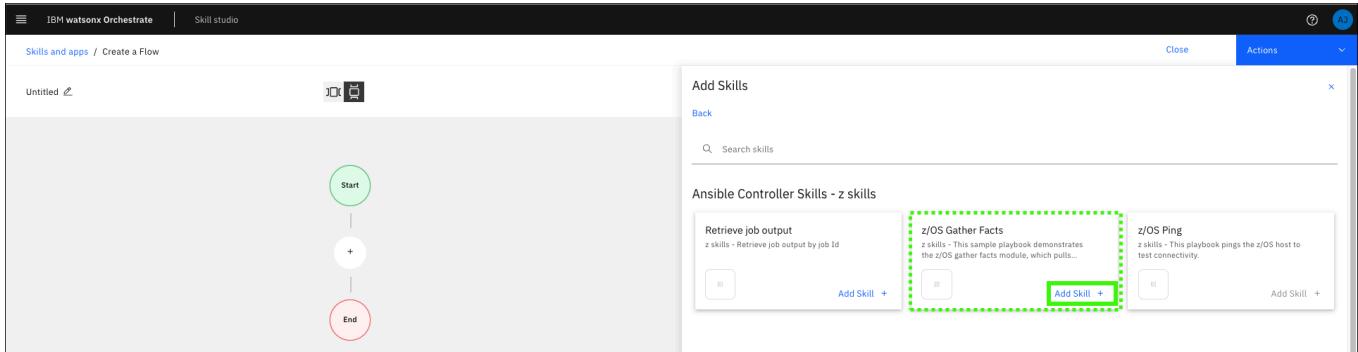
20. Expand the **Create** drop-down menu and click on **Skill flow**.

21. Click the + icon.

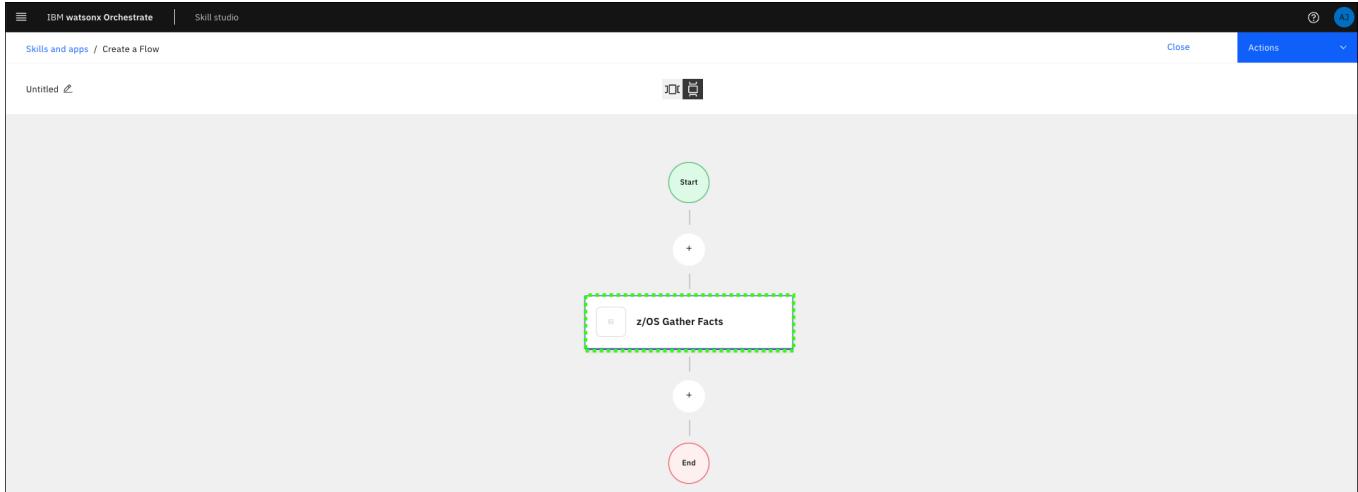
Next, you need to add the **z/OS Gather Facts** skill and the **Retrieve job output** skill to the skill flow. Use the **Search apps** function to locate the skills.

22. Search for the application name you specified earlier and click its tile.

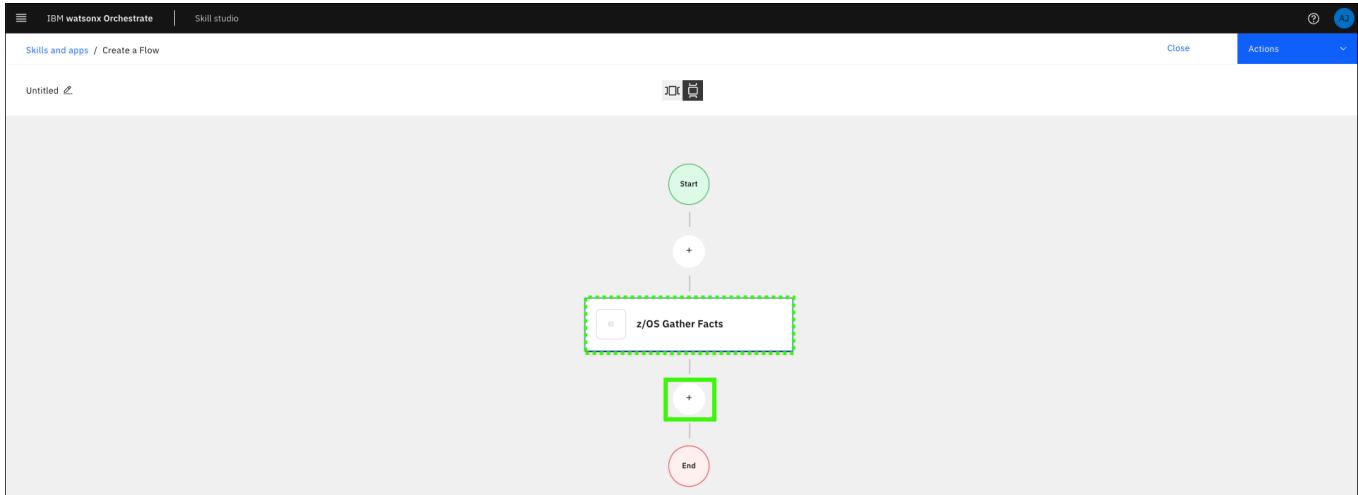
23. Click Add Skill in the **z/OS Gather Facts** tile.



24. Verify the **z/OS Gather Facts** skill is added to the skill flow.

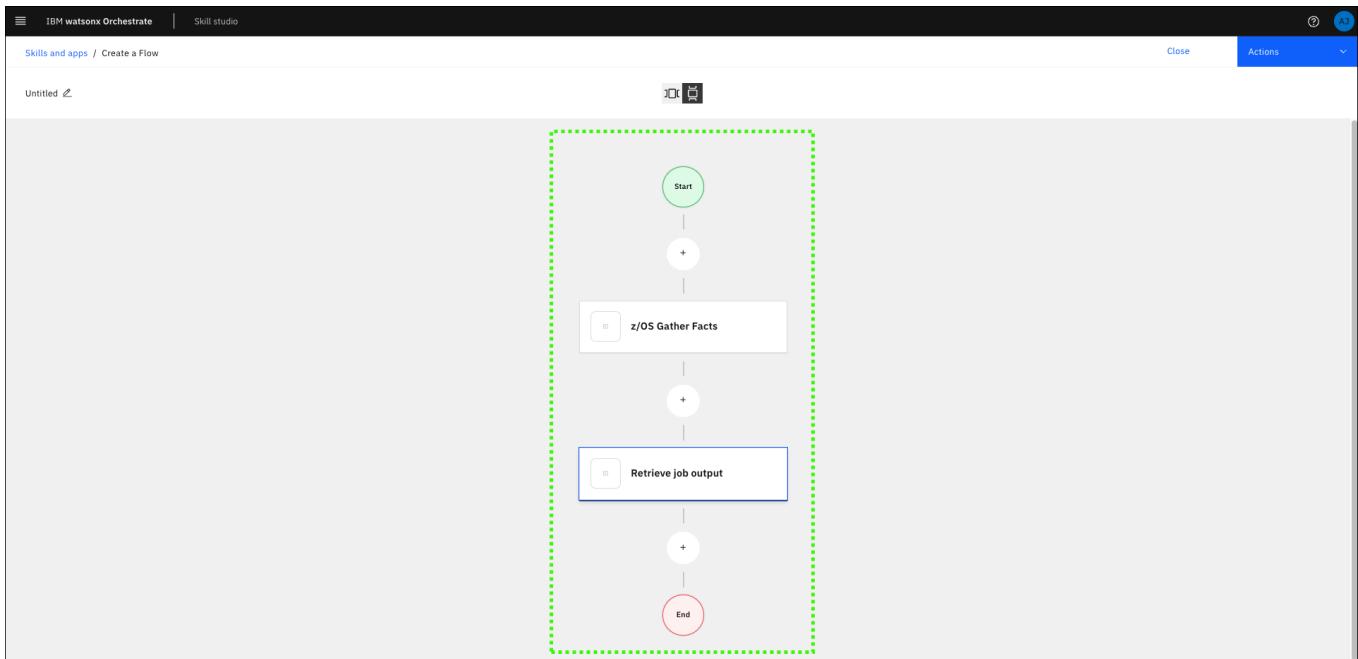


25. Click the **+** icon after the **z/OS Gather Facts** tile.



26. Repeat steps 22 and 23 for the **Retrieve job output** skill.

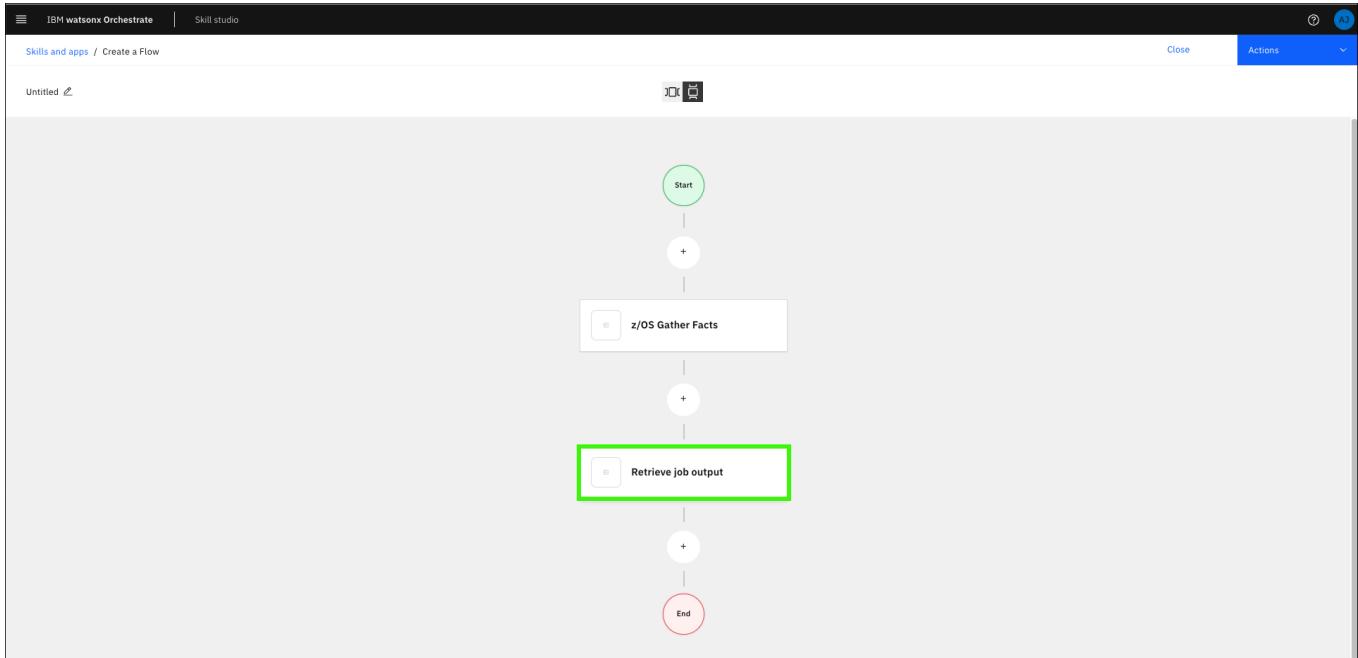
After adding the **Retrieve job output** skill, your skill flow should look like:



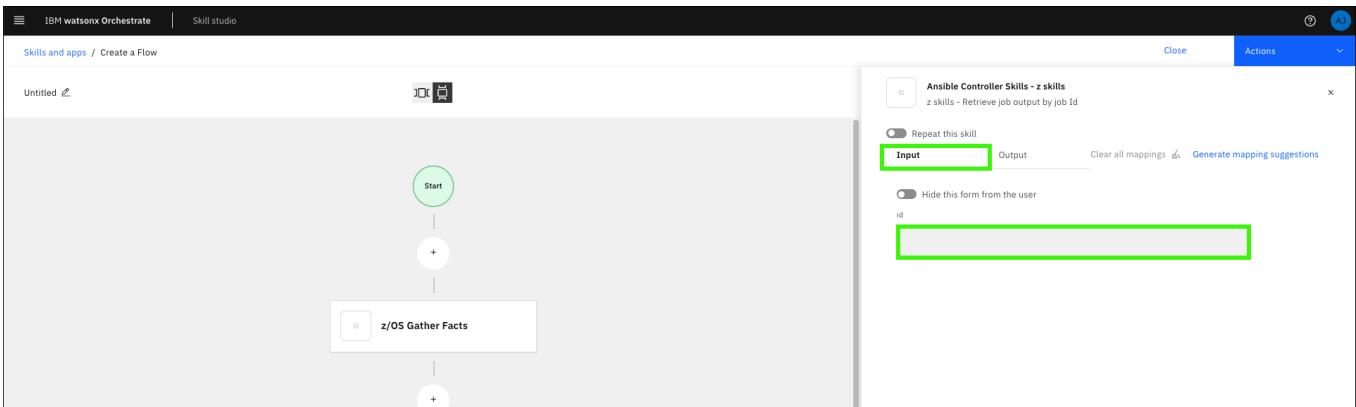
Map the outputs and inputs of the skills

Next you must map the output values of the first skill to the input of the second skill. In this case, pass the "job id" output from **z/OS Gather Facts** as an input for **Retrieve job output**.

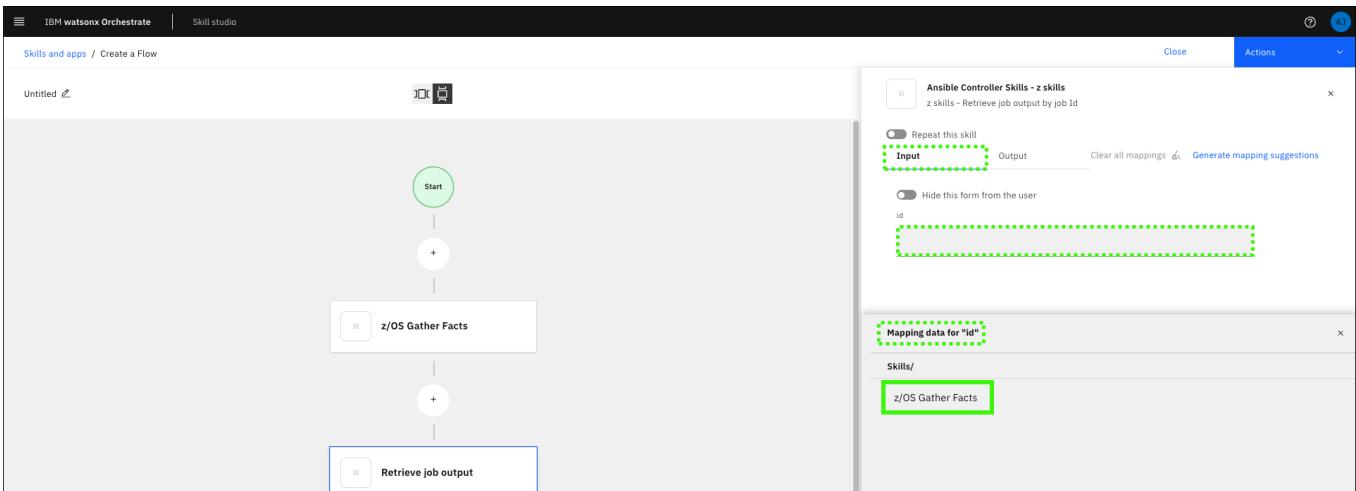
27. Click the **Retrieve job output** tile.



28. Select the **Input** tab and click in the **id** field.

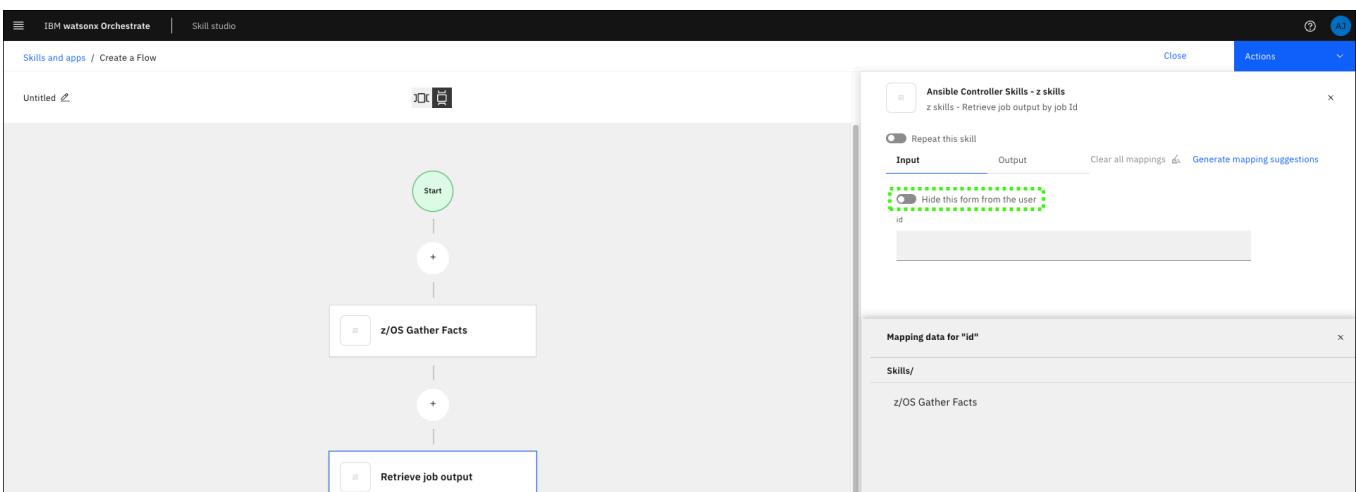


29. Click the **z/OS Gather Facts** skill in the **Mapping data for "id"** section.

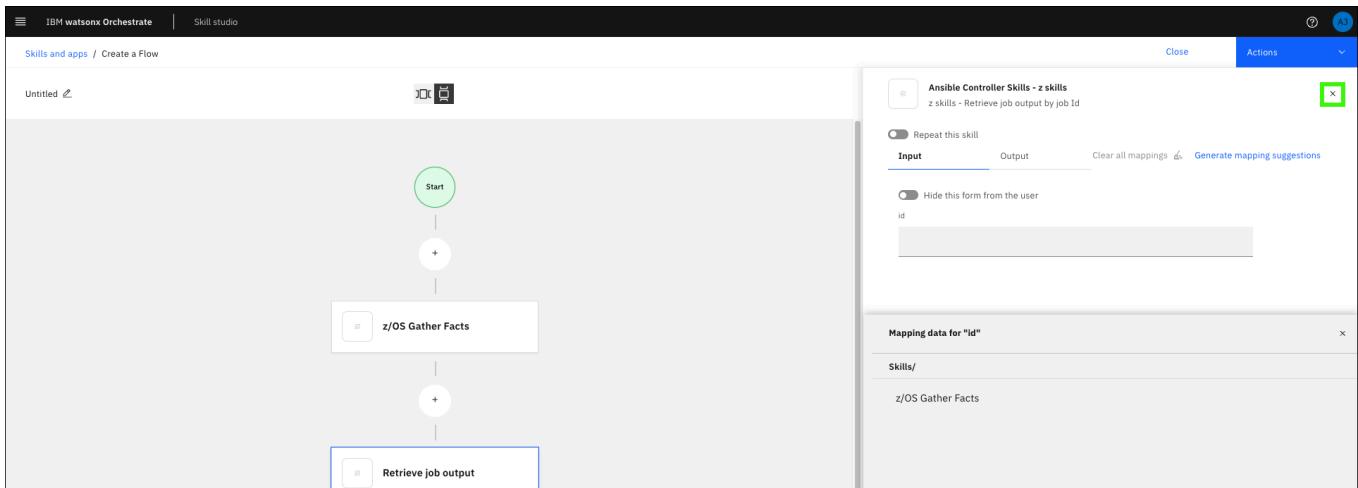


30. Optionally, toggle the **Hide this from from the user** setting.

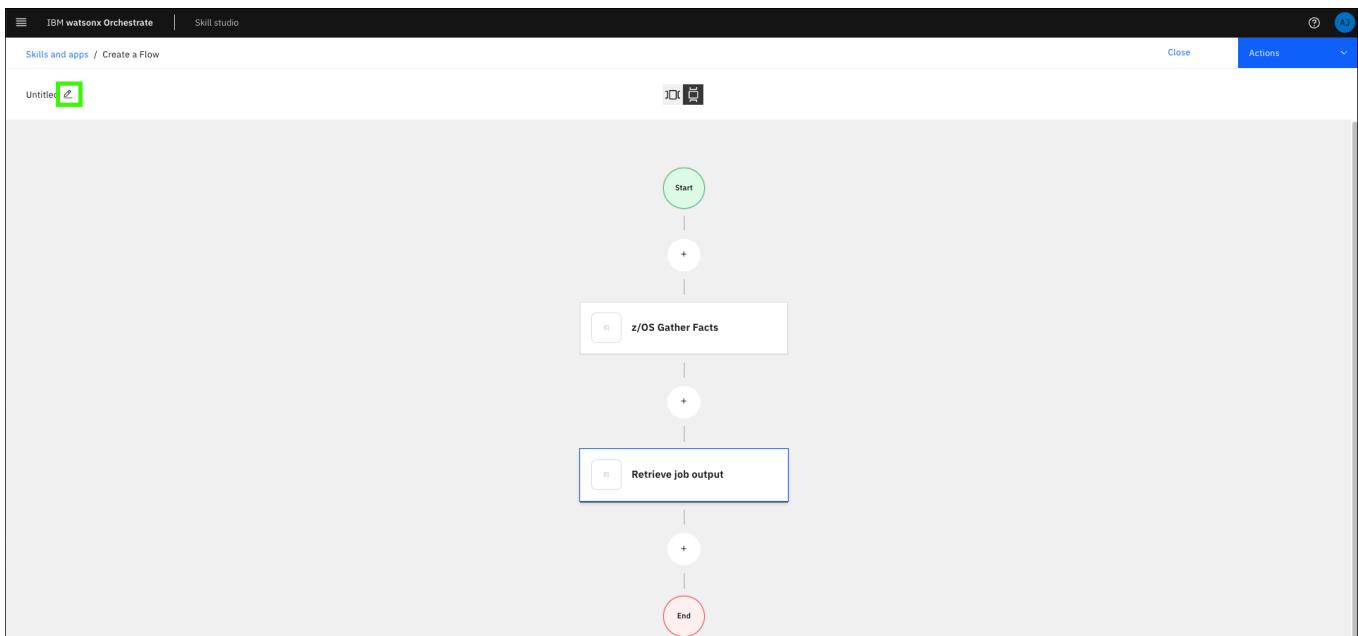
For this lab guide, this option is left disabled. Learn more about this option [here](#).



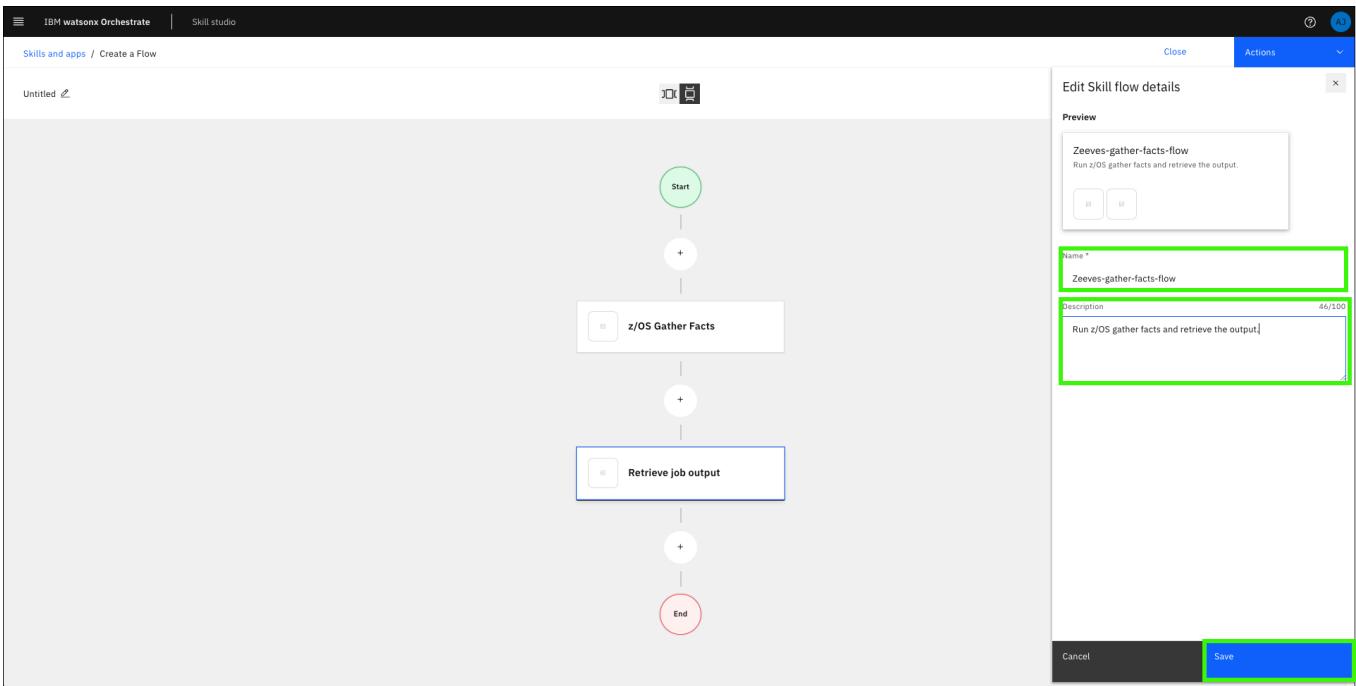
31. Click the **x** to close mapping window.



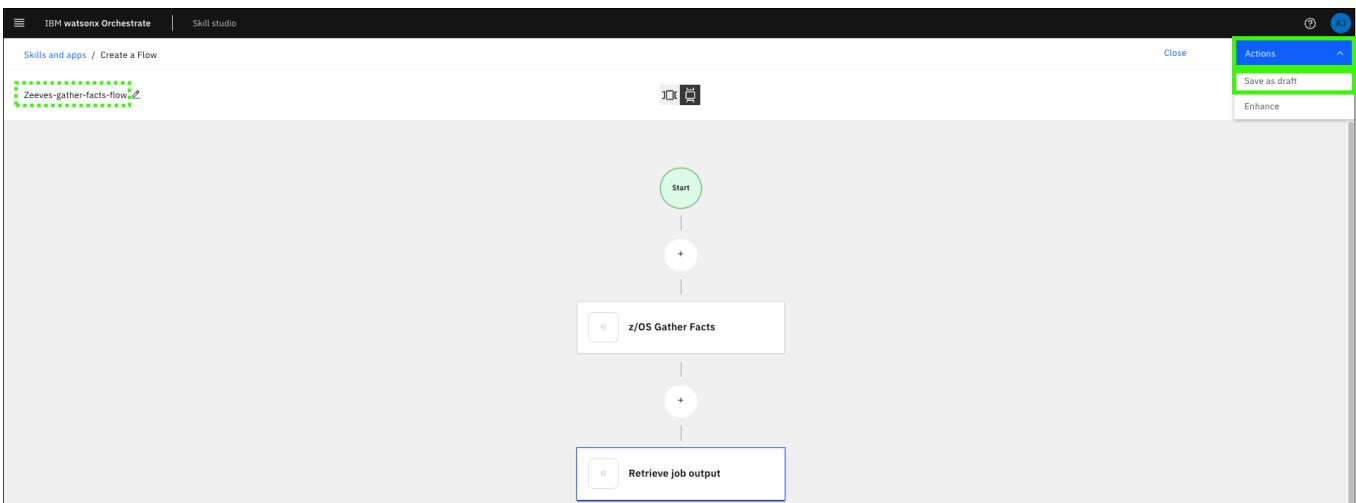
32. Click the pencil (>Edit).



33. Enter a (a) **Name** and (b) **Description** for your skill flow and then (c) click **Save**.



34. Expand the **Actions** pull-down list and click **Save as draft**.



35. Expand the **Actions** pull-down list and click **Enhance**.



Enhancing the skill flow

On the **Enhancing the skill** pages, you can: - modify the skill name, description, and version - add phrases (prompts) that will be recognized by the assistant to call the skill flow -

36. Click the **Phrases** tab.

37. Enter new **phrases** (prompts) for your skill flow and then click **Publish**.

Notice the default prompts are either not very intuitive (the skill flow name) or a bit verbose. Add a couple of other prompts that you anticipate users will enter.

Example prompts:

Show me z/OS facts

Gather and display z/OS facts

Enable the skill flow in your assistant

38. Click AI assistant builder in the main menu.

The screenshot shows the IBM Watsonx Orchestrate interface. The left sidebar has a dark theme with several sections: Chat, Skill sets, Skill catalog, BUILD (with AI assistant builder highlighted), ADMINISTER, and Access management. The main area is titled "Skill studio" and contains a message: "Objects to build and manage skill flows and connectors." Below this is a table with four rows of skill entries:

Step in the process	Status	Skill type	Author	Last edited
Ready to use	Published	Skill flow	andrew@jones-tx.com	November 19 2024
Ready to use	Published	Imported	andrew@jones-tx.com	November 19 2024
Ready to use	Published	Imported	andrew@jones-tx.com	November 19 2024
Ready to use	Published	Imported	andrew@jones-tx.com	November 19 2024

A green notification bar at the top right says "Published successful Published skill Zeeves-gather-facts-flow. 15:18:32". A "Configure prebuilt skills" button is also visible.

39. Hover over the Home () and click Actions.

The screenshot shows the "Actions" section of the AI assistant builder. The left sidebar includes Home, Build, Generative AI, and Actions (which is selected). The main area displays a series of cards with actions:

- Build actions: Enhance and improve your assistant's actions.
- Customize your greeting: Welcome your users in a unique way that aligns with your brand.
- Create a fallback plan: Train your assistant to adapt to specific situations.
- Preview & debug: Enable debug mode when you preview your assistant.
- Customize web chat: Design how your assistant shows up in your live environment.
- Set up a channel: Select the channels you want to deploy your assistant to.
- Set up live agent: Integrate live agents to transfer your users to when you edit web chat.

Below these cards is a diagram showing a "Default behavior" node branching into "General purpose" and "Conversational search" nodes, which then lead to a "Search" node.

40. Click New action.

The screenshot shows the "Actions" list interface. The left sidebar includes Actions, All items, Created by you (which is selected), Variables, and Saved responses. The main area lists actions under "Created by you /":

Name	Last edited	Examples count	Steps count	Status
z/OS Gather Facts	2 hours ago	2	0	Green circle

A green "New action +" button is located in the top right corner of the list area.

41. Click the Skill-based action tile.

Create an action

What kind of action do you want to build?

- AI-guided action**: Configure AI to generate responses based on specific use cases. Beta
- Skill-based action**: Build an action that helps perform a task.
- Custom-built action**: Design your own action step-by-step.

Not sure how to start? Try adding pre-built action templates.

42. Click the skill flow you created earlier and then click **Next**.

Build an action from a skill

Select a skill

Choose a conversational skill published as a foundation of your action.

Search a skill

Zeeves-gather-facts-flow	Retrieve job output	z/OS Ping	z/OS Gather Facts	Summarize the Webex meeting transcript
Run z/OS gather facts and retrieve the output. Last updated: 2024-11-19T21:18:31.793Z	z skills - Retrieve job output by job Id Last updated: 2024-11-19T20:08:59.538Z	z skills - This playbook pings the z/OS host to test connectivity. Last updated: 2024-11-19T15:58:20.567Z	z skills - This sample playbook demonstrates the z/OS gather facts module, which pulls z/OS-specific information from the z/OS host. Last updated: 2024-11-19T15:56:26.843Z	in watsonx.ai Last updated: 2024-11-04T10:49:16.502Z
Summarize the Box content in watsonx.ai Last updated: 2024-11-04T10:49:12.077Z	Summarize a Zendesk ticket in watsonx.ai Last updated: 2024-11-04T10:49:09.476Z	Summarize a Servicenow incident in watsonx.ai Last updated: 2024-11-04T10:49:05.828Z	Summarize a Salesforce opportunity in watsonx.ai Last updated: 2024-11-04T10:49:01.769Z	Sharepoint document summary in watsonx.ai Last updated: 2024-11-04T10:48:55.707Z
Salesloft email summary	Salesforce case summarization	Salesforce case sentiment analyze	Outlook email summary	Github issue summarization

43. Enter an example prompt for the skill and click **Save**.

You can use one of the prompts you used earlier for the skill flow.

Show me z/OS facts

Add example phrases:

Configuration

Example: I want to pay my credit card bill.

Skill Name: Zeeves-gather-facts-flow

Id: composite_skill_1.0.0_e1c26359857a488eb71e3cb541e...

Enter phrases that a customer types or says to start the conversation about a specific topic. These phrases determine the task, problem, or question your customer has.

The more phrases you enter, the better your assistant can recognize what the customer wants.

Enter phrases your customer might use to start this action

Total: 0

Example: I want to pay my credit card bill.

New action

What does your customer say to start this interaction?

Show me z/OS facts

Cancel Save

44. Enter any additional phrases (prompts) and then click the **save** (💾).

The screenshot shows the 'AI assistant builder' tab selected in the top navigation bar. The main area displays the configuration for a skill named 'Zeeves-gather-facts-flow'. In the 'Add example phrases' section, there is a text input field containing the phrase 'Show me z/OS facts'. A green dashed box highlights this input field.

45. Click Preview.

The screenshot shows the same configuration screen as above, but with a green box highlighting the 'Preview' button located at the bottom right of the interface.

46. Enter one of the prompts you specified into the assistant preview.

The screenshot shows the configuration screen with the 'Preview' button highlighted. On the right, a preview window is open, showing a conversation log. The log includes a message from a customer at 3:24 PM with the text 'Greet customer [default]' and a response from the assistant: 'Welcome, how can I assist you?'. A green box highlights the message 'Show me z/OS facts' in the input field of the preview window.

47. Click **Apply** when the first form is returned.

48. Review the results from the skill flow.

Use both scroll bars in the assistant preview to review all of the returned information.

The scenario shown above may or may not be relevant for your client's use case. It is intended to show how you to sequence skills together in a skill flow to create an action that your assistant triggers based on prompts using the pre-configured Ansible automation templates. You are encouraged to experiment with your own skill flows and prompts using other skills available within the [AAP instance](#).

Next, learn about custom-built actions.

Creating custom-built actions

To this point, you have learned how to: - import skills into Watson Assistant - add applications with those skills to your assistant - create skill-based actions for your assistant - combine skills in a skill flow

There is also the ability to create **custom-built** actions. Creating custom-built actions allows you to create a new action with different steps to take in the conversation, forming a sequence of prompts that define the conversation experience. The steps can be defined with or without conditions, which help you control the customer responses. Steps within the custom action can end with routing to conversational search, triggering another existing sub-action, and other actions. This is a powerful way of customizing the end-user's experience.

Learn more about creating custom-built actions [here](#).

Importing pre-packaged z/OS skills

Provided with Version 2 of watsonx Assistant for Z is a set of pre-packaged skills which can be used to automate various tasks on z/OS, such as running different console commands and retrieving logs from batch jobs.

The list of pre-packaged skills available include: - Authorized program list - z/OS IPL Information - Display zOS parmlib datasets - Unix System services options - Display zOS subsystems - List spool files - Retrieve dataset content - Retrieve spool file content - Retrieve z/OS Management Facility (OSMF) job status

IBM watsonx Orchestrate requires that any OSMF environment you connect to for skill execution has certificate authority (CA) signed certificates. In the case of the Ansible Automation Platform (AAP) & Wazi z/OS environment provisioned in IBM Technology Zone (ITZ), the z/OS system is not currently using CA signed certificates and therefore won't be able to execute the pre-packaged skills on your own z/OS system. For demo purposes, it is still recommended to import them so that the pre-packaged skills can be shown.

Work is in progress to modify the ITZ environments so that they are enabled for skill execution using these pre-packaged skills. In the meantime, the underlying automation for these pre-packaged skills can still be demonstrated using the Ansible skill 'z/OS Operator Command' which is available to import. This works because the pre-packaged skills are executing console commands directly using OSMF APIs which can also be run using the Ansible template skill 'z/OS Operator command'. For example, here are the console commands being used in some of the pre-packaged skills:

- Authorized Program list – `operator command -> d prog,lnklist`
- z/OS IPL Information – `operator command -> d iplinfo`
- Display zOS parmlib datasets – `operator command -> d parmlib`

You can import the pre-packaged skills into your sandbox environment by downloading the .zip file from [here](#) and following [these instructions](#).

You must extract the imbedded JSON file and modify the file for your environment by following [these instructions](#).

Publishing and deploying your assistant

To this point, acting as an Assistant Builder, you have built out the assistant, configured conversational search, and added skills and automations. While doing so, you have been testing your assistant using the **preview** capability of AI Assistant Builder. The **preview** capability is a closed environment for experimenting with prompts.

Once your assistant is finalized, you are able to publish it to make it available to end-users. Each assistant you create comes with two **environments**: *draft* and *live*. At this point, you have likely been configuring your assistant in the draft environment. Each environment has its own set of IDs, URLs, and service credentials that can be referenced by external services.

The **Environments** page in the AI assistant builder has tabs for managing the Draft environment and the Live environment:

This screenshot shows the 'Environments' page in the AI assistant builder. The 'Draft' tab is selected, indicated by a green dashed border. The 'Live' tab is also present. On the left, there's a sidebar with icons for environments, channels, and extensions. The main area has sections for 'Draft environment', 'Channels' (listing 'Web chat'), 'Resolution Methods', 'Draft content' (with a note about last edit and draft status), and 'Extensions' (listing 'Search'). A callout box highlights the 'Draft content' section.

This screenshot shows the 'Environments' page in the AI assistant builder with the 'Live' tab selected, indicated by a green dashed border. The 'Draft' tab is also present. The layout is similar to the Draft environment page, with sections for 'Live environment', 'Channels' (listing 'Web chat'), 'Resolution Methods', 'Content' (noting no published version), and 'Extensions' (listing 'Search'). A callout box highlights the 'Content' section.

The Draft environment contains all your in-progress work in the Actions, Preview, and Publish pages. Use the Draft environment tab to build out your assistant and use for internal testing before deployment. Any integrations you use (i.e. channels) for the draft environment are unique to that environment, and changes to draft integrations don't affect the live environment.

Publish the assistant

Each time that you publish, you're creating a new version of the assistant, for example V1. When you publish your content, you're creating a snapshot of the draft content, resulting in a version.



Versions do not contain integration configurations or environment settings

Published versions contain all of the content from actions, including settings and variables. **However, versions do not contain integration configurations or environment settings.** Integration configurations and environment settings must be configured manually in each environment.

For managing quality-control and versioning, the Live environment is the version of the assistant you should be giving access to the end-user.

Follow these steps to publish the first version of your assistant using Assistant Builder:

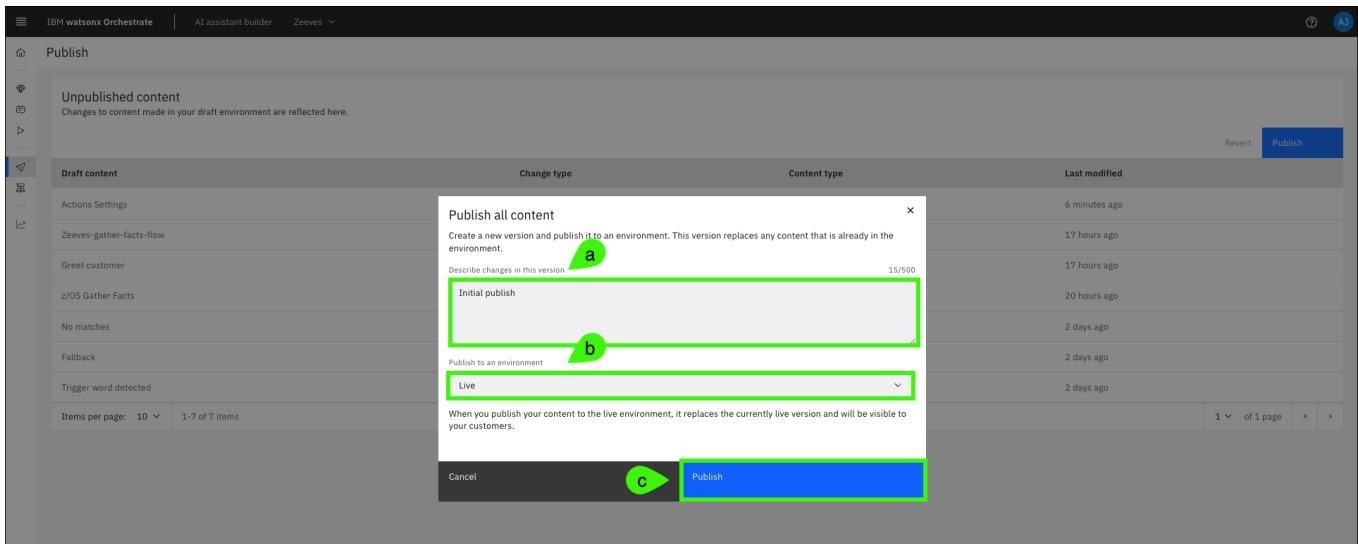
1. Hover over the **Home** icon () and click **Publish**.

The screenshot shows the IBM Watsonx Orchestrate interface with the 'AI assistant builder' tab selected. On the left, there's a sidebar with 'Generative AI', 'Actions', 'Preview', and a 'Publish' button which is highlighted with a green border. The main area displays several cards: 'Build actions', 'Customize your greeting', 'Create a fallback plan', 'Preview & debug', 'Customize web chat', 'Set up a channel', 'Set up live agent', and 'Public assist'. Below these cards is a diagram showing a 'Default behavior' node branching into 'General purpose' and 'Conversational search' nodes, both leading to a 'Search' node. At the bottom, there's a message about new Watsonx generative AI features and a 'What's new' link.

2. Click **Publish**.

The screenshot shows the 'Publish' page within the AI assistant builder. It displays a table of 'Draft content' with columns for 'Change type', 'Content type', and 'Last modified'. The table includes rows for 'Actions Settings', 'Zeeves-gather-facts-flow', 'Greet customer', 'z/OS Gather Facts', 'No matches', 'Fallback', and 'Trigger word detected'. At the top right of the table, the 'Publish' button is highlighted with a green border. The bottom of the page shows pagination controls and a note about items per page.

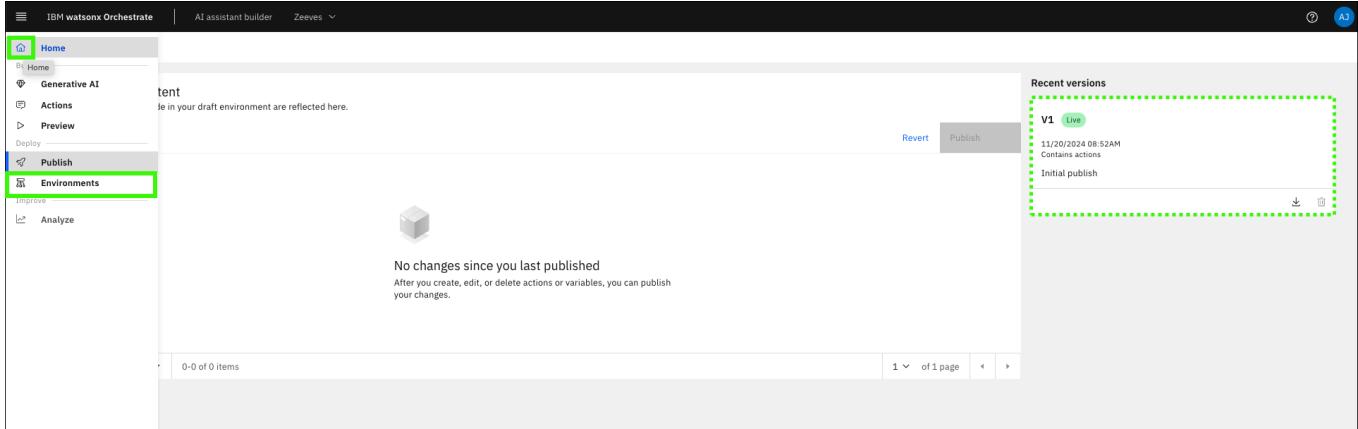
3. Enter a (a) description of the changes, set the (b) environment to **Live**, and then (c) click **Publish**.



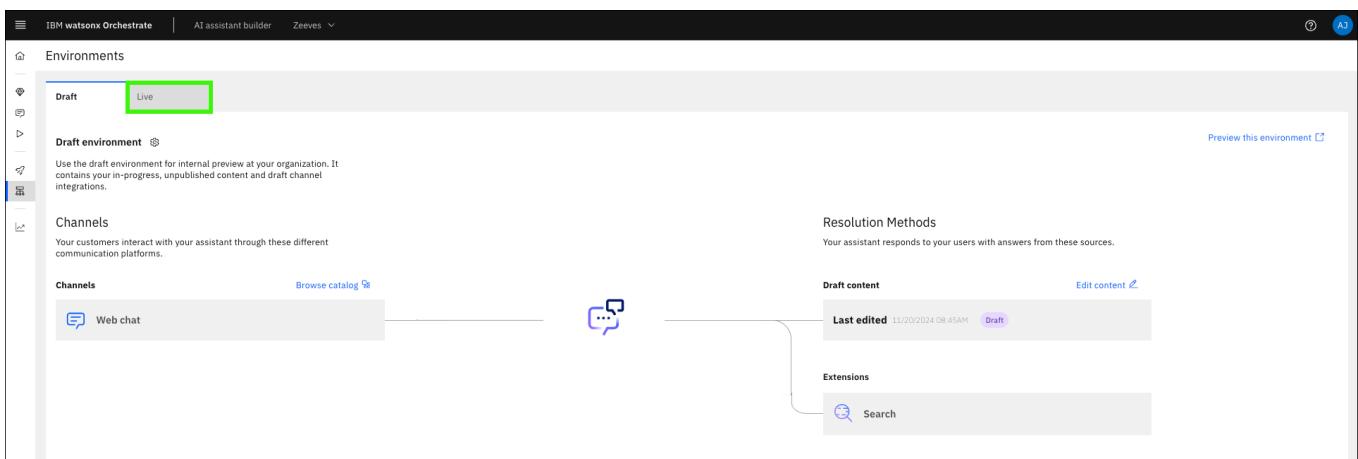
Configure the **Live** environment

When the live environment is created, the environment settings from the draft environment are not carried over (for example the configuration of the OpenSearch instance used for conversational search).

4. Hover over the **Home** icon (🏡) and click **Environments**.



5. Click **Live**.



6. Click Web chat.

The screenshot shows the 'Live' environment configuration in the IBM Watsonx Orchestrate interface. It includes sections for 'Channels' (with 'Web chat' highlighted), 'Resolution Methods', 'Content' (version V1, live), and 'Extensions' (Search). A green box highlights the 'Web chat' entry in the 'Channels' list.

7. Customize the Live assistant as you see fit.

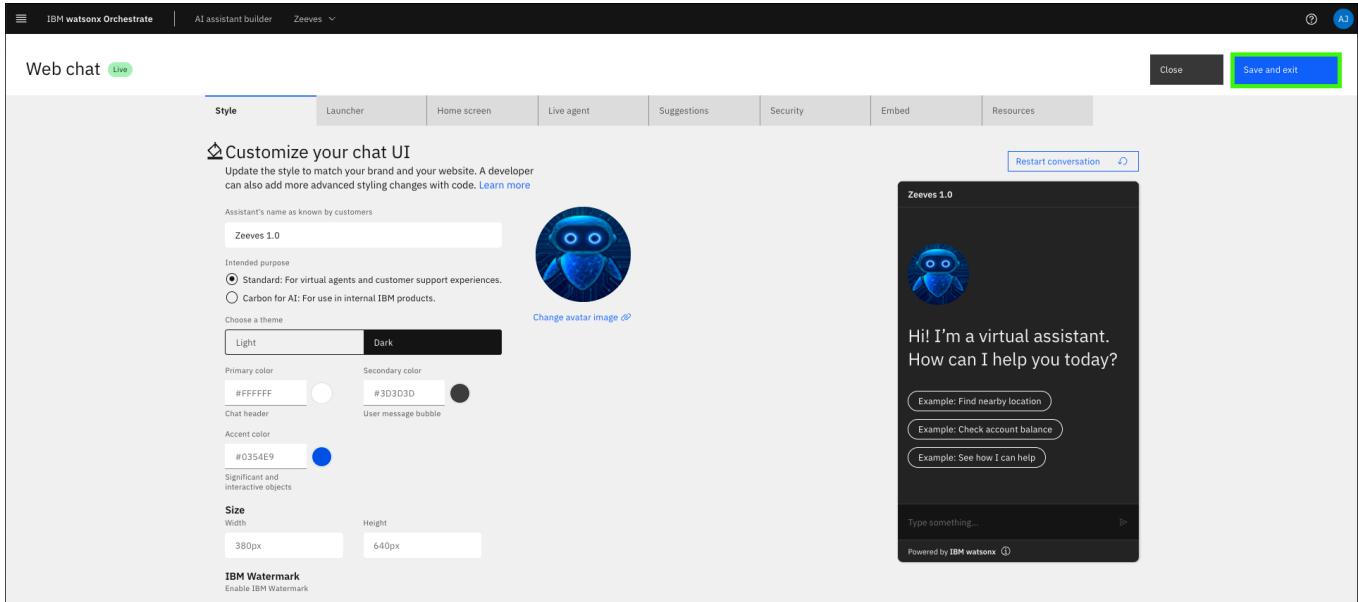
On the Style tab, you're able to set the Assistant name which will be displayed at the top of the chat window when end-users are interacting with the assistant. For pilots or demos, you may want to personalize this name for the client. Also in the Style tab, you have the ability set the themes and display settings of the chat windows, including the ability to enable the IBM Watermark and enable streaming (recommended).

On the Home screen tab, you enable and customize a default greeting message from the assistant when the user accesses the assistant chat. You're also able to set Conversation starters that will be displayed in the chat window. When selected by the end-user, the text of these conversation starters are sent as prompts, so it is important that your assistant is trained and tested to answer appropriately. It is highly recommended to remove these default conversation starters and to consider creating your own as long as they're able to be executed as actions and provide value to the end-user. At the bottom of the Home screen tab, you will also see the ability to add a Background style for the assistant chat window.

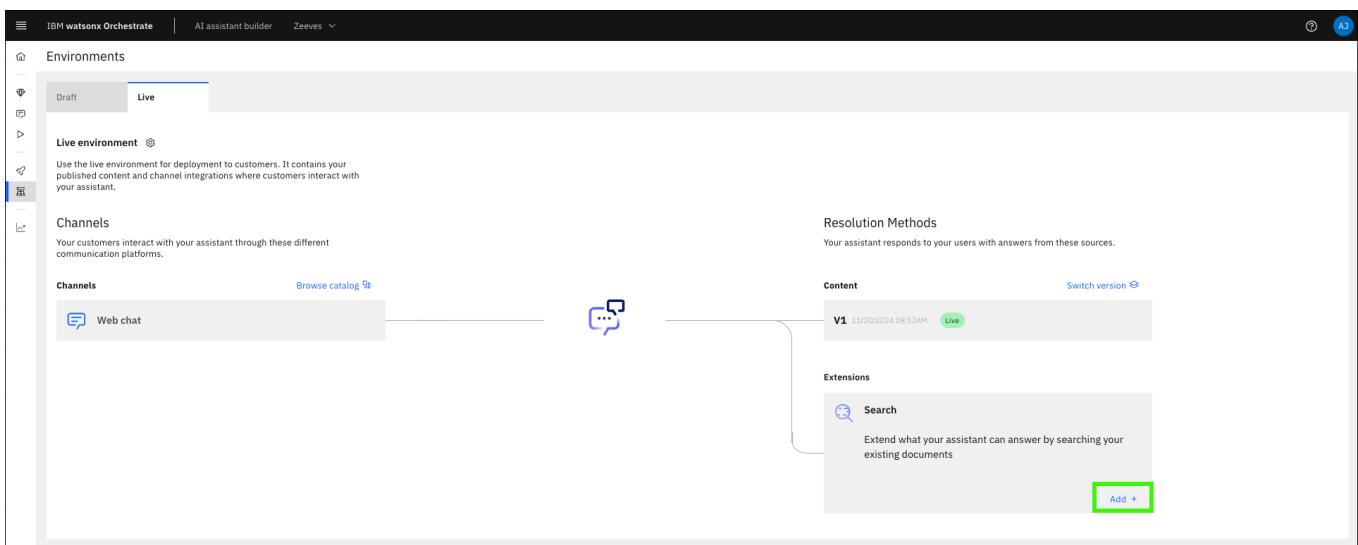
Explore all the other tabs.

The screenshot shows the 'Web chat' configuration page in IBM Watsonx Orchestrate, specifically the 'Style' tab. It includes fields for 'Assistant's name as known by customers' (Zeeves 1.0), 'Intended purpose' (Standard: For virtual agents and customer support experiences), theme selection (Light or Dark), color customization (Primary color #FFFFFF, Secondary color #3D3D3D, Chat header #0354E9, Accent color #0354E9), size settings (Width 380px, Height 640px), and watermark/streaming options. A preview window on the right shows the assistant's interface with a dark background, blue owl logo, and a greeting message: 'Hi! I'm a virtual assistant. How can I help you today?'. Buttons for 'Close', 'Save and exit', and 'Restart conversation' are visible at the top right.

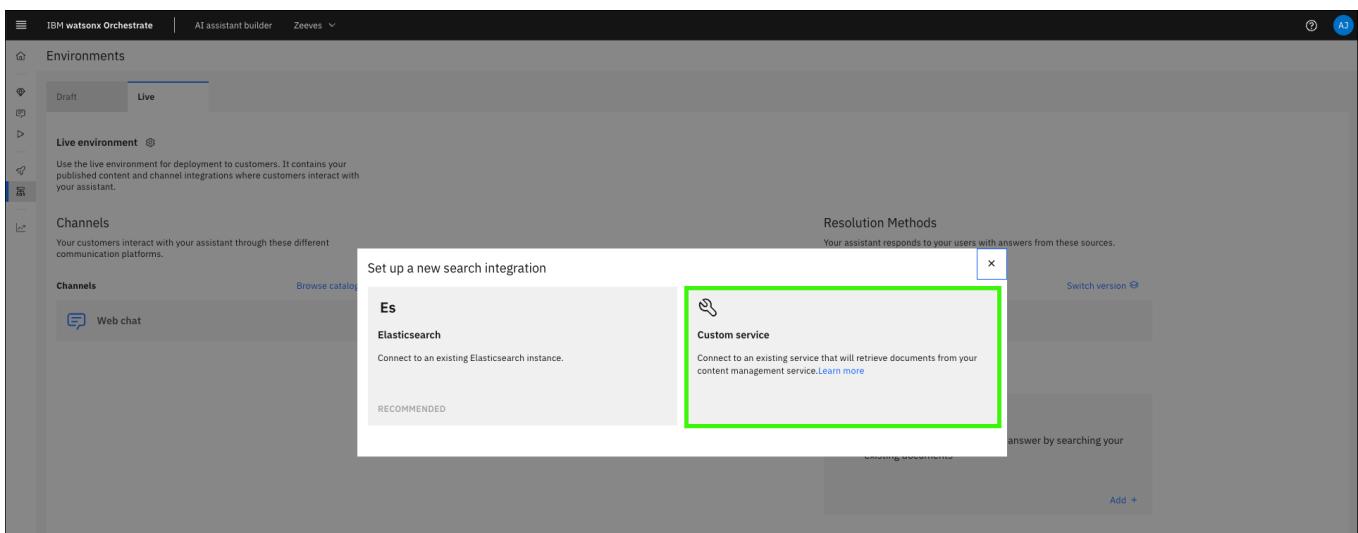
8. Click Save and exit.



9. Click Add in the Search tile.



10. Click Custom service.



11. Enter the (a) URL for your bring-your-own-search (BYOS) engine, (b) select **None** for the **authentication type**, and (c) click **Next**.



Use the correct URL!

Be sure to use the URL for your BYOS OpenSearch engine created earlier [here](#).

Custom service Live

Connect service Conversational search (optional)

Choose your search provider

By providing credentials Through your client

URL
https://wx4z-opensearch-wrapper-wxa4z-byos.apps.673cb83f0fe21364200dcb7.ocp.techz1

Choose an authentication type
None

Close Next

12. Verify conversational search is enabled and click **Save**.

Custom service Live

Connect service Conversational search (optional)

Enable conversational search (optional)

By enabling conversational search, your assistant can generate responses from supplied documents. Conversational search will use a watsonx generative AI model hosted in Dallas, TX US, to generate conversational responses. [Learn more](#)

On

By using this feature you agree to the [Pricing](#) and [Terms](#). This feature is not PCI validated. The use of this feature may impact your PCI compliance.

What's different?

When your assistant receives a message, the default search behavior returns short excerpts of the best search results. When conversational search is enabled, your assistant will instead generate a conversational response grounded in those search results.

Default search behavior (conversational search off)

Welcome, how can I assist you?
Can I use my points on airfare?

I searched my knowledge base and found this information which might be useful

We offer credit cards offer perks liability to our clients, protecting you from having to pay for unauthorized purchases. Can I use my points on airfare? Using points and miles to book your flights is a great option. Not only do you save money, but you can use your points to book...

IBM watsonx Lendy Demo Site

Close Save

13. Update the (a) conversational search settings, (b) click **Close**, and then (c) click **Save**.



Feel free to customize the settings.

This is your assistant. Feel free to customize the settings. The settings shown below reflect the changes made earlier in the lab guide to the draft version of the assistant. This includes the **Metadata** field to weigh ingested client documents higher.

Custom service

Settings **Instance**

By using this feature you agree to the [Pricing](#) and [Terms](#). This feature is not PCI validated. The use of this feature may impact your PCI compliance.

Conversational search

Use a watsonx generative AI model hosted in Dallas, TX US, to generate conversational responses.

Conversational search

On

Off

Single turn

The assistant uses only the current user input for retrieving search results and generating answers. This works well for clear, complete inputs but generally won't work with context-dependent questions like, "Why is that?" after a previous answer.

Entire conversation

The assistant uses the entire session history for retrieving search results and generating answers. This handles context-dependent questions well but may over-rely on past topics, even if the user has moved on.

Define the text for the creation title.

How do we know?

Tendency to say "I don't know"

Rarely Less often More often Most often

Generated response length

Concise Moderate **Verbose**

Default filter

Can be anything (plain text, JSON, etc.)

Metadata

```
[{"id": "id_weight", "value": "1"}, {"product_docs": "0.2", "customer_docs": "0.8", "base_indexing_idn_docs": "0.1", "customer_indices": "customer_**"}]
```

Learn more about publishing your assistant and creating live environments [here](#).

Deploy the assistant

After configuring your assistant's settings and publishing, the final step is to deploy your assistant which can be done across various channels depending on the use case.

There are several options for deploying your assistant through channels and integrations to satisfy the use cases that you might address. Learn more about all the deployment options [here](#).

For this lab, you will deploy the assistant using the web chat integration. The web chat integration provides an assistant interface that can integrate with your website and there's a lot of flexibility with how you may want to integrate it. Learn more about the web chat integration [here](#).

14. Click Web chat for the Live environment.

Environments

Draft **Live**

Live environment

Use the live environment for deployment to customers. It contains your published content and channel integrations where customers interact with your assistant.

Channels

Your customers interact with your assistant through these different communication platforms.

Channels **Browse catalog**

Web chat

Resolution Methods

Your assistant responds to your users with answers from these sources.

Content **Switch version**

V1 11/20/2024 08:52AM **Live**

Extensions

Search

15. Click the **Embed** tab.

The screenshot shows the 'Web chat' configuration page. The 'Embed' tab is selected and highlighted with a green border. On the right side, there is a preview window titled 'Zeeves 1.0' showing a dark-themed web chat interface with a blue owl logo and sample messages. The configuration options include:

- Assistant's name as known by customers:** Zeeves 1.0
- Intended purpose:** Standard: For virtual agents and customer support experiences.
- Choose a theme:** Light (selected) or Dark
- Primary color:** #FFFFFF
- Secondary color:** #3D3D3D
- Chat header:** #0354E9
- Accent color:** #0070C0
- Size:** Width: 380px, Height: 640px
- IBM Watermark:** Enabled (On)

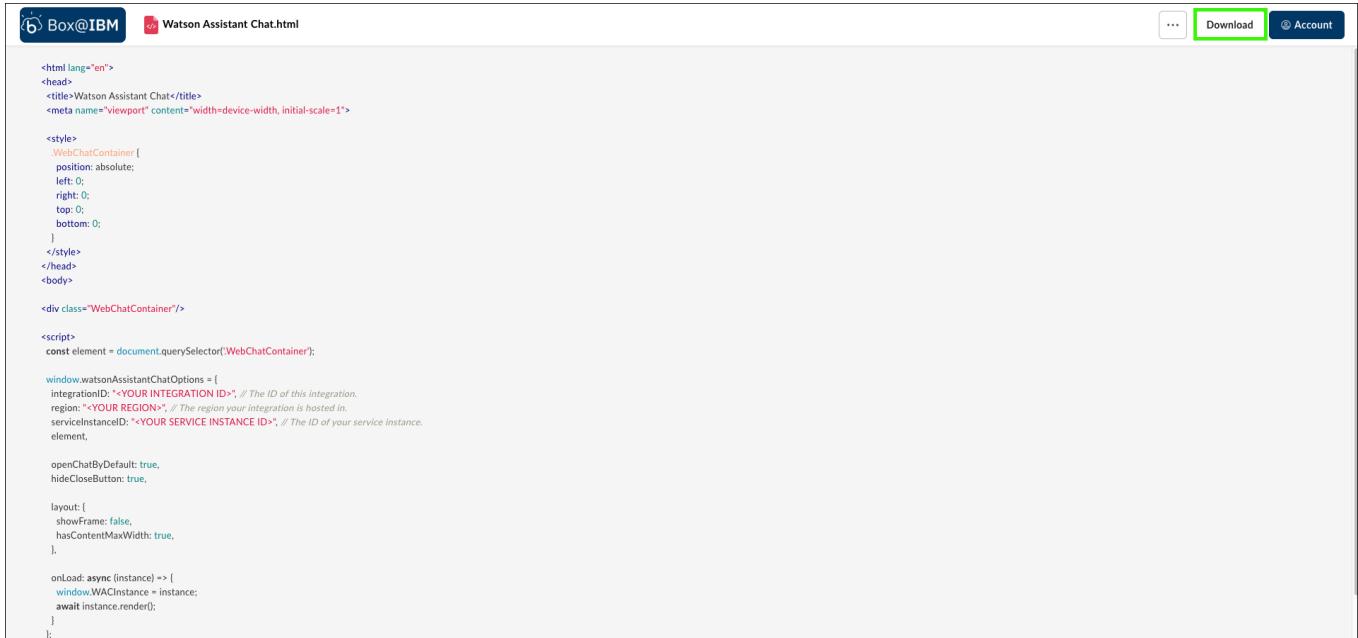
16. Copy and record the **integrationID**, **region**, and **serviceInstanceId** values.

The screenshot shows the 'Web chat' configuration page. The 'Embed' tab is selected and highlighted with a green border. A code snippet for embedding the chat is displayed in a code editor:

```
<script>
  window.watsonAssistantChatOptions = {
    integrationID: '8b4ad299-e6b9-434c-8c8c-2754d9272fe5', // The ID of this integration.
    region: 'wco-us-south', // The region your integration is hosted in.
    serviceInstanceId: '7c1e3381-8dcf-476b-bdf2-319f531087ba', // The ID of your service instance.
    onLoad: async (instance) => { await instance.render(); }
  };
  setTimeout(function() {
    const t=document.createElement('script');
    t.src='https://web-chat.global.assistant.watson.appdomain.cloud/versions/' + (window.watsonAssistantChatOptions.version || 'latest') + '/index.js';
    document.head.appendChild(t);
  });
</script>
```

17. Click the link to download a sample web chat hyper text markup language (HTML) page.

[Watson Assistant Chat.html](#)



```

<html lang="en">
<head>
<title>Watson Assistant Chat</title>
<meta name="viewport" content="width=device-width, initial-scale=1">

<style>
.WebChatContainer {
  position: absolute;
  left: 0;
  right: 0;
  top: 0;
  bottom: 0;
}
</style>
</head>
<body>

<div class="WebChatContainer">

<script>
const element = document.querySelector('.WebChatContainer');

window.watsonAssistantChatOptions = {
  integrationID: "<YOUR INTEGRATION ID>", // The ID of this integration.
  region: "<YOUR REGION>", // The region your integration is hosted in.
  serviceInstanceId: "<YOUR SERVICE INSTANCE ID>", // The ID of your service instance.
  element,
}

openChatByDefault: true,
hideCloseButton: true,

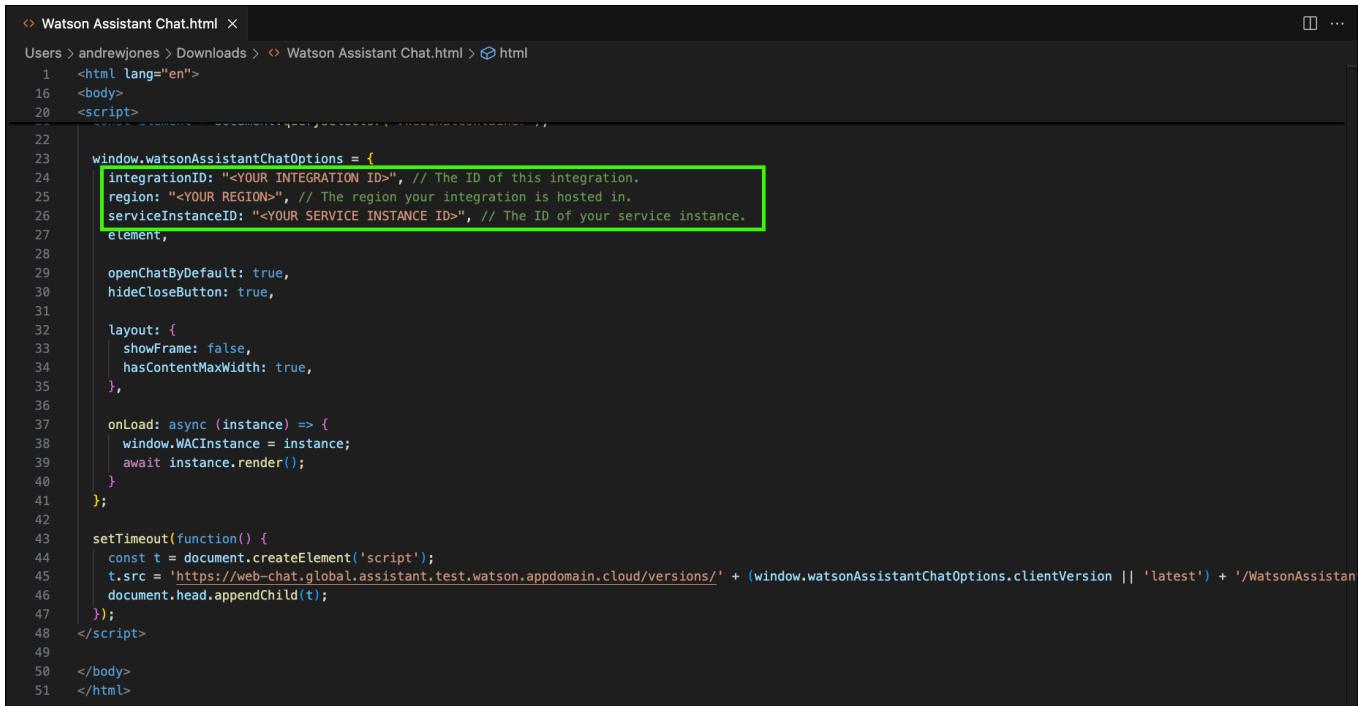
layout: {
  showFrame: false,
  hasContentMaxWidth: true,
},

onLoad: async (instance) => {
  window.WACInstance = instance;
  await instance.render();
}
</script>

</div>
</body>
</html>

```

18. Edit the Watson Assistant Chat.html file and set the values for your assistant you copied in step 16.



```

<div>Watson Assistant Chat.html <...></div>
Users > andrewjones > Downloads > <div>Watson Assistant Chat.html <...></div> > html
1  <html lang="en">
16 <body>
20 <script>
22
23   window.watsonAssistantChatOptions = {
24     integrationID: "<YOUR INTEGRATION ID>", // The ID of this integration.
25     region: "<YOUR REGION>", // The region your integration is hosted in.
26     serviceInstanceId: "<YOUR SERVICE INSTANCE ID>", // The ID of your service instance.
27     element,
28
29     openChatByDefault: true,
30     hideCloseButton: true,
31
32     layout: {
33       showFrame: false,
34       hasContentMaxWidth: true,
35     },
36
37     onLoad: async (instance) => {
38       window.WACInstance = instance;
39       await instance.render();
40     }
41   };
42
43   setTimeout(function() {
44     const t = document.createElement('script');
45     t.src = 'https://web-chat.global.assistant.test.watson.appdomain.cloud/versions/' + (window.watsonAssistantChatOptions.clientVersion || 'latest') + '/WatsonAssistan
46     document.head.appendChild(t);
47   });
48 </script>
49
50 </body>
51 </html>

```

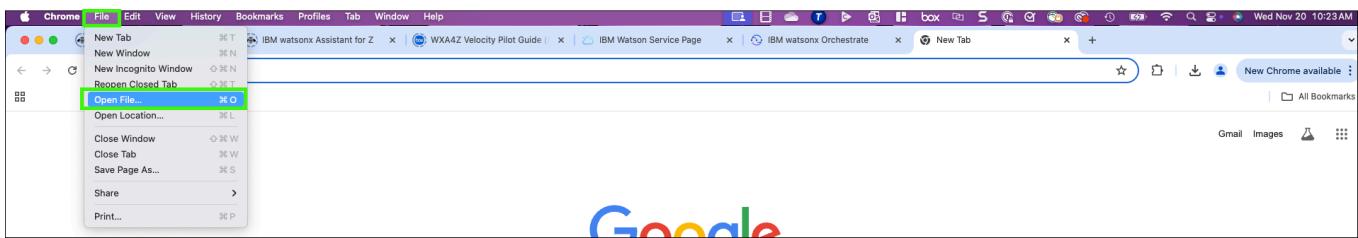
19. Save the changes to the Watson Assistant Chat.html file.

```

Users > andrewjones > Downloads > Watson Assistant Chat.html > html > body > script > layout
 1  <html lang="en">
 2  <head>
 6   <style>
 7     .WebChatContainer {
11    top: 0;
12    bottom: 0;
13  }
14  </style>
15  </head>
16  <body>
17
18  <div class="WebChatContainer"/>
19
20  <script>
21  const element = document.querySelector('.WebChatContainer');
22
23  window.watsonAssistantChatOptions = {
24    integrationID: "8b4ad299-e6b9-434c-8c8c-2754d9272fe5", // The ID of this integration.
25    region: "wxo-us-south", // The region your integration is hosted in.
26    serviceInstanceId: "c1c163381-8dcf-476b-bdf2-319f531087be", // The ID of your service instance.
27    element,
28
29    openChatByDefault: true,
30    hideCloseButton: true,
31
32    layout: {
33      showFrame: false,
34      hasContentMaxWidth: true,
35    },
36
37    onLoad: async (instance) => {
38      window.WACInstance = instance;
39      await instance.render();
40    }
}

```

20. Open the Watson Assistant Chat.html file in a web browser.



Your assistant is now live. Explore some of the earlier prompts to verify the assistant is accessing the ingested documents and your skills and skill flows are active.

Prompts to try:

What is z/OS continuous delivery?

Get z/OS facts

Show me z/OS facts

Gather and display z/OS facts

Adding other integrations

Other integrations can be added to your assistant to allow you to publish your assistant to the channels where your end-users are likely to interact with the assistant.

To do so, channel integrations must be added. By default, a web chat integration is created, allowing an assistant to be embedded in a website. Other channel integrations are available in the Integrations catalog inside the AI assistant builder interface. Learn more about adding integrations [here](#).

Within the AI assistant builder, click on the Integrations page:

Scroll down to **Channels** to view the integration deployment options for your assistant.

Clicking on any of the available channels will provide guidance on how to configure the integration.

Next steps

This IBM watsonx Assistant for Z for Technical Sales Level 4 demonstrated several of the IBM watsonx Assistant for Z capabilities. However, there is still more to learn about IBM watsonx Assistant for Z. Be sure to check out the [offering landing page](#) and the [product documentation](#).

Also, the focus of the IBM watsonx Assistant for Z for Technical Sales Level 4 Lab Guide was the *end-user* use of IBM watsonx Assistant for Z. To learn more about deploying, configuring, and customizing IBM watsonx Assistant for Z explore the other capabilities that are included in the [ITZ collection](#) [here](#). Included are guidance for running client pilots and the ability to manage an IBM watsonx Assistant for Z deployment. In the Velocity Pilot environment, you can explore all the capabilities of the offering like creating an assistant, importing documents, and adding automation skills.

Earn the badge

Finally, remember to complete the IBM watsonx Assistant for Z for Technical Sales Level 4 learning plan to earn the IBM watsonx Assistant for Z Technical Sales Advanced badge.

IBM technical sellers must complete a Stand and Deliver exercise to validate their ability to demonstrate IBM watsonx Assistant for Z to potential clients. Refer to the evaluation checklist in the [YourLearning learning plan](#) for the Stand and Deliver.

Business Partners must pass the IBM watsonx Assistant for Z for Technical Sales Level 4 quiz that is located in the [IBM Training learning plan](#). The Business Partner quiz verifies use of the demonstration environment and click-through demonstrations, and not one's general knowledge of the offering.